

MISTER MIXER

Group 14



Srinilai Kankipati, Computer Engineering
Kevin Kurian, Computer Engineering
Michael Meyers, Electrical Engineering
Juni Yeom, Computer Science
Joshua Yu, Computer Engineering

Board Committee Members

John Aedo
Mark Maddox
Saleem Sahawneh
Enxia Zhang

November 26, 2024

Contents

1	Executive Summary	1
2	Project Description	2
2.1	Motivation and Background	2
2.2	Existing Product/Past Project/Prior Related Work	2
2.2.1	USB Audio Interfaces	4
2.3	Goals and Objectives	5
2.3.1	Basic Goals	5
2.3.2	Stretch Goals	5
2.3.3	Objectives	6
2.4	Key Engineering Specifications	6
2.5	House of Quality	7
2.6	Hardware Block Diagram	8
2.7	Software Architecture Diagram	10
3	Research	10
3.1	Hardware Research and Comparisons	11
3.1.1	Audio Signals	11
3.1.2	Audio Connectors	12
3.1.3	Preamplifiers	14
3.1.4	Signal Converters (A/D & D/A)	17
3.1.5	Choosing The Processor Type: SBCs, MCUs, FPGAs	23
3.1.6	Choosing the Single-Board Computer	25
3.1.7	Choosing the Microcontroller	27
3.1.8	Power Supply	30
3.1.9	Inclusion of the Rotary Encoder	36
3.1.10	Physical USB Pin Standards Comparison	38
3.1.11	LCD Touchscreen Comparison	40
3.2	Hardware Enclosure	43
3.3	Software Research and Comparison	45
3.3.1	Audio DSP Effects	45
3.3.2	FFT Signals: Importance and Application	51
3.3.3	Comparison of Communication Protocols	51
3.3.4	Analysis of Protocols for Raspberry Pi to Mobile Data Transmission	54
3.3.5	Wireless Data Level Protocols	60

3.3.6	Real-Time Processing Considerations	63
3.3.7	Choices of Runtime Environment on the Raspberry Pi	65
3.3.8	Comparison of Protocols for Communicating Audio Parameter Changes	66
3.3.9	User Interface Frameworks, Comparison and Selection	69
3.3.10	Automixing Methods, Comparison, and Selection	71
4	Standards and Design Constraints	86
4.1	Standards	87
4.1.1	Audio Hardware	87
4.1.2	Communication Protocols	87
4.1.3	JSON Extensions	92
4.1.4	Programming Languages	94
4.2	Design Constraints	95
4.2.1	Economic Constraints	96
4.2.2	Time Constraints	97
4.2.3	Hardware Constraints	98
4.2.4	Accessibility of Datasets	98
4.2.5	Social Constraints	99
5	Comparison of ChatGPT with Other Similar Platforms	99
5.1	Limitations, Pros, and Cons	99
5.2	Examples	101
5.2.1	Automation of Boilerplate Tasks	102
5.2.2	Information Retrieval	103
5.3	Conclusion	103
6	Hardware Design	104
6.0.1	Preamplification Circuitry	104
6.0.2	LDO Circuitry	107
6.0.3	Phantom Power Circuitry	108
6.0.4	ADC Circuitry	108
6.0.5	DAC Circuitry	109
6.0.6	Differential Outputs	111
6.0.7	Headphone Amplifier	113
6.0.8	Proper Grounding	114
6.1	Overall Hardware Design Schematic	114
7	Software Design	115
7.1	Software Requirements and Constraints	115
7.1.1	User Stories and Acceptance Criteria	116
7.1.2	Unified Modeling Language (UML) Diagrams	118
7.1.3	Use-Case Diagram	118
7.1.4	User Roles	118
7.1.5	Use-Cases Correspond to Software Modules	119

7.2	Development of Software Architecture	119
7.2.1	System Topology and Deployment	120
7.2.2	Managing Data Flows	121
7.2.3	Complete Description of Architecture	124
7.3	Concrete Design and Implementation Details	124
7.3.1	State Manager Process	124
7.3.2	State Command Standard Definition	126
7.3.3	Audio Manager Process	129
7.3.4	Automix Model Design, Training, & Deployment	131
8	System Fabrication/Prototype Construction	141
8.1	PCB Design Standardization	141
8.2	Power Supply PCB	142
8.3	Mixed Signal PCB	144
9	System Testing	146
9.1	Hardware Testing	146
9.2	Software Testing	149
9.2.1	Testing Methodology	149
9.2.2	Concrete Implementation of Testing	150
9.3	Overall Integration	152
9.3.1	Hardware Integration	152
9.3.2	Software Integration	153
10	Administrative Content	153
10.1	Budget Estimates	153
10.2	SD1 Documentation timeline	155
10.3	SD2 Design Timeline	155
10.4	Work Distributions	156
11	Conclusion	157
11.1	Progress Report	158
11.2	Lessons Learned	158
11.2.1	Experience in Architecting Complex Systems	158
11.2.2	Importance of Modularization	158
11.2.3	Effective Communication	159
A	Citations	160
B	Copyright	165
C	LLM Citations	169

List of Tables

2.1	Table of Proposed Design Specifications	7
3.1	Audio Connector Comparison	14
3.2	Standards for Phantom Power: IEC 61938:2018	16
3.3	Comparison of Preamp ICs	16
3.4	ADC Comparison Chart	20
3.5	DAC Comparison Chart	21
3.6	Headphone Amplifier Comparison Table	23
3.7	Processor Comparison Table	24
3.8	Single-Board Computer Comparison	27
3.9	Microcontroller Comparison	31
3.10	LDO Comparison Chart	34
3.11	Comparison Table of Linear vs. Switching Regulators	36
3.12	Boost Converter Comparison Table	36
3.13	Caption for the table	45
3.14	Table of Communication Protocols	59
3.15	Comparison of Flutter, Slint, and Qt	70
3.16	Summary of Classical Machine Learning Methods for Automixing	76
3.17	Comparison of Deep Learning and Classical Machine Learning for Automixing	78
3.18	Comparison of Automatic Mixing Methods	85
3.19	Comparison of Various Multitrack Audio Datasets for Automixing Courtesy of Steinmetz et al.	86
4.11	Identified Constraints	96
5.1	Comparison of ChatGPT, GitHub Copilot, and Google Gemini	102
7.1	Channel Parameters	127
7.2	Equalizer Parameters	128
7.3	Compressor Parameters	128
7.4	Reverb Parameters	128
7.5	Comparison of Deployment Methods and Hardware Acceleration Techniques for the DMC Model	141
10.1	Estimated Cost of Production	154
10.2	Project Documentation Milestone	155
10.3	Project Design Milestone	156
10.4	Duty Distribution	157

List of Figures

2.1	House of Quality Diagram	8
2.2	Proposed High-level Hardware Block Diagram	9
2.3	High-Level Software Architecture Block Diagram	11
3.1	TR and TRS Wiring Key	13
3.2	XLR to TRS Connection Diagram	14
3.3	High-level Phantom Power Schematic, courtesy of Sweetwater	15
3.4	Example of Aliasing in Triangle Wave, created by Fraïssé and licensed under CC BY-SA 3.0 (https://creativecommons.org/licenses/by-sa/3.0/)	18
3.5	Architechture of N-bit SAR ADC, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.	18
3.6	4-bit SAR ADC example operation, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.	19
3.7	Binary Weighted DAC Example Circuit	20
3.8	High Level Flow of Sigma-Delta DAC, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.	21
3.9	Basic Linear Regulator, Reproduced with Permission. © Renesas Electronics Corporation or one of its subsidiaries. All Rights Reserved.	33
3.10	Example Application of TI’s Linear Regulators, Courtesy of Texas Instruments Inc	34
3.11	Example Switching Regulator Schematic, Reproduced with Permission. © Renesas Electronics Corporation or one of its subsidiaries. All Rights Reserved.	35
3.12	3D Render of the Mixer	44
3.13	The Addition Operator	46
3.14	The Multiplication Operator	46
3.15	The Delay Operator	46
3.16	A General Biquadratic Filter, reprinted with permission from miniDSP	48
3.17	Fast Fourier Transform Algorithm	51
3.18	WebSocket Connection Steps	55
3.19	HTTP/REST API Connection Steps	56
3.20	MQTT Connection Steps	57
3.21	CoAP Connection Steps	58

3.22 Wi-Fi Routing Explanation	61
3.23 BLE Connection Steps	62
3.24 Flow diagram of an adaptive audio effect	72
3.25 Architecture of [1] applying Reverb. Feature Extraction and Model training Training are Decoupled.	77
3.26 Visualization of Direct Transformation courtesy of Steinmetz et al.	82
3.27 Visualization of Parameter Estimation with Audio Loss courtesy of Steinmetz et al.	83
5.1 Prompt and Output from ChatGPT 4o	101
6.1 Schematic of Microphone/Line Preamplifier	105
6.2 Schematic of Instrument Preamplifier	106
6.3 Schematic of the three LDOs used in power circuitry	107
6.4 36V Generator and Phantom Power enabled XLR circuit	108
6.5 Schematic of one ADC	109
6.6 Schematic of DAC circuit and Post-DAC Low Pass Filter	110
6.7 Simulation Schematic used to find frequency response	111
6.8 Frequency response of Post-DAC Low Pass Filter	111
6.9 Schematic of Synthesized Differential Output	112
6.10 Schematic of Headphone Amplifier and Line/Headphone Output	113
6.11 Example Organization of Mixed Signal Circuit	114
6.12 Overall Schematic of Entire Circuit	115
7.1 Use-case diagram for mixer software	119
7.2 Deployment diagram of computing hardware	120
7.3 Modules Linked Via Publisher-Subscriber Pattern	125
7.4 Packetization of FFT Signals	130
7.5 Preprocessing Pipeline	133
7.6 Controller Network - DMC	134
7.7 Post Processor - DMC	134
7.8 Transformation Network - DMC	136
8.1 Schematic View of Power Circuitry PCB	143
8.2 3D View of Power Circuitry PCB	144
8.3 Schematic View of Main Audio PCB	145
8.4 3D view of Main Audio PCB	146
9.1 SSM2019 Preamplifier Circuit	147
9.2 SSM2019 Preamplifier Circuit Output	147
9.3 OPA2134 Preamplifier Circuit	148
9.4 OPA2134 Preamplifier Circuit Output	148

Chapter 1

Executive Summary

As students who are also in bands, we are constantly searching for ways to make the set-up process more streamlined, whether this be through buying more gear or learning new workflows. However, at the end of the day, we run into problems with pricing and modularity. Normally, an entire rig of equipment will set the musicians back multiple thousands of dollars, all while taking up large amounts of space and creating a mess of cabling. To solve these practical limitations, we envisioned a senior design project with the following goal - creating a low-cost and high performance audio mixer. Additionally, with the major boom in the usage of artificial intelligence, we were inspired to implement some novel AI features to streamline the workflow of audio mixing.

While audio mixers are ubiquitous, the most common (and subsequently lowest cost) are of the analog variety. Analog mixers are a tried-and-true technology that lack any option for configuration as well as being generally bulky. From experience, it is also relatively common to have errors with the mixing console during live settings, too, which is not ideal and introduces more problems into the loop. While analog mixers have some filters (low, mid, high) and effects (reverb, delay), anything more than that is only seen in the higher echelon professional gear. If a traveling musician needs to use their analog mixer to monitor and mix themselves, they would also need to travel with a bulky rack of pedals to achieve the various audio effects they desire on their instruments. This brings us to digital mixers, which while having more configurability, cost a fortune for even the lowest level of gear. This is not to mention that these digital mixers usually do not have touchscreen functionality and rather rely on faders entirely, as well as being run by low-powered computers. The mixers are able to work due to the use of highly specialized DSP chips, which are both difficult to interface with and configure.

With this, we seek to combine the present powers of a low-cost general computer to create an all-in-solution for audio mixing that includes AI auto-mixing features, similar to those seen in Izotope Neutron 4. This device would allow for high-level mixing capabilities at a fraction of the price while also combining hitherto separate technologies. Utilizing many off-the-shelf components, we are attempting to create an approachable and modifiable device that will be able to be more robust and feature-rich than any standard analog or digital mixer. The computer itself will be programmed using a real-time operating system to maximize data processing speed, and a custom interface will be designed to ensure a smooth experience for the end user.

To conclude, we seek to create a project that advances the world of low-cost

audio mixing while including capabilities not yet on present devices through emulation of state-of-the-art techniques currently used in the audio engineering field. Through this project, we will have a deeper understanding of the fundamentals of audio engineering along with the creation of a usable and effective product for our future selves. By combining general computing audio technologies with mixers, MISTER MIXER aims to allow ourselves and other musicians to be able to produce a powerful device at a fraction of what present consumer devices cost.

Chapter 2

Project Description

2.1 Motivation and Background

As mentioned in the executive summary, music is an expensive and somewhat cumbersome hobby to partake in; there is a lot of equipment required. This extends to all facets of playing and performing. Many beginner bands have no feasible way to mix their audio, which ultimately leads to subpar, unmixed performances that affect audience enjoyment or even playing ability. Even if they possess some sort of cheap, analog mixer, they face limitations with creative effects, parameter recall, workflow intuitiveness, and bulkiness (especially if they need a pedalboard setup on top of the mixer).

These limitations were experienced by the authors themselves. As former members of a band, they would perform in miniature concerts on top of UCF garages. When multiple bands would perform, substantial intermissions would take place between performances due to the reconfiguration of instrument connections and the analog mixer board. While some overhead was unavoidable, a lot of it could have been negated with mixer presets, instant recall, and a streamlined mixing workflow. This led the authors to realize that there is a niche in the market for an all-in-one mixing solution for fledgling bands and performers.

2.2 Existing Product/Past Project/Prior Related Work

The closest existing product to our proposed device is the Behringer Flow 8. It features 8 analog inputs including XLR and line inputs, as well as a Bluetooth streaming input. It also features wireless control over Bluetooth with a mobile app

in addition to the analog controls. The mixer has the capability to apply several musical effects on send channels [2]. The main strength of this digital mixer is the low price for the amount of functionality provided – namely, the Bluetooth capabilities and competent control scheme. Even so, there is not much more to the device in terms of workflow features and creative effects. It lacks digital parameter recall. This is partly a design choice made due to physical limitation; it lacks expensive, digitally controlled flying faders and instead opts for analog faders which do not offer any digital state control. Additionally, the library of DSP effects is not very comprehensive. There is still a clear need for creative effects and pedalboards to be applied before any instrument enters the mixer. These limitations have spurred us to create something more useful for the small performing band. In comparison, our mixer would have digital recall, many more creative effects, and an auto-mix feature.

Another product similar to our proposed device is the Mackie MobileMix 8-channel mixer. It is generally geared towards smaller performances in outdoor or other mobile settings, similar to our project. This miniature mixer features 8 channels of varied analog input (with 2 microphone preamps), Bluetooth input, monitoring outputs, and a small selection of effects (2-band EQ, reverb on a send channel). The control scheme is fully analog; the board is entirely controlled with potentiometers and buttons [3]. Like the previous example, this mixer once again has a very enticing cost to performance ratio, but the capabilities are little more than that of a typical analog mixer. The Bluetooth capability is once again a very nice feature to have, but for over \$200, we believe a mixer is capable of much more functionality. In comparison, our mixer would have digital recall, more creative effects, and the auto-mix features. Other mixers similar to the MobileMix also have very few effects other than panning and EQ functionality.

This project attempts to augment this type of pre-existing product and bring it more in line with a higher-end mixer, all the while being similarly priced (\$230 for the Mackie mixer as reference price). The features that will remain in common include volume balancing/control, panning, analog and bluetooth inputs, and portable form factor. The areas that would be expanded on will be explained in the paragraph below. Here are several commercially available digital mixers that offer a comparable feature set to what we would like to propose. They are the Allen & Heath CQ-18T and the QSC Touchmix 8.

Several standout features of these mixers include frequency monitoring, wireless control with a tablet app, digital parameter recall, multiple auxiliary mixes, rudimentary auto-mix tools ("Anti-Feedback Wizard", "Gain Assistant"), and parametric equalizers. These mixers are very feature-rich compared to both the Mackie and Behringer mixers in computational power, workflow features, signal routing, and effects, but this comes with a significant cost increase, almost \$1000 over the others. Additionally, they are less portable (particularly the Allen & Heath CQ-18T). Ideally, our mixer would be comparable to these digital mixers with the same price point and form factor as the Mackie mixer, all with the addition of the new auto-mixing features as well as a control scheme that mimics those of Digital Audio Workstations. This would create an option in the market for a mixer that is affordable, yet

extremely capable with a novel workflow that allows for greater creativity than with the typical analog mixer.

We will concede that several features found in high-end digital mixers must be omitted or pared back due to practical constraints. One of these is physical controls. Due to design time and cost constraints, almost all controls on our mixer will be through the touchscreen. There may be one rotary encoder, but we will lack the plethora of physical buttons and knobs that these competing mixers have. One other feature that would not be present is the capability to configure multiple different auxiliary and monitor mixes. We only have one line-out and headphone out, so this would not be possible to implement. Furthermore, even if we did decide to upgrade the amount of hardware outputs, we believe that would not significantly benefit the performances of small bands who do not have multiple PA systems and monitoring setups. This would simply incur extra cost with little to no benefit.

2.2.1 USB Audio Interfaces

A goal in our project is to utilize the mixer to have the ability to present a USB audio interface to the user to allow them a wider use of audio processing tools. USB audio interfaces are devices that act as a bridge between analog audio signals and a computer or digital audio workstation (DAW). DAWs are software platforms used for recording, editing, mixing, and mastering audio. The process is done by converting the analog signals into a digital format that can be processed by digital systems. The interface captures audio signals (e.g., from microphones, instruments, or mixers), converts them to a digital format using ADCs, and transmits them to the computer via a USB connection. Additionally, it can also convert digital audio data back into analog signals using DACs, allowing the computer to send audio to monitors, speakers, or headphones. They usually feature high quality audio converters and pre-amps, which can offer better sound quality compared to the built-in audio interface of a PC or digital device.

There are several products on the market that integrate a USB audio interface into their mixer, similar to the task that we are trying to accomplish.

The Tascam Model 16 features a built-in USB Audio Interface to send/receive (16-in/14-out) audio data to/from DAW [4]. One of the primary advantages of adding USB audio interface capabilities to our project, as demonstrated by the Model 16, is the seamless integration it provides with DAWs. Through a USB audio interface, audio signals from the mixer can be routed into the DAW for multi-track recording. This allows for high-fidelity recording for small bands. In addition, an attractive feature of the USB audio interface is that users can leverage the processing capabilities of the DAW to apply effects, EQ, reverb, or compression, beyond what our mixer's onboard processing may be able to provide.

The Soundcraft Signature 22 MTK offers an ultra-low-latency USB interface that flawlessly captures every channel, which can then be mixed or transferred to a DAW for further mix-down and production [5]. Another significant advantage to utilizing a USB audio interface is how it can facilitate real-time monitoring and playback of audio signals from the DAW or the computer. This allows users to

playback processed audio from the DAW through the mixer for further refinement or output to speakers, allowing flexibility between digital and analog audio processing.

2.3 Goals and Objectives

We desire to make a digital mixer that has a compact form factor, an abundance of creative effects, and an intuitive workflow that augments the creative vision of any group of performers. This will provide an all-in-one gigging solution for small bands. The goals needed to realize this vision are outlined below.

2.3.1 Basic Goals

For one, the mixer will provide basic mixing capabilities. Fundamentally, it must be able to take in audio signals, sum them, and output a mixed audio signal. On top of this basic functionality, the mixer will provide a library of creative DSP effects that can be applied to the audio signals. This will allow for a compact, all-in-one performing solution that removes the burden of carrying separate audio effects like guitar pedals.

For control, the mixer will have an intuitive touchscreen interface to easily control effects and routing. Additionally, the mixer will also be able to be wirelessly controlled through the use of a mobile app. This will allow for the mixing engineer to rapidly make mix decisions from the audience's perspective. The use of fully digital interfaces allows for retention of state (e.g. saving presets) across performances and will drastically speed up the mixing process.

Adding to the convenient control scheme, the mixer will have an easy-to-use auto-mix feature to automate setup and the mixing process. The feature will ideally give the performers a competent, great-sounding mix or, at the very least, provide a solid foundation upon which small creative adjustments can be made.

Finally, the mixer will be able to act as a USB audio interface. This will enable bands to record instruments and vocals directly into a Digital Audio Workstation. Such functionality greatly increases the value proposition of the project, eliminating the need for the band to buy a separate USB audio interface to record their performances.

2.3.2 Stretch Goals

The team has several stretch goals which may be explored after the conclusion of Senior Design. The stretch goals will increase the functionality of the mixer even further and make it a comprehensive solution for small bands. They are listed below.

One such goal is to support the use of external control interfaces such as foot switches and MIDI controllers. Normally, guitarists and other performers will adjust effect parameters on their pedalboards, keyboards, or other effect interfaces. This

mixer aims to eliminate the need for such bulky implements, although it does introduce the downside that such typical control interfaces are taken away from the performer. This feature will reintroduce such control to the performers and even augment the control performers have due to the power of digital configuration. External control will allow for performers to make adjustments on the fly and bring a new level of dynamism to their performances.

Another goal is to allow for audio connection to the mixer by Bluetooth. This will enable backing tracks to be played with ease from a performer's phone or other mobile device. This negates cumbersome routing and cable adapter setups whereby a phone must be connected through a headphone adapter, cable, and 3.5mm to 1/4" adapter into an analog input.

2.3.3 Objectives

The specific technical objectives below outline the implementation details required to achieve the desired functionality, with a focus on both essential features and additional enhancements.

There are several essential features that make up the feature set of a basic mixer. In this case, there will be 8 analog input channels with analog to digital conversion and preamplification to handle microphone, instrument, and line level inputs. For outputs, the mixer will possess both a line output to drive PA systems as well as a headphone output for personal monitoring. As for the actual mixing, the device will possess basic functions such as volume control, panning, and summing ("mixing") capabilities.

For the control scheme, as mentioned above, there will be both an onboard touchscreen LCD with a GUI as well as a wireless control scheme through a mobile app.

In order to augment the basic capabilities of a mixer, there will be a collection of DSP effects such as reverb, delay, compression, equalization, flanging, and distortion.

In order to expedite the mixing process, there will be a software auto-mix features that performs automatic volume balancing and equalization and applies creative DSP effects based on selected genre and detected inputs.

2.4 Key Engineering Specifications

Table 2.1 outlines the specifications targeted for this project. These specifications were established to ensure our digital mixer remains feature-rich compared to the consumer market options while also considering the time and money constraints of the project.

Attribute	Description
Dimensions	15" x 10" x 3"
Weight	<10 lbs
SNR (Signal to Noise Ratio)	>70 dB
Inputs	8 TRS, 4 XLR (w/ combo jack), 1 USB, Bluetooth
Outputs	1 Stereo Out, 1 Headphone-out, 1 USB
Input Gain	>50 dB
Digital Effects	Gain, Panning, Delay, Compression, Equalization, Flanger, Distortion, Reverb
Power Usage	<75W
Auto-Mixing	Automatically set levels, EQ, creative effects, subjectively improved over an unmixed performance. Qualitative listening tests, stereo loss to evaluate performance.
Cost	<\$400
Overall System Latency	<20ms

Table 2.1: Table of Proposed Design Specifications

2.5 House of Quality

House of quality is incredibly useful for comparing marketing requirements against engineering requirements. Below, the house of quality for this project is presented.

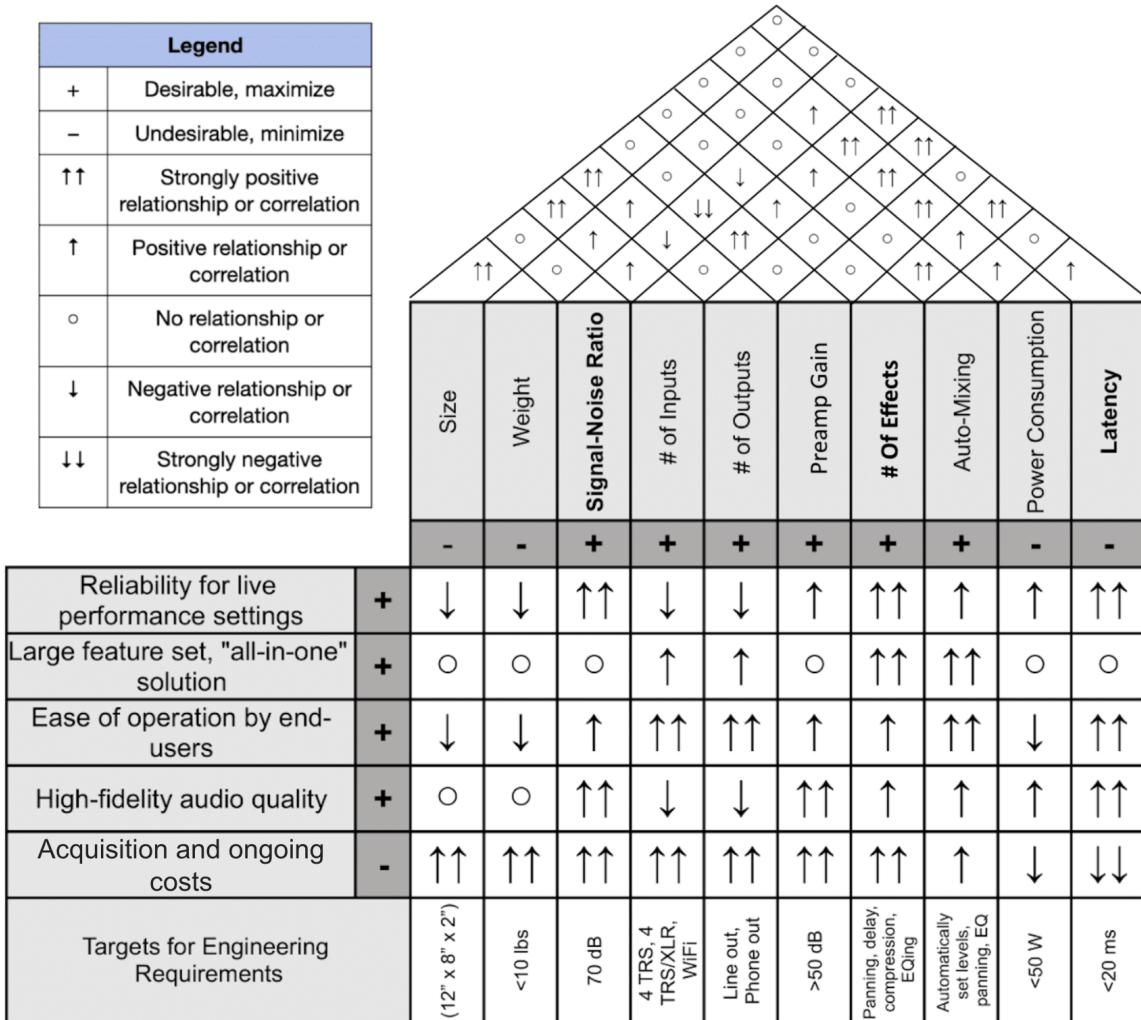


Figure 2.1: House of Quality Diagram

2.6 Hardware Block Diagram

Included below in Figure 2.2 is our proposed hardware diagram. The blocks are colored to indicate the main person responsible, and the color key is included on the diagram. The overall topology of our system is a dual-processor setup with one MCU to handle DSP for streams of audio and a Raspberry Pi 5 SBC to handle state management, wireless communications, and the AI auto-mix model.

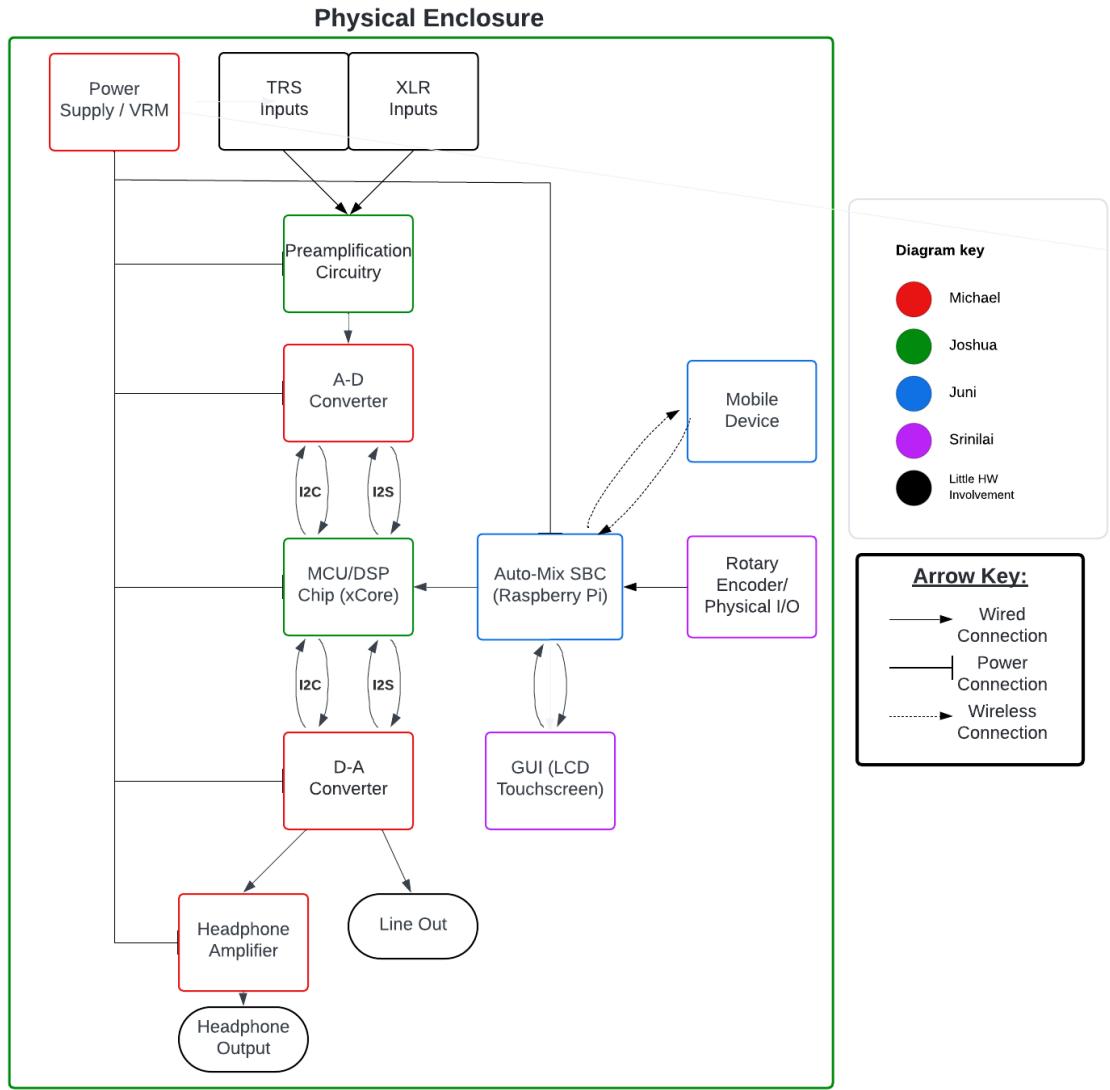


Figure 2.2: Proposed High-level Hardware Block Diagram

2.7 Software Architecture Diagram

Our prospective software architecture is illustrated in figure 2.3. Software components are distributed between the Raspberry Pi 5 and an XMOS xcore. The Raspberry Pi 5 is responsible for managing the user interfaces, state management, and higher-level audio processing tasks, while the xcore handles low-latency digital signal processing (DSP) operations for real-time audio effects.

The diagram is divided into two primary sections: the **Raspberry Pi 5** and the **XMOS xcore**. On the Raspberry Pi, there are several core modules:

- **State Manager** (blue) manages the mixer's central state information, ensuring all connected modules maintain consistency in settings like volume, panning, and effect parameters.
- **Audio Manager** (blue) oversees the flow of audio data between the Pi and the DSP system. It is responsible for handling audio streams, downsampling data for visualization, and passing raw audio to the auto-mixing model and the user interfaces.
- **Onboard Interface** (red) and **Mobile Interface** (red) are user interfaces connected to the manager modules via WebSockets connections. These interfaces allow real-time control over the mixer settings and audio effects.
- **Auto-Mixing Model** (green) automates various audio processing tasks, including volume balancing and applying effects based on AI model inferencing.

The **XMOS xcore** handles real-time audio DSP through the **DSP Manager** (yellow), which manages the chain of audio effects applied to the incoming audio streams. These effects are processed and then output to either speakers or external devices via USB. The connection between the Raspberry Pi and the xcore is established via **I²C** and **SPI**, allowing the state and audio data (respectively) to be transmitted between the two processing units in a synchronized manner.

Chapter 3

Research

All good engineering begins with research and it is the imperative responsibility of the team to research the components that make up the mixer. This chapter is split into two overarching sections, hardware research and software research. In

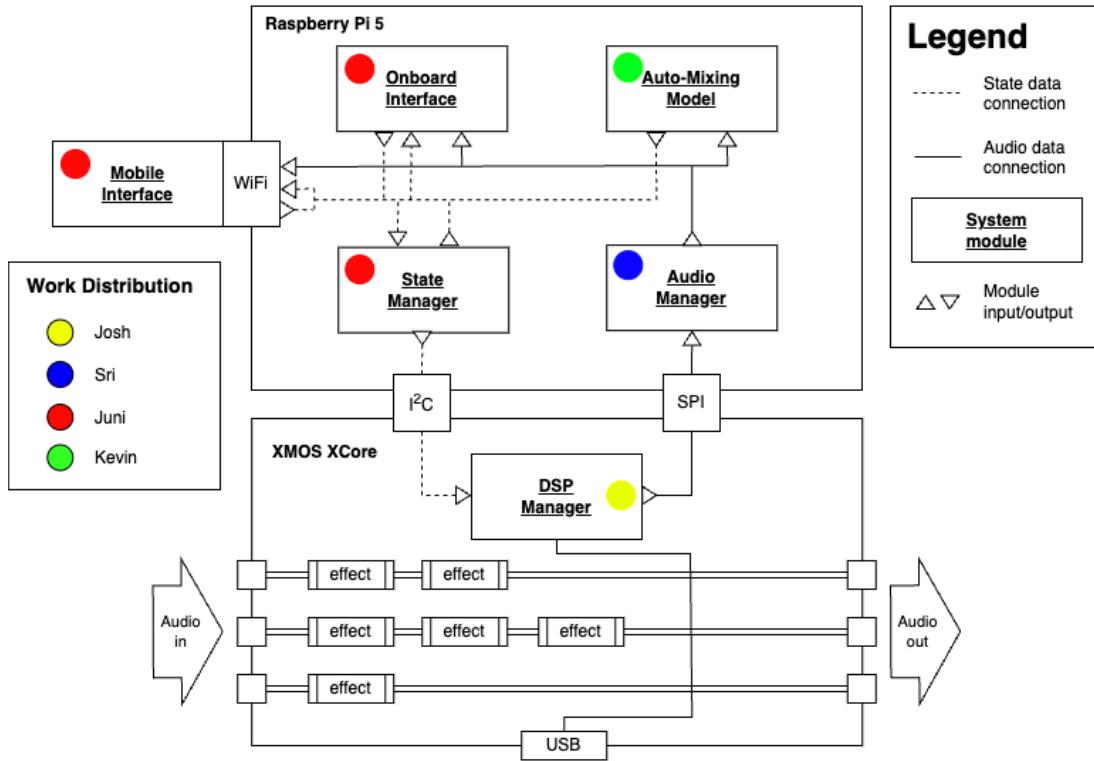


Figure 2.3: High-Level Software Architecture Block Diagram

each section, different technologies will be explained. Furthermore, comparisons will be made between specific parts. The selected parts will be highlighted and a justification will be given. The two sections are outlined below.

3.1 Hardware Research and Comparisons

3.1.1 Audio Signals

Before any circuitry work is to be done, identifying the types of analog audio connections that are required in a mixer is necessary for overall understanding. Differentiating between stereo and mono signals as well as balanced and unbalanced connections is important before the signal processing stage to ensure that the signal is transmitted properly. Additionally, we must name and differentiate the physical connectors that are to be used for the mixer, those being TS/TRS and XLR.

Signal Types: Mono vs. Stereo

There are several broad categories of audio signal types. One of these categories relates to the number of channels, that is, mono versus stereo.

Mono audio, as specified by the name, is a signal with a single channel of audio. It is by and large the standard input type for mixers and other professional audio

applications. Even in applications where stereo audio is needed, stereo audio is typically transmitted through two mono cables rather than a single stereo cable. Why this is the case will be touched on in the following paragraphs.

Stereo audio is a signal with two channels of audio, typically a left and a right channel. Stereo audio is useful because it allows an audio stream to have a sense of direction and width. Without delving too deep into psychoacoustics, stereo allows for audio to be placed in space by simply adjusting the balance of audio between the two channels. For example, if you listened to a pair of headphones, you would perceive a sound as coming from the right if it only played through the right driver. Similarly, you would perceive a sound as coming from the left if it only played through the left driver. It would come from the center if played in equal volume through both drivers.

As touched on above, the signal types that our mixer will be interfacing with are predominantly mono audio signals. The only notable exception is the main output and headphone output, which can transmit stereo. This is because the process of mixing generally follows a template where mono audio signals are "placed" into a stereo sound-field to produce an immersive listening experience. That is what our mixer will give artists the ability to do.

Signal Types: Unbalanced vs. Balanced

Another broad category of audio signal types related to noise-rejection properties, that is, unbalanced versus balanced signals.

Unbalanced signals only have two connections, a signal connection and a ground connection. The signal is simply transmitted and measured as the difference between the signal and ground terminals. Unbalanced connections are simple but prone to noise and interference. Thus, unbalanced signal cables must be restricted in length; the longer the cable, the more noise the signal will pick up.

Balanced signals have three connections, two signal connections and a ground connection. The key is that the two signal connections are wired so that they are 180 degrees out of phase with each other. At the end of the connection, the signals will then be phase-aligned again and summed. Any noise that is picked up by the two signal connections of the balanced cable will now be 180 degrees out of phase with its counterpart and removed when summed through destructive interference. This gives balanced cables excellent protection against noise; thus, they are preferred in professional audio environments. Balanced cable runs can be much longer.

3.1.2 Audio Connectors

Connector Types: TS/TRS

The most common audio connectors that are used are both the TS and TRS connections. TS connections consist of the Tip and Sleeve, whereas TRS connections have a wire for Tip, Ring, and Sleeve. As shown by Figure 3.1 featuring male TS

and TRS connectors, both connections look very similar, but TRS can be used for unbalanced stereo or balanced mono inputs whereas TS cables can only be used for unbalanced mono sound inputs [6]. These connections can either be 1/4" or 1/8", but the normal audio standard for input from instruments is 1/4".

Of note is that while TRS connectors have the necessary quantity of pins to connect to balanced microphones, they are not frequently utilized because of the segmented design of the connector. Many microphones require phantom power. If phantom power was applied across the hot and cold sections on the TRS connector, the sections would short as they were being plugged and unplugged.

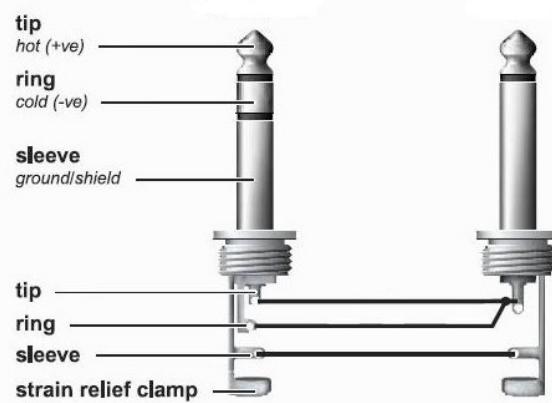


Figure 3.1: TR and TRS Wiring Key

Connector Types: XLR

The last connection of note for this project is an XLR connection, which is predominantly used for microphone signals. Similar to the TRS, it usually features 3 (but can have up to 7) pins, with the first pin always being dedicated to ground. Figure 3.2 demonstrates that the first pin of a female XLR connection is ground, the second pin is positive, and the third (bottom) pin is negative. With its three-pin topology, XLR cables are capable of carrying both balanced mono audio and unbalanced stereo audio, however, balanced mono audio is the typical use case. In male XLR connectors, pins 1 and 2 swap sides, but the connectors still have the same pattern [7]. Pins 2 and 3 on XLR connectors are capable of carrying a Phantom Power signal in devices such as a ribbon or condenser microphone that require power to operate. The separated pin design ensures that the phantom power is delivered safely without the possibility of any shorting, unlike TRS cables.

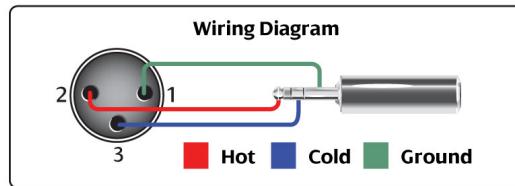


Figure 3.2: XLR to TRS Connection Diagram

Audio Connector Comparison

Below is a tabular comparison of all the covered audio connectors. All three will be utilized in the mixer's design.

Connector	TS	TRS	XLR
# Connections	3	3	2
Connection Capability	Unbalanced	Balanced	Balanced
Phantom Power	No	No	Yes

Table 3.1: Audio Connector Comparison

3.1.3 Preamplifiers

Input Levels

The process of preamplification must begin with an explanation of the purpose of preamplification as well as different input levels found in audio. Preamplification is the process of bringing up weak analog audio signals to a standardized "line" level such that they can be processed or further amplified (power amplification) by standard audio equipment. In the case of this project, all input audio must be brought up to line level so that they can be fed into the ADC circuitry.

What is line level? For audio, there are two general standards. One is consumer line level, which is typically measured to be -10dBV. Another is professional line level, which is measured to be around +4dBu [8]. As is described by the name, professional line level is what is typically used in professional production gear such as mixers and PA systems. As such, that is the level that our mixer preamplifiers must output to the ADCs. Professional line level measures on the order of magnitude of 1 volt, 1.3 volts RMS to be specific. One may notice the discrepancy in the measurement units of the consumer and professional line level. Consumer line level is expressed in dBV, while professional line level is expressed in dBu. dBV is measured in reference to 1 volt; 0 dBV is 1 volt. dBu is measured in reference to 0.775 volts; 0 dBu is 0.775 volts. This is an odd number, but makes sense in the right context. If you calculate the power dissipated by a 0 dBu signal through a 600 Ohm load (a common value for transmission lines), you end up with exactly 1mW of power [9].

Another common input level is instrument level. Instrument level is outputted by common instruments such as electric guitars and basses. There can be a wide range of levels, but a rough estimate of instrument level places it at -20 dBu [10], on the order of 100 mV. As such, the preamplifier will need to provide a moderate amount of amplification (20-30 dB) to produce an adequate signal for the ADC.

Beyond signal, another point of discussion is the input impedance. Instrument inputs are unique in that they are high-impedance, on the order of a million ohms. This is in contrast with the typical impedance of a microphone or line input, which is on the order of ten thousand ohms. As such, in order to prevent undesired loading and filtering side effects, the instrument input must have its own high-impedance signal path. They cannot just be plugged into a microphone or line preamplifier circuit with adjustable gain.

One last common input level is microphone level. Microphone level is extremely low, around -40 to -60 dBu. This is on the order of 1 mV to 10 mV. As such, the preamplifier will need to provide at least 40 dB of gain in order to amplify the microphone input to a usable signal level. This will be taken into account in the preamplifier design.

Phantom Power

One important consideration in a microphone preamplifier design is phantom power. Many microphones (typically condenser microphones) require a DC voltage to be applied from the preamp in order to operate. The most common phantom power voltage specification is 48 volts, although more specifications exist. There are also standards for 24 volt phantom power and 12 volt phantom power. The voltage is brought into the circuit with a pair of resistors defined by the specification and applied to pins 2 and 3 of an XLR connection. Below is a high-level schematic of a phantom power circuit.

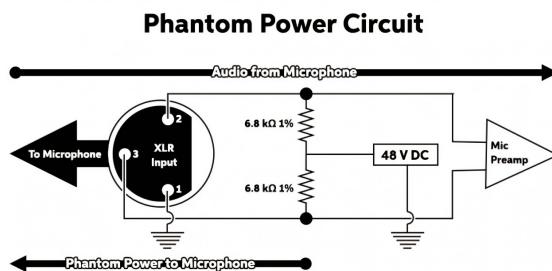


Figure 3.3: High-level Phantom Power Schematic, courtesy of Sweetwater

Most microphone designs today do not necessitate the use of the full 48 volts. In fact, most microphones step down the phantom power voltage and clamp it steady with a Zener diode at approximately 11 volts peak. However, for wider compatibility, the full 48 volts is used in most professional designs. In terms of current draw, the nominal current draw across all three phantom power specifications is quite low. For example, the 48 volt phantom power specification specifies a 10 mA

nominal current draw [11]. Below is a table summarizing the three phantom power specifications.

Phantom Pwr.	Volt., no load	I_{\max} (Shorted)	R_{feed}	Op. Current
P12	12V $\pm 1V$	35mA	680 680	17mA
P24	24V $\pm 4V$	40mA	1.2k 1.2k	20mA
P48	48V $\pm 4V$	14mA	6.8k 6.8k	10mA

Table 3.2: Standards for Phantom Power: IEC 61938:2018

One very interesting fact established through a design article is that the phantom power specification does not have to be obeyed at all. One can provide phantom power to microphones at non-standard voltages by simply adjusting the value of the resistors that bring in the DC phantom power voltage [12]. For example, one can provide a 15 volt phantom power supply by bringing the voltage in with 820 ohm resistors. This is an expedient design choice that can be exploited so that phantom power can be supplied from a wide range of voltages.

Preamplifier ICs

Operational amplifiers are the backbone of any preamplification circuit. There are a plethora of popular op-amps tailored for use in audio applications. They are particularly useful due to their differential inputs. Some popular audio pre-amps are compared below.

Preamplifier	SSM2019	THAT1512	OPA2134
# Channels	1	1	2
Max. Gain (dB)	70 dB	64 dB	120 dB
Supply Voltage (V)	$\pm 5V$ to $\pm 18V$	$\pm 5V$ to $\pm 20V$	$\pm 2.5V$ to $\pm 18V$
Noise f=1 kHz, G=1000)	1 nV/Hz	1 nV/Hz	8 nV/Hz
THD+N (f=1kHz, G=1000)	0.017%	0.008%	0.00008% (G=1)
Slew Rate (V/μs)	16 V/ μ s	19 V/ μ s	20 V/ μ s
CMRR (G=1000)	130 dB	120 dB	100 dB
Supply Current	± 4.6 mA	6 mA	4 mA
Output Swing ($\pm V_{out}$)	$\pm 13.9V$	$\pm 13.3V$	Approx. $\pm 16V$
Price	\$5.14	\$6.46	\$5.24

Table 3.3: Comparison of Preamp ICs

Preamplifier ICs have stellar performance across the market. They are usually not the bottleneck in any analog design. As such, the specifications are considered; however, the driving factor in preamplifier IC selection is the availability of

reference designs and robust documentation. An excellent reference design has been found that utilizes a dual-IC topology with an Analog Devices SSM2019 and a Texas Instruments OPA2134. The preamplifier design will be adapted from the reference design. As such, they are the preamplifier ICs of choice.

3.1.4 Signal Converters (A/D & D/A)

Arguably, the most important part of any audio equipment is the Signal Conversion chain – both converting the analog audio signals into digital ones for processing as well as recreating the output audio signal from digital to analog. Thus, understanding of the entire signal chain is necessary to the proper functionality of our digital mixer.

Analog-to-Digital Converters

In our mixer, the analog inputs in the system will be balanced Line Out signals TRS signals (mono) and XLR signals (stereo). After these are sent through the preamplifier circuitry (to allow the dynamics of the signals to be better identified), they must be sent into an ADC for Analog-to-Digital conversion. Since most signals coming into the mixer will be stereo signals (one channel for each side), we will be using an ADC that has stereo capabilities (both a left and right input channel).

Before talking about the technologies used in ADCs, the two most important parameters involved must be addressed: Signal-to-Noise Ratio (SNR) and bandwidth. SNR, as the name suggests, represents the ratio between the input signal (non-noise bits) and latent noise in the signal pathway, which is measured in dB. The higher the SNR, the more accurate the overall conversion process, and thus the audio signal will be less lossy. A typical SNR for a mixer will be ≈ 70 dB. Bandwidth is the range of frequencies that ADCs can accurately process. Due to the Nyquist-Shannon sampling theorem, which states that the maximum bandwidth cannot be more than one half of the sampling rate of the device without causing aliasing (distortion in the sampled signal) [13], as shown in Figure 3.4 with the triangle wave input. The higher the sampling rate, the higher the bandwidth, and the sampling rate is known as the Nyquist rate. For audio applications, the audible range is around 20-20000Hz, which means the sampling rate needs to be at least 40kHz to cover the entire frequency range, but the standard sampling rate is 44.1kHz to have a buffer for any extra frequencies. High fidelity audio (HiFi) nearly doubles this frequency to a total of 96KHz sampling rate, but there is not a real reason to do that with an audio mixer that is intended to be used live due to distortion from sources like speakers.

The bit number (degree of specificity for A/D conversion, given by a 2^n number of combinations) of an ADC heavily affects the SNR, with a higher number of bits giving a higher SNR. Sampling rate also affects the SNR, but much less than the number of bits. This is seen where the number of bits adds ≈ 6 dB per bit to the SNR, whereas doubling the sample rate only adds ≈ 3 dB to the SNR.

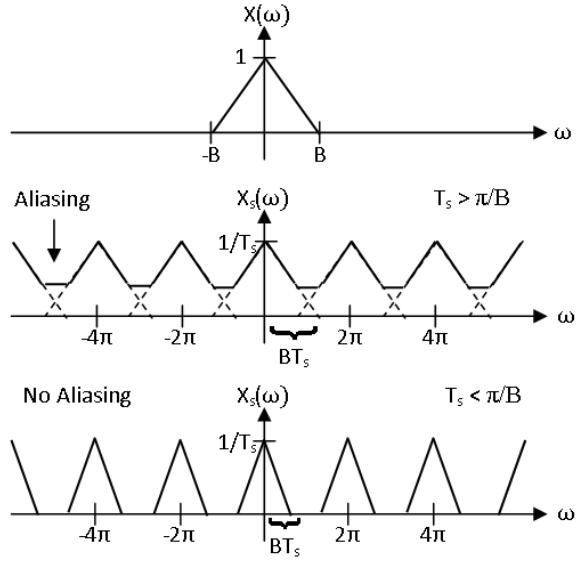


Figure 3.4: Example of Aliasing in Triangle Wave, created by Fraïssé and licensed under CC BY-SA 3.0 (<https://creativecommons.org/licenses/by-sa/3.0/>)

There are two popular methods of A-D conversion: Successive Approximation (SAR) and Sigma-Delta Conversion. SAR is an older technology that uses comparators between the input signal and internal DAC to approximate (hence the name) the voltage, resulting in an identified voltage of specified resolution depending on the number of bits [14]. This method is inherently noisier than Sigma-Delta conversion, though it can be done at low power in a small form factor.

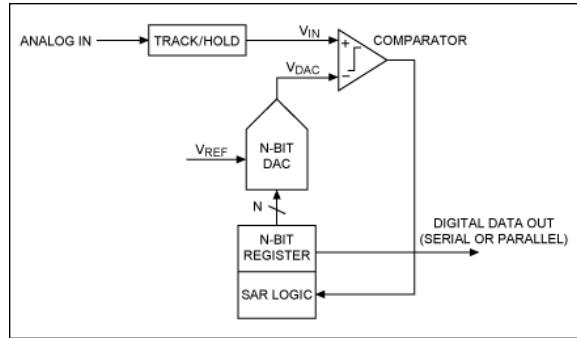


Figure 3.5: Architechture of N-bit SAR ADC, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.

As shown in Figure 3.5, the architecture inside of a SAR ADC is relatively simple, with a reference voltage being used to supply the DAC, which each bit adding another comparison operation to the ADC. Due to the intrinsic reliance on the internal DAC for conversions, it must be properly calibrated for accurate functionality. Successful operation of voltage comparisons of a 4-bit SAR can be shown below in Figure 3.6 [15]. Since these devices require successive operations, they are relatively slow compared to their sample rates (in the magnitude of megasamples per

second), but this does not affect our operations as audio information is relatively slow.

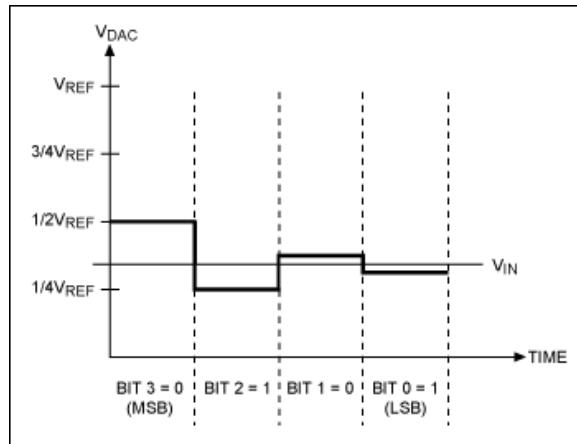


Figure 3.6: 4-bit SAR ADC example operation, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.

The internal operations of Sigma-Delta converters are relatively the same, using an input DAC with reference voltages as feedback to be summed with the input signal via a switched capacitor circuit [16]. Sigma-Delta converters offer higher resolution capabilities in addition to less noise during conversion at the cost of reduced bandwidth. Normal Sigma-Delta converters for audio applications may only use around 22KHz in bandwidth yet have sampling rates of magnitudes more, due to oversampling methods (usually 16 times the required sampling rate). Although the total usable sample rate is lower, there is no requirement for anti-aliasing and noise does not have much of an affect on the conversion (thus higher SNR). Because of its advantages with regards to signal integrity, our device will be using Sigma-Delta converters to maximize the overall SNR.

ADC Comparison

In choosing an ADC, we desired a resolution of at least bits and minimum sampling rate of 44.1K while having an SNR of over 90. Due to the mixing needing 8 audio inputs, finding an ADC that had maximal input channels while supporting I²S was paramount. In the commercial industry, Cirrus Logic ADCs are used most of the time. With that being said, due to the cost requirements, we chose the PCM1865 to utilize for the Analog-to-Digital conversion. This ADC also includes build in preamplification gain, which allows for more configurability of the signals post-preamp, and can operate on a single 3.3V supply. Daisy-chaining two of the PCM1865's together using I²C allows the two devices to communicate via TDM, allowing eight channels of audio to be sent to a single input pin on the DSP chip. The low cost and configurability made it an easy choice to choose for this project.

Part	Resolution	Sampling Rate	SNR	Price
PCM1865PWG4	24 bits	96kHz	110 dB	\$5.10
PCM1851APJT	24 bits	96k	101 dB	\$3.21
PCM1822IRTER	32 bits	8k - 192k	111 dB	\$3.67
CS5334-KS	20 bits	2kHz - 50kHz	100 dB	\$3.96
CS5335-KS	20 bits	2kHz - 50kHz	105 dB	\$3.96
CS5330A-KSR	18 bits	2kHz - 50kHz	90 dB	\$6.03
AD1877JR-REEL	16 bits	48k	90 dB	\$9.21

Table 3.4: ADC Comparison Chart

Digital-to-Analog Converters

After signal processing has taken place inside the microcontrollers, the digital audio signals must be converted back to analog signals to be able to be amplified and ultimately heard. Fortunately, the inner workings of DACs are relatively simpler than ADCs (as evidenced by the fact that many use internal DACs in comparator circuitry) and are robust for audio applications.

One of the simplest and textbook DAC solutions is using a resistor ladder e.g. $R, 2R, 4R$ in a summing amplifier to obtain the output voltage, where the most significant bit would be at the resistor corresponding to R (and whether or not it is an open or closed circuit depends on the bit value being 1 or 0 at that point). Figure 3.7 shows this simple method of D/A conversion, but it is easy to see the quick downsides to this method. The most glaring issue is that if each subsequent bit requires 2x the previous resistance, it exponentially increases to a point where it becomes unfeasible to process more than 4 bits of information – while this does provide a resolution of 2^4 or 16 possibilities, it is nowhere near sufficient for any audio application. Additionally, resistor tolerances make it so that getting exact ratios between the base and subsequent resistors is nigh impossible, let alone the fact that standardized resistor values usually do not come in multiples of 2.

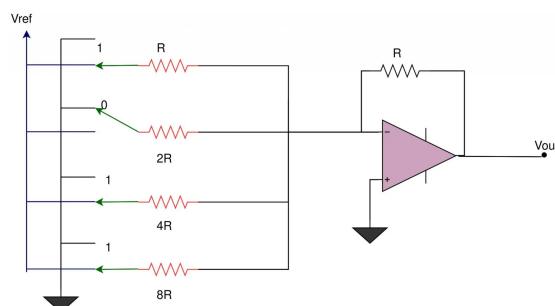


Figure 3.7: Binary Weighted DAC Example Circuit

Realistically, most DACs, like ADCs, will utilize Sigma-Delta techniques as it

provides a high SNR and is not reliant on a physical comparison using a resistor ladder. As shown by Figure 3.8, Sigma-Delta DACs are heavily digitized, first going through an interpolation filter and then modulating followed by conversion. The overall process is designed to minimize noise (interpolation filter padding most of the data with zeroes, Sigma-Delta modulator acting as a low-pass filter to data, and external LPF). This is also aided by the oversampling method mentioned in the ADC section. Just as most modern ADCs utilize Sigma-Delta conversion, most DACs also utilize Sigma-Delta conversion due to the high quality functionality it provides [17].

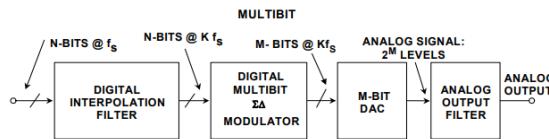


Figure 3.8: High Level Flow of Sigma-Delta DAC, Analog Devices, Inc. (“ADI”), © 2024. All Rights Reserved. These images are reproduced with permission by ADI.

DAC Comparison

The main desires for our chosen DAC are a SNR over 90 dB, a resolution of 24 bits, and an output of 1 stereo audio channel as the signal can be split for line-out and headphone channel because we are using a headphone amplifier. Due to only needing one DAC, budgetary concerns were not particularly present and thus most of the decision went into the quality of the product as well as ease of use. Cirrus Logic DACs are also used most of the time for commercial applications, but the PCM1780 aligns more closely with the other audio circuitry used (made by Texas Instruments) and offers ease of use coupled with low price. It is with that we chose the PCM1780 to utilize for the Digital-to-Analog conversion.

Part Number	Resolution	Sampling Rate	# outputs	SNR	Price
WM8521CH9GED	24 bits	192KHz	1	98 dB	\$1.11
CS4340-KS	24 bits	96KHz	1	101 dB	\$2.00
PCM1780	24 bits	192KHz	1	106 dB	\$2.43
PCM1748KE	24 bits	100KHz	1	106 dB	\$2.69
CS4340A-KS	24 bits	192KHz	1	101 dB	\$2.74
CS4391A-KS	24 bits	192KHz	2	94 dB	\$3.08
PCM1742E	24 bits	200KHz	2	117 dB	\$3.18
CS4340-BSR	24 bits	96KHz	1	91 dB	\$3.73

Table 3.5: DAC Comparison Chart

Headphone Amplifier

Once the audio signal is converted back to an analog signal, one must amplify it to hear through headphones. To do this, a circuit similar to a pre-amplifier is used to boost the output voltage signal into an acceptable level for headphones to be used. This circuit also helps feed current to the headphones to power them (which is how the headphones output audio). While a headphone amplifier can be as simple as a biased transistor circuit, the highest-performance and easiest to use consist of Headphone Amplifier ICs. These ICs are specifically made for audio in mind and do not require much external circuitry, all while allowing reliable operation at specified ranges.

There are a few parameters that play important roles in choosing a headphone amplifier: load impedance and maximum output power. Load impedance is the impedance of the headphones that will be used for audio playback, and the higher the impedance, the lower the output current will be (which means that it will be easier for the amplifier to drive them). Headphones with lower impedances will be able to utilize more current, which ties into the power requirements. Headphone amplifiers are capable of supplying enough power for the audio up to a point, usually less than 100mW. This means that when choosing an amplifier, one must know the general range of the headphone impedances followed by the desired gain of the output voltage signals. An important specification to look at is the Total Harmonic Distortion (THD) vs output level curves, which provide a more accurate representation of the overall performance of a specific amplifier instead of the total maximum power output. There is also the topic of headphone sensitivity (usually in dB/mV) but this ranges greatly between different types of headphones and it is better to overestimate the headphone requirements than underestimate them. As such, choosing a headphone amplifier with a comfortable driving range is necessary [18].

Headphone Amplifier Comparison

The most important factors when choosing a headphone amplifier are both the maximum output power and the THD + Noise ratio for the output signal. As most amplifiers above 40mW are sufficient for headphones with 16-32 ohm impedances, we chose the TPA6100A2DR for offering a very low THD + Noise factor in addition to having an acceptable maximum output power, while being very inexpensive, shown in Table 3.6.

Part	Max Output Power	THD + Noise (%)	Price
LM4808MX	105 mW	0.05%	\$0.97
TPA6100A2DR	50 mW	0.04%	\$1.05
TPA6132A2RTER	25 mW	0.025%	\$1.30
LM4911MM	145 mW	0.1%	\$1.65

Table 3.6: Headphone Amplifier Comparison Table

3.1.5 Choosing The Processor Type: SBCs, MCUs, FPGAs

Prelude: Processor Requirements

In order to fulfill the feature requirements of the mixer, the processor must meet several key requirements. For one, it must have enough processing power to power DSP effects across at least eight channels of audio in real time. Additionally, it must have enough processing power and memory to run our AI-automix model. Most of the candidate models are quite sizeable (upwards of 2 million parameters) and require significant memory and processing capabilities to perform inference in an adequate timeframe. Additionally, the processor must feature adequate I/O to interface with at least eight channels of AD/DA conversion circuitry. Lastly, the processor must possess the ability (in and of itself or with peripherals) to communicate with another device over Bluetooth and/or WiFi protocols.

One large consideration is the use of SBCs (single-board computers), MCUs (micro-controller units) versus FPGAs (field-programmable gate arrays). Which is best suited for this project?

Single-Board Computers

Single-board computers are full-fledged, miniature computers. They contain one of multiple central processing units (CPUs), random-access memory (RAM), and a whole array of peripheral connectivity. They typically run full operating systems that manage the onboard resources, allowing for parallel operation of many programs at once. They are also typically much more powerful than microcontrollers or FPGAs.

Microcontrollers

Microcontrollers are essentially small computers. They contain one or multiple central processing units (CPUs), random-access memory (RAM), and input and output peripherals all in a single integrated circuit (IC). Functionally, they are designed to perform single, repeated tasks and they offer limited support for parallelism; however, this can be expanded with the use of real-time operating systems (RTOS). In terms of programming, microcontrollers typically utilize high-level system languages like C and C++ [19].

Field-Programmable Gate Arrays

FPGAs are integrated circuits that offer a collection of configurable logic blocks (CLBs) and a way to configure the interconnects between these logic blocks. Fundamentally, field-programmable gate arrays offer the developer the ability to reconfigure the hardware, offering increased flexibility and performance over microcontrollers. This also allows for massive parallelism due to the configurability of the logic blocks.

However, this flexibility comes at a cost in regards to programming complexity. FPGAs are not programmed in a high-level language like C or C++ but rather a hardware design language like VHDL, Verilog, or SystemVerilog. These languages are much more difficult to program with; they fall under a completely different level of abstraction compared to high-level languages. One has to think in terms of low-level digital design, which is a completely different skill-set than procedural or object-oriented programming.

Additionally, FPGAs are often not complete systems. Peripheral interfaces are not integrated with the chip and thus must be manually developed/configured. The development cycle involves a lot more work.

SBCs, MCUs, and FPGAs: Conclusion

In order to facilitate rapid prototyping and development, microcontrollers are the preferred choice for this project for audio DSP. In addition to being easier to program, microcontrollers offer integrated peripherals that are vital to our project such as serial communication interfaces for AD/DA converters and GPIO for external device control.

There are some concerns with parallelism and performance due to our multi-channel audio requirements; however, our decision is bolstered by the ubiquity of microcontrollers for DSP in the professional audio sphere.

Additionally, because of the power requirements of the auto-mix model, a single-board computer will be necessary. A single-board computer will also provide a simpler development environment with the availability of Python interpreters and high-level development environments. The findings are summarized in a table below.

Processor	Single-Board Computer	Microcontroller	FPGA
Computation	Very powerful	Low Power	Moderate Power
Parallelism	Yes	Limited	Yes
Software	Operating System	Bare-metal, RTOS	Bare-metal
Programming	Any	C, C++, Assembly	HDL

Table 3.7: Processor Comparison Table

3.1.6 Choosing the Single-Board Computer

Raspberry Pi 5

The Raspberry Pi 5 is a single-board computer (SBC) that runs Raspberry Pi OS (a Linux distribution) in its default configuration. It is technically not a microcontroller due to the fact that it runs an operating system and features consumer-facing interfaces such as HDMI and USB. We are including it in our considerations due to the significant computational power necessary to run our auto-mix algorithm.

The Raspberry Pi 5 is the latest and most powerful model from the popular line of Raspberry Pi microcontrollers and SBCs. It features a Broadcom BCM2712 CPU with four Arm Cortex cores that run at up to 2.4 GHz. The Raspberry Pi 5 is available in up to 8 gigabytes of LPDDR4X RAM and features expandable storage through microSD and NVMe interfaces. The Raspberry Pi 5 also includes native support for Bluetooth and WiFi connections with its integrated hardware as well as its software libraries. Some of the relevant IO includes 28 GPIO (general-purpose IO) pins, 9 SPI interfaces, 6 UART interfaces, 7 I²C interfaces, and 2 I2S audio interfaces.

In terms of power, the Raspberry Pi 5 benchmarks well, scoring 72059 on the Embedded Microprocessor Benchmark Consortium's (EEMBC) CoreMark test [20]. This type of power is necessary to handle the auto-mix algorithm which may include auto-encoder and convolutional neural network models with upwards of 2 million parameters. Additionally, the extra power will be useful in developing a fluid GUI (graphical user interface). We desire to use a unified code base for the mobile and embedded GUIs, and thus the computational parity will ensure that performance on each platform will be similar.

Unfortunately, the Raspberry Pi 5 cannot handle all the computational needs of the digital audio mixer. For one, the Raspberry Pi 5 is already a complete system; there is little room to design a PCB to meet the requirements of senior design. Additionally, on a technical level, The Raspberry Pi 5 does not have the necessary amount of I2S connections to take in 8 channels of digital audio from our ADC. Thus, it is not a viable option for audio DSP. This leaves room for another microcontroller to handle that.

Overall, the Raspberry Pi 5 fulfills enough features for it to be considered for the mixer project. Its ample power as well as its wireless communication capabilities allow it to cover both the AI auto-mix functionality and the wireless communication functionality.

BeagleBone AI-64

The BeagleBone AI-64 is an extremely capable single-board computer from BeagleBoard. It features a Texas Instruments TDA4VM SOC that contains dual Arm Cortex A72 cores that run at up to 2.0 GHz as well as a plethora of accelerators. Among these are a C7x floating point vector DSP chip, a deep-learning matrix multiply accelerator, vision processing accelerators, depth and motion processing accelerators, and more. The main A72 cores benchmark at 24 million DMIPs

(Dhrystone Million Instructions Per-Second). The BeagleBone AI-64 is available with 4 gigabytes of LPDDR4 RAM and comes with 16 gigabytes of flash storage which can be expanded through microSD. The BeagleBone AI-64 does not feature any wireless connectivity, however, WiFi capability can be added through the PCIe expansion slot. Some of the relevant IO includes 10 GPIO (general-purpose IO) pins, 11 SPI interfaces, 12 UART interfaces, 10 I²C interfaces, and 12 Multichannel Audio Serial Port (MCASP) interfaces supporting up to 16 channels [21].

The Beaglebone AI-64 is a commendable choice due to its processing power and comprehensive physical IO. It would be an excellent choice for audio DSP. However, similarly to the Raspberry Pi 5, it is a complete system that allows for little room to be done in the realm of PCB and hardware design. Additionally, the lack of wireless connectivity makes it a poor choice for communications. Another consideration is the price. The Beaglebone AI-64 costs approximately 230 dollars [22]. The incorporation of one would put the project well over its intended budget.

Nvidia Jetson Nano

The Jetson Nano is a specialized SBC from Nvidia for use in AI inference. It features a quad-core CPU with Arm Cortex A57 cores, and, crucially, a dedicated 128 core GPU accelerator with 128 Maxwell cores. The Jetson Nano is available with 4 gigabytes of LPDDR4 RAM and comes with 16 gigabytes of flash memory. The Nvidia Jetson Nano does not come with any wireless functionality, however, it can potentially be expanded with wireless adapters through a PCIe connections. Some of the relevant IO includes 2 SPI interfaces, 3 UART interfaces, 4 I²C interfaces, and 2 I2S interfaces.

A particular draw of the Jetson Nano is its GPU accelerator. It is the only SBC in this lineup that comes with NVIDIA Cuda cores; as such, it should be by far the best in terms of AI performance. This makes it an attractive option to run the AI auto-mix models for the mixer.

The Jetson Nano is a solid option for the mixer project. However, it lacks in several key points. It does not include first-party wireless communication solutions. This makes it useless for the communications requirement of the project. Additionally, the physical audio IO is lacking. 2 I2S interfaces only allow for 4 channels of audio at best, which makes it an unviable solution for audio DSP. As such, there is not a use case for the Jetson Nano to cover within the project.

Final Comparison and Conclusion

Below is a tabular comparison of all the relevant specifications of the compared single-board computers.

SBC	Raspberry Pi 5	BeagleBone AI-64	Nvidia Jetson Nano
Clock Speed	2.4 GHz	2.0 GHz	1.43 GHz
# Cores	4	2	4
Other Processors	N/A	Accelerators	CUDA Cores
Core Type	Arm Cortex A76	Arm Cortex A57	Arm Cortex A57
RAM	Up to 8 GB	4 GB	4 GB
Flash Storage	External	16 GB	16 GB
# Audio Inputs	2	12	2
Price	\$80	\$228	\$129

Table 3.8: Single-Board Computer Comparison

The Raspberry Pi 5 was chosen for the mixer project. It was chosen for its processing capabilities as well as its wireless communication capabilities.

3.1.7 Choosing the Microcontroller

ESP32-S3

The ESP32-S3 is a low-cost yet richly-featured microcontroller from Espressif with several attractive features. The ESP32-S3 includes WiFi and Bluetooth capabilities as well as a competent set of peripheral interfaces. For processing, the ESP32-S3 features a dual core Xtensa LX7 microprocessor with clock speeds of up to 240 Mhz. The ESP32-S3 also comes with 512 KB of SRAM (static random access memory) and an onboard ROM of 384 KB. Both the SRAM and ROM are expandable with external memory devices through protocols such as multi-channel SPI and QPI.

Some of the relevant IO includes 45 GPIO (General-Purpose Input/Output) pins, 4 SPI interfaces (although 2 are reserved for external memory), 3 UART interfaces, 2 I²C interfaces, and 2 I2S audio interfaces.

As for processing power, The ESP32-S3 scores 1181 on the Embedded Microprocessor Benchmark Consortium's (EEMBC) CoreMark test [23]. This is meager compared to the aforementioned Raspberry Pi 5 and also does not fare particularly well against other microcontroller platforms (which will be expounded on in following sections).

One of the attractive development features of the ESP32 series of microcontrollers is its cross-compatibility with Arduino libraries and the Arduino IDE. Arduino is one of the largest embedded platforms, and cross-compatibility with Arduino grants the ESP32 access to many high-quality DSP and networking libraries that would facilitate rapid development.

Unfortunately, the ESP32-S3 has severe deficiencies that make it unsuitable for the project. Similarly to the Raspberry Pi 5, the ESP32 does not have enough

channels of I2S to handle all the incoming audio streams from the ADC. This makes it a bad choice for DSP, although given its low cost, it may be viable to use multiple ESP32-S3 units in parallel to perform DSP, with each module handling two channels. Unfortunately, this would complicate our system and increase the scope of our PCB design beyond what is reasonable.

As for graphical performance, the ESP32-S3 does not come with any graphical accelerators (like STM microcontrollers) or bespoke graphical libraries. Given our desire to develop a rather demanding GUI, the ESP32-S3 is not a great choice to power this aspect of the project.

Ultimately, the only real draw of the ESP32-S3 is its price as well as its wireless capabilities. It is not a particularly attractive choice for DSP or GUI. Given that the Raspberry Pi 5 is already covering wireless communications, there is no real place for the ESP32-S3 in our project.

STM32H747

The STM32H747 is a powerful, full-featured microcontroller from STMicroelectronic's flagship lineup of STM32H7 MCUs. The STM32H747 features a dual-core architecture with an ARM Cortex-M7 running at 480 MHz as well as an ARM Cortex-M4 running at 240 Mhz. The STM32H747 has 1 Mbyte of RAM and supports up to 2 Mbytes of flash memory through a Quad-SPI memory interface that runs at 133 MHz.

Of note is the form factor. The STM32H747 comes in a ball-grid array package. These types of packages are impossible to solder by hand and must be assembled with an automated pick-and-place machine. This complicates the development process; if a problem is found with the board, the whole MCU and board must be thrown out; it is not feasible to remove and reattach the MCU to a different PCB.

Concerning processing power, the STM32H747 packs a real punch, scoring 3224 on CoreMark, or an equivalent 1327 DMIPs [24]. Additionally, the MCU possesses 4 DMA (direct memory access) controllers to offload memory operations from the CPU. This is crucial to real-time audio processing. For graphics, the STM32H747 possesses an integrated LCD controller, a Chrom-ART graphical hardware accelerator to reduce CPU load, and an accelerated hardware JPEG codec. Overall, the MCU is a solid option for both DSP and GUI purposes.

Some of the relevant IO includes up to 168 GPIO (General Purpose Input/Output), 6 SPI interfaces with 3 multiplexed I2S channels, 4 UART interfaces, 4 I²C interfaces, and 4 SAI (Serial Audio Interface) interfaces. Of note is the fact that SAI interfaces are stereo, so 4 channels of SAI can theoretically handle 8 mono audio streams. This is in addition to the 3 I2S audio interfaces. As such, the STM32H747 features adequate IO for audio DSP purposes.

Unfortunately, there is still the concern of raw processing power and massive parallelism with eight-plus channels of real-time audio processing. The STM32H747 is only dual-core, and it would be a difficult task to optimize the DSP effects to operate across 8 channels in real time with minimal audio artifacts. Again, it would be possible to run several MCUs in parallel for DSP, but that is a complicated task

that is beyond the scope of this project.

In conclusion, the STM32H747 is a solid choice to power the GUI portion of this project. However, the Raspberry Pi 5 is already handling the GUI, so there is no need for the STM32H747 in our project. If this project were to be expanded on in the future, it may be a good idea to combine the GUI capabilities of the STM32H747 with the wireless capabilities of the ESP32-S3 to provide a cheap yet powerful alternative to the Raspberry Pi 5.

XMOS XU316-1024-TQ128

The XMOS XU316-1024-TQ128 is a full-featured, powerful MCU from XMOS's xcore.ai lineup of microcontrollers. The XMOS possesses a unique, multi-core architecture with 16 logical cores split into two "tiles" (groups) and a deterministic, hardware scheduler called xTIME that mimics the behavior of a real-time operating system (RTOS). The logical cores have high-speed interconnects both between the cores on the chip and with cores on other SoCs. The logical cores themselves are custom RISC cores. The XMOS possesses 1 MB of onboard SRAM. The XMOS has no onboard flash memory, but external memory can be connected through Quad-SPI and SPI protocols.

Again, the MCU package is of note. The XMOS XU316-1024-TQ128 comes in a TQFP-128 package with a 0.4mm pin pitch. The package is leaded unlike the ball-grid array of the STM32H7 microcontroller; however, the minuscule pin pitch and high pin count still make it difficult for hand-soldering. This high pin-density design may necessitate the development of a breakout board for the chip to be later connected to the main PCB. This alleviates troubleshooting woes due to removing and reattaching the MCU.

Another consideration is the power topology and overall integration complexity. The XMOS is unique in that it requires a 1v rail for core functionality, as well as 1v8 and 3v3 rails for peripherals. Additionally, the XMOS has power sequencing requirements. The peripheral rails must be brought up before the core 1v rail; then, the reset can be released the the XMOS can boot. This introduces extra complexity to our design. Fortunately, the XMOS features extensive documentation and even full schematics of their development boards intended to be used as reference designs. The XMOS may feature extra complexity, but it compensates with robust documentation.

Concerning processing power, the MCU runs at 600 MHz and provides 2400 MIPs of total processing throughput. This is substantially greater than all the microcontrollers covered thus far and provides an adequate amount of headroom for multi-channel audio DSP. Of course, there is also the aforementioned multi-core architecture of the XMOS MCU. This multi-core architecture means that each audio stream can have a dedicated processing pipeline. This is an incredible feature and will make multi-channel audio programming much easier; there will be no need to manually implement a scheduling algorithm to process each channel of audio in sequence. This makes the XMOS MCU an outstanding option for audio DSP. The extra cores can be used to handle IO, which may include communications with

the Raspberry Pi 5 as well as other hardware peripherals directly connected to the XMOS.

In addition to its raw processing power and multi-core architecture, the XMOS features a hardware vector unit for the acceleration of AI (Artificial Intelligence) tasks such as classification and inference. The XMOS was considered to run the AI auto-mix algorithm instead of the Raspberry Pi 5. Unfortunately, with some rudimentary size calculations, the model size would be in the realm of 10 Mbytes. The XMOS only comes with 1 MByte of RAM, so most of the model would have to be offloaded to external flash memory which operates at a slower pace. It remains to be seen whether or not the XMOS would process the AI auto-mix algorithm at a reasonable pace, but for now, the Raspberry Pi 5 has been chosen to run the algorithm. If the XMOS proves to be sufficient to run the auto-mix algorithm, a future MCU topology could include the XMOS MCU for DSP and AI purposes, an STM32H7 MCU for GUI and state management, and an ESP32 or dedicated wireless module for communications.

In regards to IO, the XMOS is also unique in that IO pins can be configured in software to run different serial communication protocols. The available IO is defined by the developer and not by the hardware, which makes it incredibly useful for our multi-channel audio DSP purposes. It is possible to implement eight-plus channels of digital audio input through I2S with the first-party serial communication libraries.

Overall, the XMOS XU316-1024-TQ128 is an incredibly compelling MCU for audio DSP purposes and may even be a suitable option for AI auto-mix processing. Its raw power as well as its multi-core architecture will make multi-channel processing much easier. As such, the team has decided to incorporate this MCU into the design alongside the Raspberry Pi 5.

Final Comparison and Conclusion

Below is a tabular comparison of all the relevant specs of the compared microcontrollers.

The final processor selection is a dual-processor topology with a Raspberry Pi 5 and an XMOS XU316 MCU. The Raspberry Pi will handle the GUI, wireless communications, and the AI auto-mix algorithm due to its immense power and ample memory. The XMOS XU316 will handle audio DSP due to its ample power and audio I/O as well as its easy-to-program, multicore architecture.

3.1.8 Power Supply

Prelude: Power Requirements and Topology

A major part of powering a mixer is providing sufficient voltage levels to the various components within the board. There are several key requirements that the power supply system must meet, dictated by the components selected for the mixer.

Processor	ESP32-S3	STM32H747	XMOS XU316
Clock Speed	240 MHz	480 MHz, 240 MHz	600 MHz
# Cores	2	2	16
Core Type	Xtensa LX7	ARM Cortex M7, M4	Custom RISC
RAM	512 KB	1 MB	1 MB
Flash Storage	384 KB, expandable	1 MB, expandable	External
# GPIOs	45	168	78
# Audio Inputs	2	7+	User-defined
Package	Module	BGA	TQFP
Price	\$1.85	\$16.53	\$15.67

Table 3.9: Microcontroller Comparison

For one, a high-current 5 volt rail is required to power the Raspberry Pi 5. This voltage rail will also power the audio circuitry on the board in tandem with a voltage inverter that outputs -5 volts.

One may ask why a dual-ended supply is needed. Most preamplifier ICs (and operational amplifiers in general) operate off a dual-supply system. It is possible to operate them off a single supply, but it falls on the designer to create a robust virtual ground solution which significantly increases the design complexity of the system. Virtual grounds are easy to make in theory with a resistor divider; however, any significant current draw biases the virtual ground away from the center of the single supply range such that the amplifier may clip. More robust virtual ground solutions utilize dedicated rail splitter ICs, however, they are limited in the amount of current they can provide. A dual-ended voltage rail allows for the audio signals on the board to oscillate around ground, negating the need to float the audio signals with a DC offset and simplifying the overall design.

As for other necessary voltage rails, 3.3 volt and 1.8 volt rails must be supplied to power the numerous peripherals that are necessary for the mixer's operation, things such as the AD/DA circuitry and IO control. One additional requirement imposed on the power requirements is the necessity of a 0.9 volt rail to power the XMOS xcore microcontroller. The last requirement is a high-voltage rail to provide phantom power to microphones. The high-voltage phantom rail can come in a variety of different voltages and will be expanded upon in later sections.

Such is the voltage topology of the project. Now that the voltage requirements have been established, it is necessary to discuss the implementation of the voltage supplies. Several components and technologies are involved. The following sections will expound on the power supply unit (PSU) and different types of voltage regulators, linear regulators and switching regulators.

Power Supply Unit

The first component to consider in powering a piece of electronic equipment is the DC power supply unit. Power supply units are confusing to define because they are distinct from voltage regulators, but they may or may not contain voltage regulators. In the context of consumer electronics, DC power supply units convert the high voltage output of wall outlets into a voltage that can be utilized by the piece of equipment, usually with a transformer. This voltage may be further converted by voltage regulators within the device or simply clamped steady.

For the conversion from wall power, the 120V AC signal must be converted to DC and stepped down. This requires the use of a transformer followed by rectification— where the transformer (using different coil and winding ratios) steps the voltage up or down depending on the purpose and subsequently adjusts current output. Rectification is the method that is used to turn AC power into pseudo-DC power, most easily done with a full-bridge rectifier of four diodes.

The voltage output of the DC power supply in the case of the mixer will likely be 15 or 12 volts. These are a reasonable voltagese to deliver the amount of power needed while keeping the overall current draw to a reasonable quantity. For the mixer, the power requirement is estimated to be under 75 watts. When delivering power with 15 volts, for example, the current draw is under 5 amps, which is reasonable for a consumer power supply. Once the high AC voltage wall has been stepped down to a reasonable operating voltage, the voltage must be further converted to to power the different components of the system.

Upon initial research, it was discovered that there are a plethora of triple output supplies that can supply a dual-ended voltage rail on the order of 10-15 volts and a high current 5 voltage rail. This provides a great solution to the power requirements of the mixer. A minimal amount of conversion is needed; only the 3.3 volt, 1.8 volt, and 0.9 volt rails need to be provided with further voltage regulation

Below is a comparison of some power supplies that meet the required specifications.

Part	Power (W)	Rails	Efficiency	Price
Mean Well RT50-B	50	5V @ 4A, 12V @ 2A, -12V @ 0.5A	77%	\$32.94
Mean Well RT50-C	50	5V @ 4A, 15V @ 1.5A, -15V @ 0.5A	78%	\$33.91
Mean Well RT65-B	65	5V @ 5A, 12V @ 2.8A, -12V @ 0.5A	77%	\$38.89

For the purposes of this project, the Mean Well RT65-B is the most suitable. It has the highest power output and the greatest amount of current available from its

voltage rails. Its 5 volt 5 amp rail is sufficient for the Raspberry Pi and the 2.8 amps available from the 12 volt rail should be sufficient to power our 8+ preamp circuits.

It is worthwhile to mention that an alternative to this is to procure an off-the-shelf consumer power supply for existing mixers. Some of them are able to provide both 5V and $\pm 15V$ at different amperages, such as the power supply of the aforementioned TouchMix mixer. This would negate the problem of not being able to efficiently supply the positive and negative rails of the operational amplifiers OPA2134's and cut down on the overall footprint of the PCB, making the design of the circuit significantly smaller. The problem with this approach is that it can be very costly, with most power supplies of this ilk costing upwards of \$100. With the cost comes less complexity on the overall board itself, but a detraction from the general usability of the mixer itself. The user must carry around a bulky, custom external supply versus plugging a power cord into the mixer and utilizing a relatively cheap, built-in power supply.

Linear Regulators

Once the high-voltage wall power has been stepped down to a usable DC voltage, the voltage must be further stepped down and regulated to provide power for certain components in the system. One of the technologies with which this is accomplished is linear regulators.

Linear regulators utilize a very simple concept to step voltages down—dissipate any unwanted output voltage as heat, which can be done through a variable resistor. As such, linear regulators are only capable of lowering the output voltage rather than also being able to raise the voltage. As expected with this type of hardware, it is very simple, requiring very little besides a few capacitors and resistors (plus transistors and Op-Amps) and thus very cost-efficient, being ideally less than \$1. Figure 3.9 exemplifies its simplicity, with very little in the way of active components.

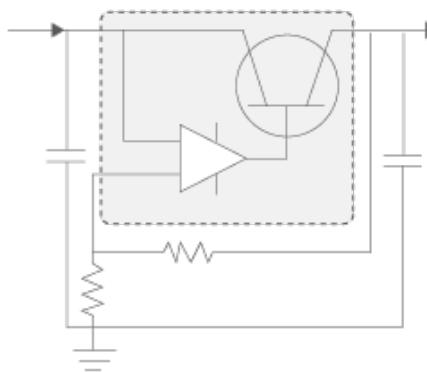


Figure 3.9: Basic Linear Regulator, Reproduced with Permission. © Renesas Electronics Corporation or one of its subsidiaries. All Rights Reserved.

In addition, linear regulators have very little noise present when being used, which is especially crucial when using analog signals.

In this case, since we would be using an IC as the regulator (e.g. TPS7A53B), it would reduce the overall complexity of the device with better performance compared to the basic transistor and Op-Amp setup. In figure 3.10, one can see that the only additional parts added to the IC are decoupling capacitors and a resistor ladder that can alter the output voltage [25]. This is effective because of the very low power required to use such an IC, but the efficiency decreases drastically the further away the desired output voltage is compared to the input voltage.

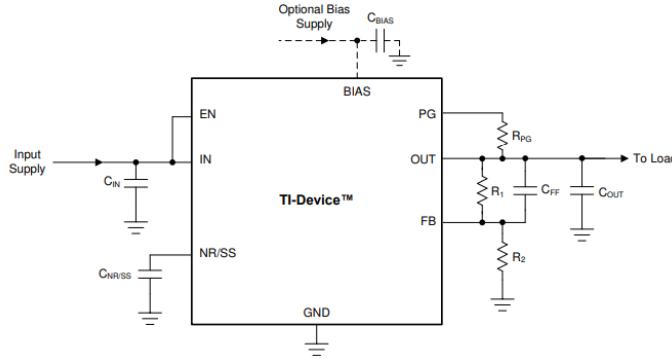


Figure 3.10: Example Application of TI's Linear Regulators, Courtesy of Texas Instruments Inc

Linear Regulator Comparison

In choosing a linear regulator, the most desirable characteristic is the minimal number of external components required. There are numerous linear regulators that have the same external current output and voltage output, but as shown in table 3.10, we chose the NCP1117DT50RKG due to the number of external components required being only two. This is additionally the case for the 3 other LDOs, as they do not require external feedback resistors to allow for a specified output voltage.

Part	Output Voltage	Output Current	Price	Ext. Parts
NCP1117DT50RKG	5V	1A	\$0.60	2
TLV73333PDBVR	3.3V	0.3A	\$0.13	2
TPS73618DBVTG4	1.8V	0.4A	\$2.34	2
TLV71209DBVT	0.9V	0.3A	\$1.06	2
SPX1117M3-L-5-0/TR	5V	800mA	\$0.36	4
LD1117S50TR	5V	800mA	\$0.44	3
TLV1117-50CDCYR	5V	800mA	\$0.46	4

Table 3.10: LDO Comparison Chart

Switching Regulators

Another type of technology used to regulate voltages is switching regulators. As stated before, switching regulators require more complicated circuitry, but are more efficient than linear regulators. As the name suggests, switching regulators achieve the desired output voltage by switching between off and on and using an output inductor to supply the energy through the circuit when off. This can be done either synchronously or asynchronously using FETs. As shown in figure 3.11, both the internal and external circuitry is more robust than that of a linear regulator.

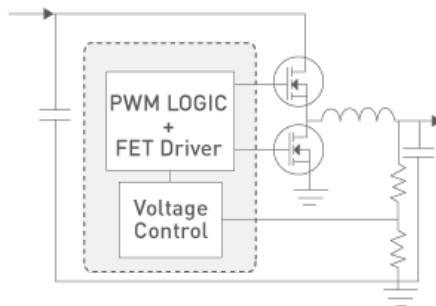


Figure 3.11: Example Switching Regulator Schematic, Reproduced with Permission. © Renesas Electronics Corporation or one of its subsidiaries. All Rights Reserved.

Unlike linear regulators, switching regulators are capable of outputting either a lower (Buck converter) or higher voltage (Boost converter), with some being able to do Buck and Boost conversion. While linear regulators become more inefficient the further away the desired output voltage is compared to the input voltage, switching regulators do not have this problem. This subsequently allows the desired output voltage to be much lower and/or higher than the input voltage. Additionally, switching converters are able to supply much more power through the devices, such as 2A through the TPS62822 [26]. The main downsides of using these types of regulators, though, are noise due to switching (much higher than any noise created by a linear regulator) and the cost being higher than a linear regulator due to complexity [27].

VRM Type	Linear Regulator	Switching Regulator
Output Voltage	Lower (Buck)	Lower (Buck) Higher (Boost) Lower/Higher (Buck/Boost)
Cost	Low	Medium
Noise	Low	Medium
Complexity	Low	Medium (Depends on IC)
Efficiency	Medium (↓ as $V_{in} - V_{out}$ ↑)	High

Table 3.11: Comparison Table of Linear vs. Switching Regulators

After deliberation about the stability of using high frequency switching converters, we felt it would be more proper to exclusively use LDOs exclusively for power regulation down-converting. This would allow for less noise in the power lines, and a simpler, more effective design of the overall power circuitry. While buck converters will not be used, we will still use a singular boost converter for amplification of the 12V input into 36V output for phantom power.

Switching (Boost) Converter Comparison

Because we are only using one boost converter to get from 12V to 36V, there were minimal restrictions guiding our decision. As such, choosing an established converter that can output a somewhat high current (100-200mA) was the main decision, and price coming in as the second most important point. Another point of concern was switching frequency, so we opted for a switching regulator that was not in the order of MHz.

Part	Current Limit	Switch Freq.	Voltage Range	Price
LT8338	1.2A	300KHz-3MHz	3-40V	\$6.48
TPS61175PWPR	3.8A	2.2MHz	2.9-38V	\$3.23

Table 3.12: Boost Converter Comparison Table

3.1.9 Inclusion of the Rotary Encoder

The inclusion of a rotary encoder on our mixer is essential for providing precise, intuitive, and versatile user control over various audio parameters. Rotary encoders are critical components in professional audio equipment, offering a tactile interface for real-time adjustments to settings such as volume, balance, equalization, and other effects. For our project, where both hardware and app-based control are

integrated, the rotary encoder enhances the physical interface of the mixer, complementing the touchscreen and app while ensuring that the system remains highly functional and user-friendly in diverse operating scenarios. The rotary encoder provides fine-grained control over audio parameters, which is particularly important in professional and live audio settings. Unlike a simple potentiometer, the incremental nature of a rotary encoder allows for consistent and precise adjustments without drift over time. This makes it ideal for tasks such as setting audio levels, tuning equalizer bands, or adjusting gain, where precision directly impacts the quality of the audio output.

For our project, we are focusing on utilizing the PEC11R-4030F-N0024, which is a high-quality, low-cost incremental encoder, meaning it outputs a series of pulses that indicate the direction and amount of rotation. It does not have an absolute position memory, so it measures changes in position rather than the position itself. For the rotary encoder, the shaft type is flat, shaft diameter: 6 mm, shaft length: 30 mm. The flat shaft design provides an easy and secure connection to knobs or other control elements. This outputs two square wave signals, A and B, in a quadrature pattern. This allows detection of both the direction of rotation and the number of steps moved. The encoder includes a push switch feature that can be actuated by pressing the shaft. It has a switch rating of 10 mA at 5 VDC. This functionality adds versatility by allowing the encoder to serve dual purposes, such as volume adjustment and mode selection. Operating voltage is typically 5 VDC, while the output type is digital (quadrature signals). The rotational life is rated for 100,000 cycles, ensuring durability over extended use. There is a detent mechanism which includes 30 detents, providing tactile feedback for each step of rotation.

The 24 PPR resolution, combined with 30 detents, ensures accurate control and tactile feedback, enhancing the user experience in applications requiring fine adjustments. The built-in push switch adds extra control capability without the need for additional components, saving space and simplifying design. With a lifespan of 100,000 rotational cycles, it is highly reliable for long-term use in both consumer and industrial applications. Its small size and through-hole mounting make it easy to integrate into compact PCBs and control panels. It operates reliably in extreme environmental conditions, making it versatile for various use cases.

While 24 PPR is adequate for many applications, it may not be sufficient for extremely high-precision requirements. As an incremental encoder, it cannot track absolute position without an external reference, which might require additional software or hardware for certain applications.

The PEC11R-4030F-N0024 is a robust, versatile, low-cost rotary encoder designed for a wide range of applications. Its combination of precision, tactile feedback, and dual-functionality push switch makes it an excellent choice for our audio project requiring reliable and precise user input.

3.1.10 Physical USB Pin Standards Comparison

Prelude: USB Pin Standard Necessity

In the development of a USB audio interface as part of our project, understanding the physical USB pin standards is crucial. USB (Universal Serial Bus) is a widely adopted interface for data transfer and power delivery between devices. The physical pinout of a USB connector varies based on the USB standard (e.g., USB 2.0, USB 3.x, USB-C) and the corresponding connector type (e.g., Type-A, Type-B, or Type-C). This section provides a detailed explanation of the physical pin standards for the commonly used USB interfaces, emphasizing their role in facilitating data transfer, power delivery, and compatibility, and why we decided to go with a USB-C 2.0 protocol.

USB 2.0 Pin Standards

USB 2.0, introduced in 2000, is one of the most widely used USB standards, offering a maximum data transfer rate of 480 Mbps. The physical pinout for USB 2.0 is consistent across Type-A, Type-B, and Micro-USB connectors, with slight variations in pin arrangement. VBUS (+5V): Provides a 5V power supply to the connected device. This pin is responsible for powering peripherals, such as the USB audio interface. D- (Data -): Transmits differential data signals in the negative polarity. D+ (Data +): Transmits differential data signals in the positive polarity. The combination of D- and D+ enables bidirectional communication using a differential signaling method, which minimizes noise and interference. GND (Ground): Provides the ground return path for the power supply and data transmission. The VBUS and GND pins supply power to the USB audio interface, while the D- and D+ pins handle the audio data transfer between the host device (e.g., Raspberry Pi or computer) and the USB audio interface.

USB 2.0 offers simplicity and wide compatibility, with a straightforward 4-pin configuration and a data transfer rate of up to 480 Mbps, which is sufficient for most standard audio streaming needs, including real-time data transfer. It is universally supported across devices and is cost-effective to implement, making it an appealing option for budget-conscious projects. However, its limited data transfer speed may struggle to handle high-resolution, uncompressed, or multi-channel audio data, and it lacks advanced features such as Power Delivery (PD) or multi-lane data transfer, which restrict its versatility for more modern applications.

USB 3.x Pin Standards

USB 3.x, including USB 3.0, 3.1, and 3.2, introduces a significant increase in data transfer speeds (up to 5 Gbps for USB 3.0 and up to 20 Gbps for USB 3.2 Gen 2x2) while maintaining backward compatibility with USB 2.0. USB 3.x connectors have additional pins to support SuperSpeed data transfer and dual-bus architecture. The pin configuration for USB 3.x Type-A and Type-B connectors: USB 3.x

connectors have 9 pins, with the 4 USB 2.0 pins (VBUS, D-, D+, GND) and 5 additional pins for SuperSpeed signaling: SSTX+ and SSTX- (SuperSpeed Transmit Pair): Differential signaling pair used to transmit SuperSpeed data from the host to the device. SSRX+ and SSRX- (SuperSpeed Receive Pair): Differential signaling pair used to receive SuperSpeed data from the device to the host. GND (Super-Speed Ground): Ground return for SuperSpeed signals. The additional pins in USB 3.x enable faster data transfer, which is essential for high-resolution audio streaming. For instance, transmitting uncompressed or multi-channel audio requires the high bandwidth provided by USB 3.x.x.

USB 3.x introduces significantly higher data transfer speeds, ranging from 5 Gbps with USB 3.0 to 20 Gbps with USB 3.2 Gen 2x2, allowing for seamless handling of uncompressed and multi-channel audio. Additionally, its dual-bus architecture supports simultaneous high-speed and low-speed data transmission, while maintaining backward compatibility with USB 2.0. However, the 9-pin configuration of USB 3.x increases complexity, power consumption, and implementation costs compared to USB 2.0.

USB Type-C Pin Standards

USB Type-C is the most recent evolution of USB connectors, designed to be reversible and universal. Type-C supports USB 2.0, USB 3.x, and USB4 standards, as well as other protocols such as Thunderbolt 3. It also supports Power Delivery (USB-PD) for higher power outputs, up to 100W (20V, 5A). With 24 pins, the pin configuration is similar to the 3.x. VBUS (4 Pins): Four pins are dedicated to providing power. Type-C supports variable power delivery up to 20V and 5A, making it suitable for devices with higher power requirements. GND (4 Pins): Four ground pins for stable power delivery and signal integrity. TX/RX Differential Pairs (4 Pairs): These pins support high-speed data transmission and reception for protocols like USB 3.x, USB4, and Thunderbolt. Each pair includes: TX1+/TX1- and RX1+/RX1- for SuperSpeed lanes 1. TX2+/TX2- and RX2+/RX2- for SuperSpeed lanes 2 (used in USB4 and Thunderbolt). D+ and D- (USB 2.0 Pins): Used for backward compatibility with USB 2.0 devices. CC1 and CC2 (Configuration Channel): These pins are used for identifying cable orientation, negotiating power delivery, and determining the role of the device (host or peripheral). SBU1 and SBU2 (Sideband Use): Reserved for alternate modes such as audio or video signals (e.g., DisplayPort). The versatility of USB Type-C makes it an excellent choice for modern USB audio interfaces.

USB Type-C stands out for its modern design and versatility. Its symmetrical, reversible connector improves usability and eliminates orientation issues, while supporting multiple standards, including USB 2.0, USB 3.x, USB4, and Thunderbolt. Furthermore, USB Type-C supports Power Delivery (PD), allowing for power and data transfer through a single cable, and it is increasingly adopted across industries, making it a future-proof choice. Despite these advantages, USB Type-C's 24-pin configuration adds complexity to hardware and software implementation, and its connectors and cables are more expensive than traditional USB options.

While USB Type-C supports older protocols like USB 2.0, physical backward compatibility requires adaptors for devices with legacy connectors.

Final Comparison and Conclusion

Ultimately, we chose USB Type-C with the USB 2.0 protocol for our project to balance functionality, cost-efficiency, and scalability. The USB 2.0 protocol provides a sufficient data transfer rate of 480 Mbps, meeting our real-time audio streaming requirements without the need for the higher speeds offered by USB 3.x or USB4. The simplicity of USB 2.0 reduces complexity and implementation costs, while USB Type-C's versatility ensures compatibility with modern devices and provides a pathway for future upgrades. The Power Delivery capabilities of USB Type-C enable the audio interface to be powered through the same cable used for data transfer, simplifying the design and eliminating the need for additional power connectors. Additionally, adopting USB Type-C ensures the long-term viability of the project, as it is increasingly becoming the standard for new devices. By combining the simplicity of USB 2.0 with the modern features and adaptability of USB Type-C, we achieve a design that is practical, efficient, and forward-compatible.

3.1.11 LCD Touchscreen Comparison

Prelude: LCD Touchscreen Necessity

In our project, the inclusion of an LCD touchscreen on the Raspberry Pi is essential for creating an intuitive and accessible user interface, enhancing the functionality and versatility of the system. While the mobile app provides a remote interface for controlling the mixer, the LCD touchscreen serves as a direct physical interface for users to interact with the mixer itself. This dual-interface design ensures that the system can cater to diverse user preferences and operating scenarios, offering flexibility and redundancy.

The LCD touchscreen allows users to interact with the mixer physically and directly, providing a tactile and visual experience that complements the app-based control. This feature is particularly beneficial in professional or live performance settings, where users may need quick access to controls without relying on an external device. By integrating the touchscreen into the mixer hardware, we ensure that essential functions, such as adjusting audio parameters, selecting inputs, or visualizing real-time data like FFT signals, can be performed seamlessly on the device itself. This direct interaction not only streamlines the user experience but also improves reliability, as users can operate the mixer independently of the mobile app.

Raspberry Pi Official LCD Touch Screen

The Raspberry Pi official LCD touchscreen is a 7-inch capacitive touch display specifically designed to integrate seamlessly with Raspberry Pi devices. It features a resolution of 800x480 pixels and connects via the DSI (Display Serial In-

terface) port. This direct connection reduces the need for additional adapters or converters, simplifying hardware integration. The display is supported natively by the Raspberry Pi OS, which includes built-in drivers that eliminate the need for manual installation or configuration. The capacitive touch technology supports 10-point multitouch, enabling precise and intuitive interactions.

The key benefits of the official touchscreen include its plug-and-play compatibility, which drastically reduces the development time and complexity of setting up the display. Its native support in Raspberry Pi OS ensures that the display functions reliably without requiring additional software or troubleshooting. The capacitive touch technology allows for smooth and responsive operation, making it ideal for tasks that require quick and accurate inputs, such as selecting audio settings or interacting with real-time visualizations. Additionally, the use of the DSI interface reduces cabling compared to HDMI or USB-based displays, improving power efficiency and reducing the chance of connection errors.

Despite these advantages, the Raspberry Pi official touchscreen has some limitations. The 7-inch size may not be adequate for applications requiring larger screens, such as detailed audio monitoring or interfaces with extensive menus. The resolution of 800x480 pixels, while sufficient for most applications, is relatively low compared to HDMI-based touchscreens that can offer 1080p or higher resolutions. These constraints make it less ideal for use cases where screen size and resolution are critical factors. However, for our project designed around quick user interaction, its simplicity and reliability outweigh these drawbacks.

HDMI-Based LCD Touch Screens

HDMI LCD touchscreens are highly versatile and come in a wide variety of sizes, ranging from 3.5 inches to 15 inches or more, with resolutions that often reach 1080p or even 4K. These displays connect via the HDMI port for video signals and typically require a separate USB connection for touch input. The availability of high-resolution options makes HDMI-based touchscreens ideal for applications requiring detailed visuals, such as our project, where precise graphical representations of data like FFT signals are important.

The primary benefit of HDMI touchscreens lies in their broad range of sizes and resolutions, allowing developers to select a display tailored to their specific requirements. Larger screens and higher resolutions enhance the user experience by providing ample space for detailed interfaces and clear visualizations. HDMI-based displays are also widely compatible with devices beyond the Raspberry Pi, making them a flexible choice for multi-platform projects. This versatility makes them suitable for professional setups that demand greater visual fidelity.

However, HDMI touchscreens have notable drawbacks. They require multiple connections, as the HDMI port handles video signals while USB or another interface is needed for touch functionality. This increases cabling complexity, especially in compact systems, and can make the overall setup less tidy. HDMI displays also tend to consume more power than DSI-connected screens, which can be a challenge in systems with power constraints. Furthermore, some HDMI-based

touchscreens require additional driver installation, which may complicate setup and introduce compatibility issues. While they excel in providing high-resolution visuals, their increased power consumption and connection complexity make them less suitable for our project, which doesn't require such a high resolution visualization.

GPIO-Based LCD Touch Screens

GPIO-based LCD touchscreens are compact, cost-effective displays that connect directly to the Raspberry Pi's GPIO pins, often using communication protocols such as SPI or I2C. These displays are typically small, with screen sizes under 5 inches, and are designed for simple and portable applications. Due to their direct GPIO connection, they do not require video ports like HDMI or DSI, making them a practical option for our project with limited hardware resources.

However, GPIO-based touchscreens have several limitations that restrict their use in our project. Their small size and low resolution — often under 480x320 pixels — make them unsuitable for interfaces requiring detailed visuals or larger screen real estate. Many GPIO screens also use resistive touch technology, which supports only single-point touch and lacks the responsiveness of capacitive touch screens. Additionally, the need to manually configure SPI or I2C settings and install drivers adds complexity to their setup. While GPIO-based screens are excellent for basic, low-cost projects, they lack the performance and user experience required for more sophisticated systems like ours.

USB-Based LCD Touch Screens

USB-based LCD touchscreens connect to the Raspberry Pi via a single USB port, handling both video signals and touch input through the same interface. These displays are often designed as portable or secondary screens, with sizes ranging from small 3.5-inch panels to larger desktop displays. The simplicity of their USB connection makes them a plug-and-play option for many systems.

The primary advantage of USB-based touchscreens is their ease of use. With a single connection for both video and touch input, they reduce the need for multiple cables and adapters. This makes them an appealing option for systems where minimizing connections is important. USB touchscreens are also widely compatible, as most operating systems recognize them as standard input and display devices, simplifying integration into multi-platform projects.

However, USB-based displays have notable drawbacks, particularly regarding data bandwidth and performance. Since USB connections are not optimized for high-speed video transfer, these displays often suffer from lower resolutions or reduced frame rates compared to HDMI or DSI alternatives. Additionally, many USB touchscreens require proprietary drivers for touch functionality, which can introduce compatibility issues or complicate the setup process. Their limited resolution and responsiveness make them less suitable for performance-critical applications like real-time FFT data visualization.

Third Party DSI Touch Screens

Third-party DSI touchscreens use the same DSI interface as the official Raspberry Pi display but are offered by third-party manufacturers. These displays provide greater flexibility, as they are available in a broader range of sizes and resolutions, including larger screens and higher-definition options. Some third-party DSI screens also offer additional features, such as integrated enclosures or extended touch capabilities.

The main benefit of third-party DSI touchscreens is their customization. We can select screens that better match their specific needs, whether that involves a larger display, higher resolution, or unique form factors. They maintain the advantages of DSI connectivity, such as reduced cabling, lower power consumption, and efficient communication between the display and the Raspberry Pi.

However, third-party DSI screens come with potential drawbacks. Many require custom drivers, which can complicate setup and create compatibility issues. The build quality of these screens may also vary significantly between manufacturers, leading to concerns about reliability and durability. While they offer more customization options than the official Raspberry Pi display, their variability in quality and the need for additional configuration make them less dependable for critical projects.

Final Comparison and Conclusion

Each LCD touchscreen type has unique strengths and weaknesses, making them suitable for different applications. The Raspberry Pi official LCD touchscreen stands out for its plug-and-play compatibility, responsive capacitive touch technology, and seamless integration with the Raspberry Pi, making it the most reliable and practical choice for our project. HDMI-based displays excel in resolution and size options but introduce complexity with cabling and power demands. GPIO-based screens are cost-effective and compact but lack the resolution and touch responsiveness required for advanced interfaces. USB-based touchscreens are simple to use but suffer from bandwidth and resolution limitations. Third-party DSI screens offer customization but present risks in terms of driver support and build quality. Ultimately, the Raspberry Pi official LCD touchscreen provides the best balance of compatibility, performance, and ease of use, aligning perfectly with the needs of our project.

3.2 Hardware Enclosure

It is of course necessary to package all the electronic components into a hardware enclosure both for protection and usability. In this section, several hardware enclosure options will be discussed.

The first option is a bespoke, manufactured metallic hardware enclosure with custom measurements to fit the PCB and all hardware components. As part of the early research for the project, a 3D rendering was created to envision the actual physical form factor of the device. The rendering is shown below.

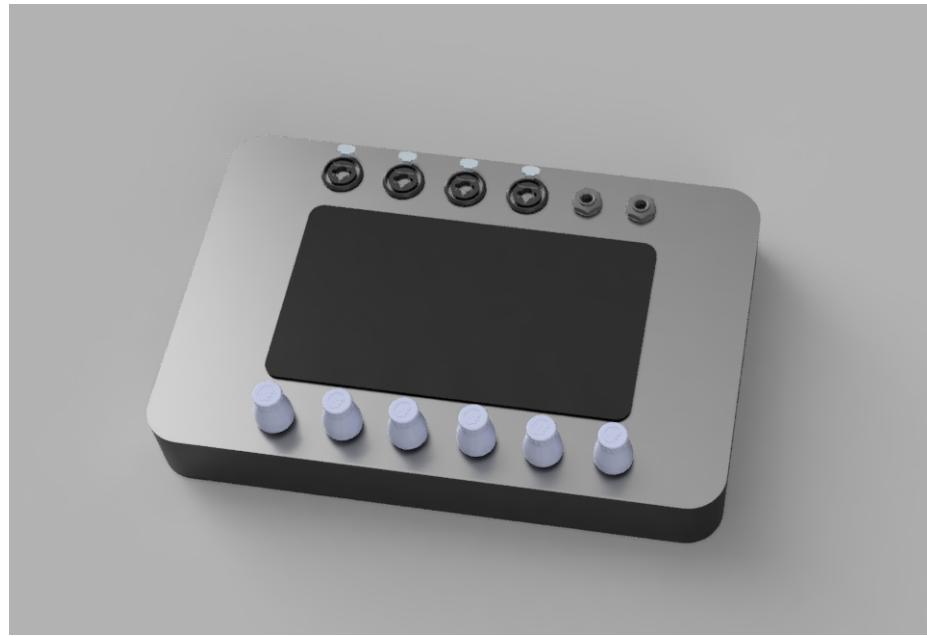


Figure 3.12: 3D Render of the Mixer

There is a slight discrepancy with the input selection due to it being an early design, however, the form factor would have been excellent. It is slim, portable, and durable with the metal exterior. However, it would have been a lot of extra work to design and manufacture. A bespoke case would have required machining of metal parts, which is not a skill that anyone on the team possesses. Furthermore, the team does not have the resources required to manufacture such an enclosure. As such, it is not really a viable option.

Another option is to use acrylic panels to create a "sandwich" style enclosure. This only requires the use of two acrylic panels and some posts to provide space between the panels for the PCB and other components. The cost is relatively low and the assembly is not difficult. However, this style of design is not particularly robust because of the lack of side panels. Liquid and particulate matter can easily enter the case which is not ideal. The design does not align with the design vision of the mixer; it should be portable and durable.

Another option is to use an off-the-shelf metal enclosure. There are several viable enclosures, however, there are surprisingly few metal project boxes that match the desired dimensions of the mixer. One viable option is the Hammond 1456PL3BKBU. The Hammond 1456PL3BKBU is a console style project box with dimensions of 14.2" x 11.8" x 3". This fits the dimension requirements of the engineering specifications and provides ample room for the PCB and other components. The only downside is the price. The price is over 50 dollars, which is quite expensive for a metal enclosure. However, it is the best option out of the selections available. A comparison of the enclosure options is shown below.

Enclosures	Custom Metal	Acrylic Sandwich	Hammond 1456PL3BKB
Dimensions	Custom	Custom	14.2" x 11.8" 3"
Durability	High	Low	High
Cost	High	Medium	\$54.02

Table 3.13: Caption for the table

3.3 Software Research and Comparison

The software systems of the digital audio mixer will span several key areas, namely, audio DSP effects, inter-system communications, app design, and machine learning in the form of the auto-mix algorithm. The research is presented here.

3.3.1 Audio DSP Effects

Prelude: Digital Audio Representation

Digital Signal Processing involves the processing of sampled, quantized signals. Analog signals are sampled at a constant rate and quantized to a certain granularity (e.g. 10 bits of quantization → 1024 possible values).

One of the basic principles underlying digital signal processing is the Nyquist frequency. Any analog signal can be adequately sampled as long as the sampling frequency is twice that of the highest frequency present in the analog signal. In the context of audio, a typical sampling frequency is 44.1 kHz. This is due to the fact that the range of human hearing lies approximately in the range of 20 Hz to 20 kHz. In order to adequately sample the audio, the sampling frequency must be above the Nyquist frequency of $2 * 20 \text{ kHz} = 40 \text{ kHz}$. For this project, the minimum common sampling frequency of 44.1 kHz is selected to reduce processing requirements.

Another convention in audio DSP relates to the range of possible sample values. In most applications, audio is presented as a sequence of values ranging from -1 to 1 representing the displacement of the audio driver. The granularity of the quantization within that range is dependent on the ADC.

In summary, the overall software representation of audio in the context of this mixer project will be a sequence of displacement values ranging from -1 to 1 that is sampled at 44.1 kHz (44,100 times per second).

Prelude: DSP Operators

Given a sequential array of float values sampled at a constant rate, there are several key DSP operators that can be used to modify the signal and achieve desired audio effects.

One operator is addition, or summing. Different streams of audio can be added together with simple addition. It is represented with the symbol shown below.

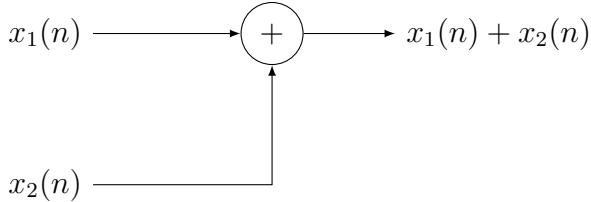


Figure 3.13: The Addition Operator

Another operator is the multiplication operator. The multiplication operator involves the multiplication of values in a stream by some constant. The operator is represented by the triangle symbol shown below. The value shown in the triangle is the multiplication constant.

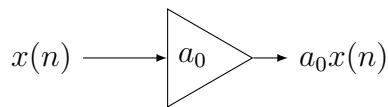


Figure 3.14: The Multiplication Operator

One last operator is the delay operator. The delay operator involves the delaying of a sample by a sampling period. The operator is represented by the symbol shown below. The exponent value of the z symbol represents the number of sampling periods of the delay. It is typically 1.

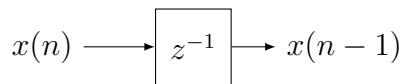


Figure 3.15: The Delay Operator

Together, these operators perform all the transformations on audio signals necessary to produce higher-level audio effects. The audio effects shown below use such operators to achieve the desired output.

Gain

Gain (specifically, digital gain) is the simplest DSP audio effect and is used as a standalone effect and as a building block of other effects. It involves a simple scalar multiplication of an audio stream. It is important to make a distinction between digital gain and analog gain. Analog gain is the electrical gain of the preamp system applied to the analog audio signal before it is converted to a digital signal by the ADC. It is controlled by hardware rather than by software. It is also important to note that analog gain may introduce irreversible clipping from the operational amplifiers which cannot be fixed by simply bringing down the digital gain.

Digital gain may be applied in many places. It may be applied to the signal path at the end of the signal chain or it may be applied by other effects modules (e.g. a compressor "makeup gain"). It is simple but useful.

Pan

Panning is a DSP audio effect that adjusts the balance between left and right audio channels so that the sound is placed in a certain place within a stereo field. The effect is simple; gain is applied to the left or right channel to impact the channel balance and perception of location. However, there are several details to be considered.

One must consider how the adjustment of balance operates so that the overall perception of volume is the same. If the left and right channels linearly changed in volume as the panning swept across the sound field, there would be an increase in volume at the center position compared to a pan of left or right due to the summation of the two channels. In order to alleviate this issue, mixers and audio workstations implement something called a "pan law." A pan law dictates the attenuation of an audio signal as it sweeps across the stereo field in order to maintain a constant sound pressure level. A common pan law dictates that both audio channels will be attenuated by 3 dB at the center panning position compared to a hard pan of left or right [28].

Equalizer/Filter

Equalizer effects (also referred to as filters) are essential audio DSP effects. Filters modify the frequency content of audio, allowing for certain frequencies to be suppressed or boosted. Filters can modify the frequency content of signals in several ways.

For one, they may completely cut off sounds that are beyond a certain frequency threshold. These are referred to as cut or pass filters. For example, a high-cut filter will cut off all frequencies higher than a certain threshold. A high-pass filter will cut off all frequencies lower than a certain threshold.

Filters may also boost or lower sounds beyond a certain threshold without wholly cutting them off. These are referred to as shelf filters. For instance, you may have a high-shelf filter that boosts or lowers the frequencies above a certain threshold. You may also have a low-shelf filter that boosts or lowers frequencies above a certain threshold.

Filters may also boost or lower sounds at a specific frequency. These are known as notch filters and are very useful for cutting out resonant, harsh frequencies present in certain sounds. Notch filters can be located anywhere in the audio frequency spectrum. One attribute that a notch filter possesses is a Q value. The Q value of a filter determines how wide or narrow it is.

The actual digital implementation of these filters is realized through something called a biquadratic filter. An example biquadratic filter is shown below.

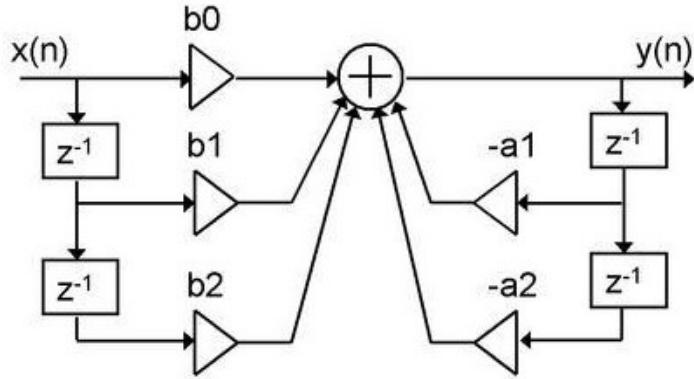


Figure 3.16: A General Biquadratic Filter, reprinted with permission from miniDSP

Of note is all the coefficients required to implement the filter. The coefficients determine the behavior of the filter, including what type of filter it is (cut/pass, shelf, notch) as well as the filter frequency. The general filter properties and values do not map onto the coefficients in any intuitive way. The actual selection of the coefficients is a complicated mathematical exercise that is beyond the scope of this segment. However, suffice it to say that there are a plethora of audio "EQ Cookbooks" which translate the filter values into coefficients that realize the filter. That is sufficient for the developers to work with. The difference equation of the biquadratic filter is shown below.

$$y[n] = b_0x[n] + b_1x[n - 1] + b_2x[n - 2] - a_1y[n - 1] - a_2y[n - 2] \quad (3.1)$$

Compressor

Compressor effects fall under the category of dynamics controllers. They control the volume envelope of the sound. They are essential to modern music in controlling the dynamic range of sounds and ensuring that everything remains at a level, pleasant volume.

Compressors typically have several key parameters which can be adjusted. One of these is the threshold. The threshold is the volume level above which the audio signal will be attenuated. Another parameter is the ratio. The ratio is how much the audio signal will be attenuated by in reference to the threshold. For instance, if an audio signal is 4 dB above the threshold and the ratio is set to 4:1, then the audio signal will be attenuated by 3 dB to 1 dB. Two more parameters that are closely related are attack time and release time. Attack time specifies how long it takes for the signal to be attenuated. Release time specifies how long it takes for the signal to stop being attenuated. The adjustment of both parameters affects the volume envelope of the sound and helps to control the transients of the audio.

The digital implementation of a compressor is relatively straightforward. An unaltered copy of the signal is used to drive an envelope detector. This envelope detector uses the volume of the signal to calculate an output gain based on the

signal itself as well as the compressor parameters such as threshold, knee, and ratio. This output gain value is then fed to a digital gain controller which effects the original signal such that it has been dynamically "compressed." After that, a simple gain effect may be applied to "make-up" the signal such that it retains the volume of the original signal.

Delay

Delay is a basic creative effect that takes a section of audio, delays it, and plays it again. It is akin to an echo effect, but it can be quite versatile as things such as feedback can be introduced. A typical delay effect has several parameters which can be adjusted. One is the delay time. This is the amount of time that the audio is delayed by before being replayed. Another is the feedback level. This feedback level is expressed in percent and controls how much of the delayed signal is again fed-back into the audio stream to echo and eventually die out. Finally, some delay effects may also have parameters which modulate the delay. For example, a common delay modulation is called ping-pong delay. The delayed sound will alternate between the left and right stereo channels, producing a bouncing sensation in the stereo field.

Delay is digitally implemented with something called a digital delay line. It is essentially a large buffer that stores audio samples. The samples enter the buffer and come out after a certain number of samples of delay. Afterwards, the samples are summed into the original audio processing buffer. The difference equation for the delay effect is as follows, where D is the number of samples to be delayed by and fb is the feedback coefficient.

$$y[n] = x[n - D] + fb \cdot y[n - D] \quad (3.2)$$

Reverb

Reverb is an effect that gives a sound a sense of space. Most people intuitively understand reverb. A sound played in a small bathroom will sound much different than a sound played in a cathedral. There are quite a few parameters attached to reverb, as well as a plethora of ways to implement the effect.

There are two broad categories of reverb effect implementations. One of them is algorithmic reverb. Algorithmic reverbs, as implied by the name, perform digital transformations of the sound using a plethora of other digital building blocks such as delay lines, filters, and more. Algorithmic reverbs allow for direct control of many of the reverb characteristics such as the decay time, early reflections, timbre, and more. Algorithmic reverbs are generally computationally inexpensive albeit at the cost of realism. However, inversely, algorithmic reverbs are useful for otherworldly reverb sounds that are impossible in the real world.

The other category of reverb implementation is the convolution reverb. Convolution reverbs use the mathematical operation of convolution to convolve the audio signal with the impulse response of a space. Through the use of real, captured im-

pulse responses, incredibly realistic simulations of real spaces can be added to an input signal. The downside is that convolution reverb is quite a computationally expensive effect. Furthermore, it can be difficult to customize the reverb effect due to its mathematical tie to the properties of the used impulse response. It is possible with extra methods to implement parameters on top of the reverb, however, the default implementation allows for one reverb sound per impulse response. For the purposes of the mixer project, it is practical to implement an algorithmic reverb due to their low computational requirements.

Flanger

The flanging effect is a creative effect that is heard in many records. It was discovered in the 1960's when studio engineers physically de-synced two tape machines by placing a finger on the flange (the ring that holds the tape) of one machine's reel. This created a swirly whoosh effect that became a mainstay in popular music and is still used today in many creative ways.

Flanging is considered a modulated delay effect. Flanging utilizes a delay line where the delay time is not static, rather, it is modulated by a low frequency sine wave. This modulated delay emulates the physical de-syncing of the tape machines and produces the whoosh-y flanging sound.

Distortion/Overdrive/Saturation

Distortion, overdrive, and saturation are all effects which shape the volume envelope of a sound in a non-linear fashion. Typically, they are associated with a heavy, blown-out sound that is common in rock and heavy-metal, among other genres. Distortion is incredibly common and is used both creatively and technically. In a creative capacity, it can create heavy, distorted sounds that artists across many genres love. In a more technical capacity, saturation can be used to introduce gentle harmonics which boost the upper-frequency presence of a sound and increase clarity.

For brevity, these collection of effects (distortion, saturation, overdrive) will be referred to as waveshapers because they shape the volume envelope of sounds. The fundamental operation behind all of these effects is a non-linear mathematical function that operates on the input sound as a function of the volume. Based on the function, the sound may be hard-clipped, soft-clipped, or something in between.

Hard-clipping is the process by which a sound is harshly limited at a certain threshold. For instance, one may set a hard-clip threshold at 0.6 (out of the +/- 1 digital audio representation). Any sample which crosses this threshold is simply replaced with this value of 0.6. This creates a hard-clipped signal with an incredibly flat volume envelope resembling a square wave. This process introduces a large amount of harmonics and overtones which gives the sound a distorted quality.

Soft-clipping is less extreme. Soft-clippers may use some non-linear mathematical function to slowly clip the volume once it passes a certain threshold. One common function is the hyperbolic tangent function. With the hyperbolic tangent

function, the volume of each sample is compressed into a smaller range, reducing the dynamic range while also increasing the harmonic content of the signal. The increase in harmonics is not nearly as drastic, and as such, the sound is gently shaped but not distorted like the output of the hard-clipper.

3.3.2 FFT Signals: Importance and Application

While searching for a means to visualize our audio signals, we found that the industry standard is to use Fast Fourier Transform (FFT) signals. FFT signals represent a transformation of an input signal from the time domain into the frequency domain through the application of the Fast Fourier Transform algorithm. The FFT algorithm decomposes a signal, which has been sampled over a specific period of time, into its constituent frequency components.

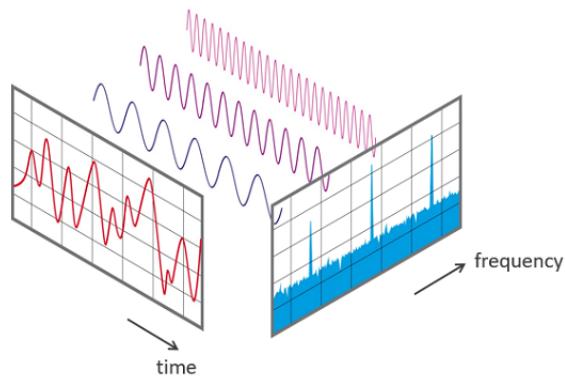


Figure 3.17: Fast Fourier Transform Algorithm

This transformation enables the visualization and analysis of the frequency content of a signal, allowing a user to see which frequencies are present as well as their associated intensities. By mapping frequency against volume in this way, we are able to observe and display the frequency spectrum of a piece of audio.

3.3.3 Comparison of Communication Protocols

Prelude: Communication Protocol Requirements

In the context of real-time data transmission between the xcore microcontroller and the Raspberry Pi, the choice of an appropriate communication protocol is paramount. A few requirements must be met. To start, real-time data synchronization is crucial. As audio gets sent from the xcore to the Raspberry Pi, it must be analyzed in real-time with little to no jitters or drop-outs. In addition, high-speed data transfer rates are necessary, as large amounts of data will be sent from the xcore to the Raspberry Pi. SPI (Serial Peripheral Interface), I²C (Inter-Integrated Circuit), UART (Universal Asynchronous Receiver-Transmitter), and Ethernet are the four widely used communication protocols that are examined in depth in this

section. The choice of SPI and I²C for the current project is then supported by an analysis of the particulars of each protocol, including latency, data transfer speed, system complexity, and communication range.

Serial Peripheral Interface (SPI)

The main factors in choosing SPI as our primary communication protocol came from the fact that SPI is a synchronous, full-duplex communication protocol that can be employed for high-speed, short-distance data transfer between integrated circuits in embedded systems. Seeing as audio information will be communicated between the xcore and Raspberry Pi in real-time, it was of the utmost importance that the protocol we utilized would be able to support this data transfer with minimal to no loss. The protocol operates on a master-slave architecture, where the master (Raspberry Pi) will control the slave device (xcore). Communication is facilitated through four dedicated lines: MOSI (Master Out, Slave In), MISO (Master In, Slave Out), SCLK (Serial Clock), and CS (Chip Select). The clock line is driven by the master to synchronize data transfer, while the Chip Select is used to select the active slave device at any given time.

The full-duplex nature of SPI allows data to be transmitted and received simultaneously, which allows for efficient real-time communication. This characteristic is critical for our application of real-time Fast Fourier Transform (FFT) data transfer, where continuous communication between devices is necessary to ensure timely processing and display of data.

One of the most significant advantages of using SPI is its high communication speed. This is necessary for our application, as we will require low-latency, high-throughput communication for the real-time processing and transmission of FFT data. However, a significant aspect of our project is organizing different data transfers to not have an overlapping issue when multiple forms of data is sent over a communication protocol. In our case, we want to make sure we have real-time, bidirectional data transfer as well as a simpler, lower-speed means of data transfer, which is where I²C is utilized.

Inter-Integrated Circuit (I²C)

The Inter-Integrated Circuit (I²C) protocol employs a half-duplex, synchronous communication method, where data is transmitted using two bidirectional lines: SDA (Serial Data Line) and SCL (Serial Clock Line). I²C supports multiple masters and multiple slaves on the same bus, enabling simplified wiring and reduced pin usage.

The decision to use I²C in our project stems from its ability to facilitate efficient one-way data transmission in systems with multiple peripherals. This makes I²C particularly advantageous for interfacing with devices like the LCD screen or other user input peripherals, as it allows multiple slave devices to communicate with the Raspberry Pi (acting as the master) over the same bus.

For our project, I²C will primarily be utilized to send data from the LCD screen, mobile app, or other user-operated devices back to the Raspberry Pi. For instance,

users may interact with the LCD screen to configure mixing settings. These interactions will generate data that needs to be relayed to the Raspberry Pi for processing and integration into the system's overall functionality. By employing I²C for these tasks, we ensure a streamlined, low-overhead communication method that complements the high-speed demands handled by SPI.

While I²C's simplicity in wiring is advantageous in systems with limited I/O pin availability, this protocol suffers from certain drawbacks that make it less suitable for high-speed, real-time applications such as ours. First, I²C is significantly slower than SPI, with typical data rates ranging from 100 kbps in standard mode to 400 kbps in fast mode, and up to 3.4 Mbps in high-speed mode. Even at its maximum speed, I²C cannot match the high throughput and low latency provided by SPI. This limitation is particularly hindering to our needs.

Moreover, I²C introduces additional overhead due to its addressing scheme. Each I²C transaction involves the transmission of an address byte, which identifies the target slave device. This additional data increases the communication overhead and reduces the effective throughput of the protocol. Additionally, I²C is a half-duplex protocol, meaning that data can only flow in one direction at a time. Due to these reasons, while I²C is advantageous for simple data transfer when users interact with the system, it fails with the larger data transfers of FFT signals from the xcore to the Raspberry Pi, which is why it is used for the simpler means of communication from user to Raspberry Pi.

Universal Asynchronous Receiver-Transmitter (UART)

The Universal Asynchronous Receiver-Transmitter (UART) protocol was another possible communication option. UART operates in an asynchronous, full-duplex manner, meaning that it does not require a shared clock line between devices. Instead, it relies on start and stop bits to synchronize data transmission.

UART is typically used for point-to-point communication, where two devices exchange data using two communication lines: TX (Transmit) and RX (Receive). UART suffers from significant limitations in terms of speed and efficiency. Standard UART implementations typically operate at speeds of up to 1 Mbps, though some high-speed variants can achieve higher data rates. Nevertheless, even at its highest speeds, UART is considerably slower than SPI, making it less suitable for FFT data transmission.

Furthermore, UART introduces additional overhead due to the inclusion of start and stop bits for each byte of data transmitted. This overhead reduces the effective throughput of the protocol and increases latency, which is undesirable for our system. Given these factors, UART was not a viable option.

Ethernet

Ethernet is a high-speed, packet-based communication protocol. Ethernet operates in a full-duplex, synchronous manner and supports data transmission rates of

10 Mbps, 100 Mbps, and 1 Gbps, which makes Ethernet an attractive option for our system which requires a large amount of data to be transferred.

However, Ethernet's complexity and reliance on packet-based transmission make it less suitable for real-time communication. Ethernet introduces variable latency's due to the nature of its packet-switched network architecture, where data is broken into packets and reassembled at the receiving end. While this architecture is highly effective for long-distance communication and networking, it can introduce unpredictable delays that are unacceptable in our use case of audio processing.

Furthermore, Ethernet requires a more complex hardware and software stack, which adds unnecessary complexity where simpler protocols like SPI can achieve the desired performance. Given the short-distance, high-speed communication requirements of our task at hand, Ethernet's advantages in long-distance communication are not applicable.

Justification for Selecting SPI

After a thorough evaluation of the communication protocols, SPI was selected as the most appropriate choice for this project due to its superior speed, low-latency characteristics, and ability to handle real-time data transmission. The high-speed communication provided by SPI ensures that the large volumes of FFT data and raw audio generated in this project can be transmitted efficiently between the xcore and the Raspberry Pi. Additionally, the full-duplex nature of SPI allows for simultaneous data transmission and reception, which is essential for our real-time audio processing needs.

While other protocols, such as I²C, UART, and Ethernet, offer certain advantages in specific use cases, they each present limitations that make them less suitable for this particular application. I²C's lower data rates and half-duplex operation introduce unacceptable delays, while UART's slower speed and additional overhead make it impractical for real-time communication. Ethernet, although highly capable in terms of speed, introduces unnecessary complexity and latency due to its packet-based nature.

3.3.4 Analysis of Protocols for Raspberry Pi to Mobile Data Transmission

In the design and implementation of our digital mixer, choosing the appropriate communication protocol to transmit data between the Raspberry Pi and mobile application is crucial. This decision hinges on several factors including the need for real-time data transmission, bandwidth constraints, network reliability, and ease of integration. This section provides a comprehensive comparative analysis of the most relevant communication protocols that can be utilized to transmit Fast Fourier Transform data between the Raspberry Pi and a mobile application, detailing the strengths, limitations, and ideal use cases of each protocol, and why we ultimately chose WebSockets.

WebSockets

WebSockets are a full-duplex communication protocol that provides a persistent connection between a client and a server, enabling continuous data transfer. In the context of our project, WebSockets are particularly well-suited for real-time communication between the Raspberry Pi and the mobile application due to their ability to transmit data bidirectionally without the need to repeatedly establish new connections [29]. WebSockets establish a connection between the server (Raspberry Pi) and the client (mobile app) through an initial HTTP handshake, after which the connection is upgraded to a WebSocket protocol. Once established, this connection remains open, allowing both the server and the client to continuously send and receive data without re-establishing the connection. For our project, this enables the Raspberry Pi to stream FFT data to the mobile app as soon as new data becomes available [30].

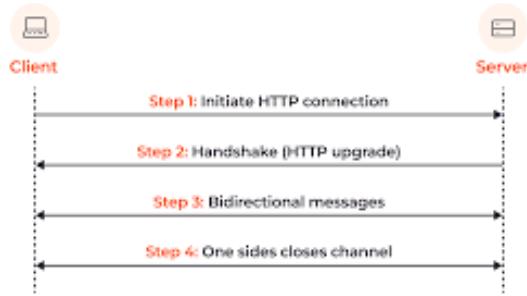


Figure 3.18: WebSocket Connection Steps

WebSockets excel in low-latency environments where real-time data transmission is required, which makes it ideal for our purposes of transmitting FFT signals in real-time to the app. In addition, since the connection remains open, WebSockets minimize the overhead typically associated with protocols that require the establishment of new connections for each data transfer, such as HTTP/REST APIs.

HTTP/REST APIs

HTTP/REST APIs are a commonly used communication protocol for client-server interactions. In particular, HTTP is a request-response protocol that is stateless, meaning that each request is independent of previous requests. This architecture is simple and widely supported, making it a popular choice for many applications [31]. In an HTTP/REST setup, the mobile application would send an HTTP GET request to the Raspberry Pi, which acts as a server, to retrieve the FFT data. The server processes the request and responds with the requested data in the form of a JSON. The mobile app would then display this data after processing the JSON parameters.

WHAT IS A REST API?

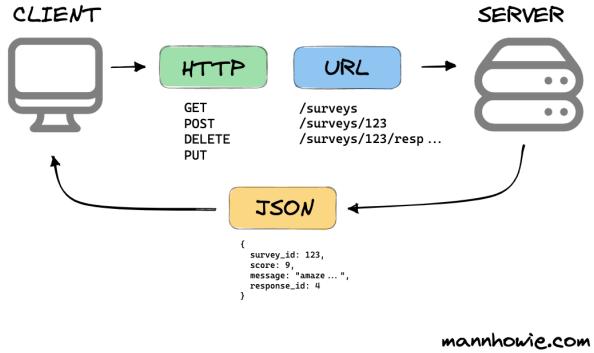


Figure 3.19: HTTP/REST API Connection Steps

HTTP is universally supported and easy to implement, which made it an attractive choice for our project initially. Since REST APIs are commonly used, they can be easily integrated with existing tools and frameworks, simplifying the development process further. However, the largest flaw we observed in HTTP/REST APIs was the fact that each data transfer requires a new connection to be established, leading to increased latency compared to protocols like WebSockets. In addition, HTTP's request-response model is less efficient for real-time applications where data needs to be streamed continuously, making it nonoptimal for our purposes, where we needed a constant, open communication stream.

Message Queuing Telemetry Transport (MQTT)

MQTT is a lightweight, publish-subscribe protocol designed for low-bandwidth, high-latency, and unreliable networks. It is widely used in Internet of Things applications, where devices need to send data over constrained environments [32]. In an MQTT setup, the Raspberry Pi publishes the FFT data to an MQTT broker under a specific topic. The mobile app, acting as a subscriber, listens for messages published to this topic. When the Pi sends new data, the broker pushes it to all connected subscribers. The broker manages all incoming messages by receiving them, applying filters, and determining the intended recipients based on their subscriptions. It then forwards the relevant messages to the subscribed clients. Additionally, the broker maintains session data for persistent clients, which includes tracking their subscriptions and storing any messages they missed during disconnection.

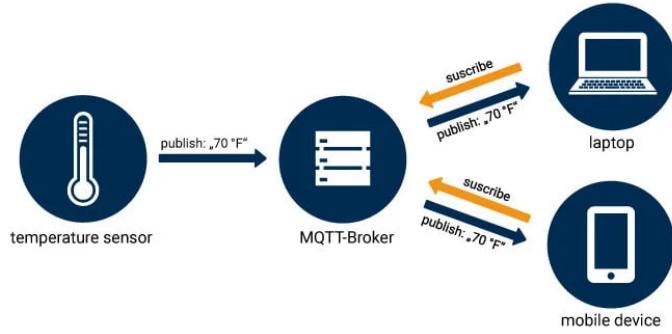


Figure 3.20: MQTT Connection Steps

MQTT is optimized for low-bandwidth environments, which makes it ideal for mobile applications such as ours that may experience intermittent connectivity. In addition, since MQTT implements a publisher-subscriber model, it enables scalable communication between multiple devices. MQTT supports different levels of delivery guarantees, ensuring that messages are delivered, which gives us a sense of security knowing that our data will be transmitted over. However, by implementing MQTT, we would need to implement the MQTT broker, which would add complexity to our project, which we are trying to minimize as much as possible. In addition, MQTT is optimized for small messages, which poses an issue for us, as FFT data is normally quite large.

Constrained Application Protocol (CoAP)

CoAP is another lightweight protocol designed for devices, such as the Pi, similar to HTTP but optimized for Internet of Things use cases. CoAP uses UDP instead of TCP, reducing overhead and increasing speed in certain environments. In addition, CoAP employs a request-response model, similar to HTTP, except it makes modifications specifically for a constrained network such as an Internet of Things device. CoAP employs the concept of observe, which allows devices to subscribe to changes that occur, instead of continuously polling for updates and wasting resources and wasting battery [33]. CoAP operates over UDP, providing faster communication. For example, the Raspberry Pi would act as a CoAP server, and the mobile app would send CoAP requests to fetch FFT data.

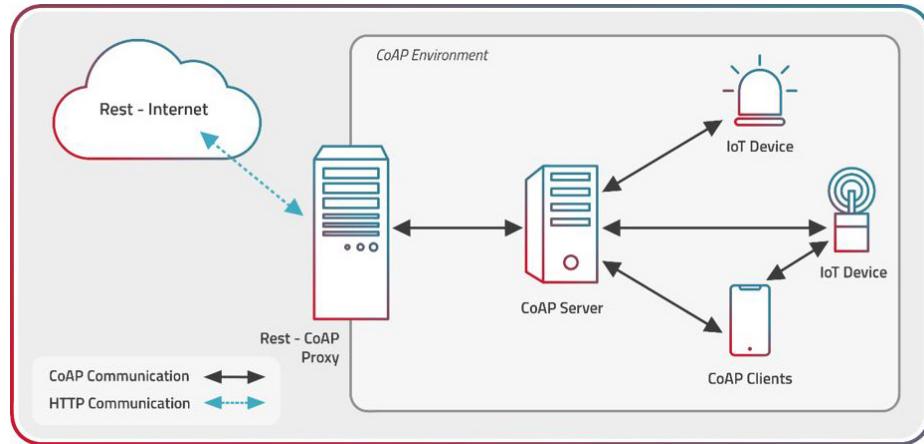


Figure 3.21: CoAP Connection Steps

CoAP's use of UDP instead of TCP makes it faster and more efficient in constrained environments. In addition, since it was designed specifically for low-power, low-bandwidth IoT devices, CoAP has minimal overhead compared to more heavyweight protocols like HTTP or WebSockets. However, since CoAP uses UDP, it does not guarantee delivery, making it less reliable than TCP-based protocols. CoAP is also less widely supported compared to WebSockets, which can make integration with existing platforms more challenging and eliminated CoAP as a viable option. Of particular importance to us is that our data is received by the client (mobile application / on-board LCD), or we may see inaccurate visualizations of the audio signals.

Justification for choosing WebSockets

After conducting comprehensive research to come up with a decision for this project, the best option will be WebSockets for real-time, low-latency communication, especially for streaming FFT data continuously from the Raspberry Pi to the mobile app. HTTP/REST could be used in simpler cases where periodic updates are sufficient, while MQTT offers a lightweight solution for environments with constrained resources or bandwidth limitations.

Protocol	Description
WebSockets	WebSockets provide a continuous, full-duplex communication channel between a client and a server, ideal for real-time applications. After an initial HTTP handshake, the connection is upgraded to WebSocket protocol, which remains open and supports bidirectional data flow. It ensures low-latency transmission of FFT signals, facilitating real-time data exchange between the Raspberry Pi and the mobile app.
HTTP / REST APIs	HTTP/REST APIs are popular for client-server communication due to their stateless request-response model and broad compatibility. In our setup, the mobile app would request FFT data from the Raspberry Pi via HTTP GET, receiving JSON responses. However, each request requires a new connection, adding latency and making HTTP inefficient for continuous data streaming, which is essential for our real-time needs.
Message Queuing Telemetry Transport	MQTT is a lightweight, publish-subscribe protocol ideal for IoT due to its efficiency in low-bandwidth and high-latency environments. In our project, the Raspberry Pi would publish FFT data to an MQTT broker, and the mobile app would receive updates as a subscriber. The broker manages message delivery and reliability options, making MQTT scalable and resilient for intermittent connections. However, the protocol's design for small messages challenges our larger FFT data needs, and setting up an MQTT broker adds complexity to the project.
Constrained Application Protocol	CoAP is a lightweight protocol designed for IoT applications, offering a request-response model similar to HTTP but optimized for constrained devices by using UDP instead of TCP. This reduces overhead, making it faster and more efficient for low-power, low-bandwidth environments like IoT. CoAP supports an "observe" feature, allowing clients to subscribe to updates without needing to poll continuously, conserving resources and battery life. However, since CoAP lacks TCP's delivery guarantees, it is less reliable, which can result in missed data and inaccurate visualizations. Additionally, CoAP's limited support on platforms compared to more widely adopted protocols like WebSockets reduces its viability for our project, where reliable, real-time data transfer is essential.

Table 3.14: Table of Communication Protocols

3.3.5 Wireless Data Level Protocols

In our project, Wi-Fi serves as the essential means of connectivity between the Raspberry Pi and the mobile application by operating at the physical and data link layers, providing the underlying infrastructure that allows data to be transmitted over a local network. Wi-Fi enables wireless communication between devices on the same network, creating a seamless, cable-free setup. When the Raspberry Pi and the mobile app are both connected to the same Wi-Fi network, they can interact as if they were directly linked, using Wi-Fi's capabilities to handle the physical transport of data packets between them. By leveraging Wi-Fi, the devices can communicate over more considerable distances within the network's reach, as long as both devices remain connected to the same network infrastructure. While WebSockets operates at the application layer, establishing the actual data-exchange protocol for real-time, bidirectional communication, it relies on Wi-Fi to provide the physical path for data transport. WebSockets alone cannot transmit data; it requires an underlying transport layer, such as Wi-Fi, to facilitate this. Wi-Fi establishes the initial network connection and transmits packets across it, supporting WebSocket connections that, in turn, create a persistent link for sending data from the Raspberry Pi to the mobile app and vice versa.

Wireless Fidelity (Wi-Fi)

Wi-Fi operates as a wireless local area network (WLAN) technology that enables devices to connect over a shared network, typically within a range of approximately 30 to 50 meters indoors. Wi-Fi's underlying structure involves a physical and data link layer within the OSI model, which manages wireless packet transmission across a robust infrastructure. A primary benefit of Wi-Fi is its ability to handle high data transfer rates, which, depending on the Wi-Fi standard (such as Wi-Fi 5 or Wi-Fi 6), can range from hundreds of megabits per second to over a gigabit per second. This high data throughput is particularly advantageous when transmitting large amounts of data, as is common in applications that require real-time data visualization, such as FFT signal analysis in the mobile app.

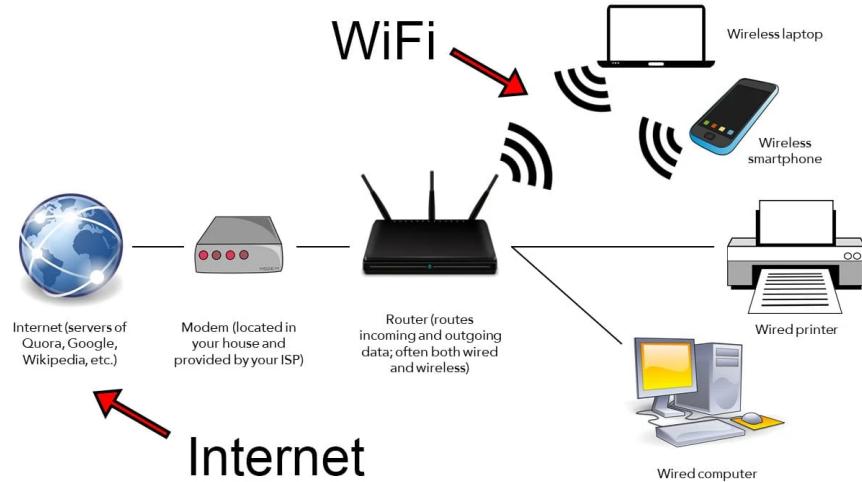


Figure 3.22: Wi-Fi Routing Explanation

Wi-Fi also benefits from its capability to support a longer transmission range compared to Bluetooth. Since Wi-Fi networks are designed to cover an entire building or large indoor space, it enables the Raspberry Pi and the mobile app to communicate across significant distances without connectivity issues, as long as both devices are within range of the same Wi-Fi network. Additionally, Wi-Fi supports network-based infrastructure, meaning multiple devices can connect to the same network simultaneously, allowing for a more scalable solution if additional devices need to be integrated. However, one notable drawback of Wi-Fi is its relatively high power consumption. The Raspberry Pi must maintain a continuous connection to the Wi-Fi network, which consumes more power compared to Bluetooth, especially in applications where long-term, battery-operated devices are used.

Bluetooth Low Energy (BLE)

Bluetooth Low Energy (BLE) is a short-range wireless communication protocol that is optimized for low-power consumption. BLE is commonly used for data transfer between critical devices that are within close proximity and can only afford to use very little energy [34]. In this setup, the Raspberry Pi acts as a BLE peripheral, transmitting data to a mobile app (which acts as a BLE central device) over a Bluetooth connection.

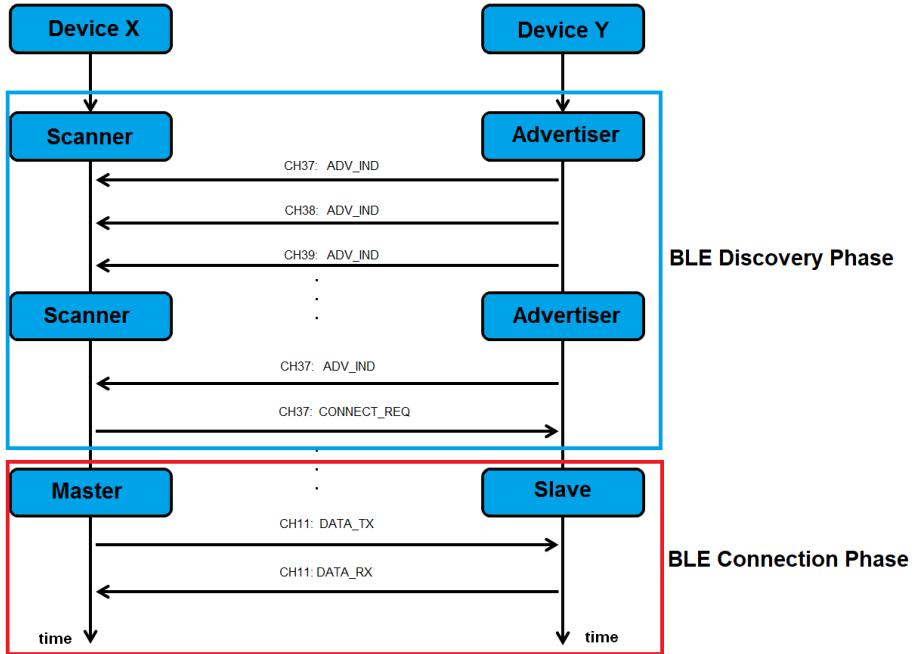


Figure 3.23: BLE Connection Steps

BLE is ideal for Internet of Things devices such as the Raspberry Pi, which has limited resources and battery life. Also, BLE excels at close-range communication, making it an efficient protocol for transmitting small amounts of data between nearby devices, such as the Raspberry Pi and mobile tablet. However, BLE is not designed for large, continuous data transfers, which would significantly hinder the ability to send FFT data from the Raspberry Pi to the mobile app.

Justification for choosing Wi-Fi

In our project, Wi-Fi is ultimately the preferred communication protocol over Bluetooth, primarily due to its superior data transfer speed, scalability, and suitability for real-time requirements. Our project requires transmitting FFT data continuously from the Raspberry Pi to a mobile application, where high-speed data exchange is crucial to ensure smooth, low-latency visualization. Wi-Fi excels here, offering significantly higher data transfer rates than Bluetooth. While Bluetooth Low Energy (BLE) generally supports speeds around 1 Mbps, and classic Bluetooth reaches approximately 3 Mbps, Wi-Fi can handle transfer rates of hundreds of megabits per second depending on the standard used, such as Wi-Fi 5 or Wi-Fi 6. This high bandwidth allows for efficient real-time streaming, reducing lag and ensuring that the FFT data displayed in the app is timely and accurate.

Additionally, Wi-Fi provides a network-based infrastructure that supports simultaneous communication across multiple devices, which is advantageous for scalability. This setup eliminates the need for direct pairing, as required by Bluetooth, making it easier to integrate additional devices into the system. Multiple devices can connect and communicate over the same Wi-Fi network, which allows for sim-

plicity, should multiple users want to connect to the mixer. Conversely, Bluetooth's direct device-to-device connection model is less suited for scenarios where multiple devices need to interact concurrently within a shared network environment.

Real-time data transfer is also crucial in our project, as we aim to stream FFT data continuously for immediate updates on the mobile app. Wi-Fi, with its persistent high-throughput connection, allows for efficient real-time data streaming, a task that would be challenging to accomplish with Bluetooth's limited bandwidth. Leveraging WebSockets over Wi-Fi offers a stable, low-latency connection ideal for applications where timely data updates are essential, ensuring that the mobile app's visualization of FFT data is in sync with the audio being processed by the Raspberry Pi.

While Bluetooth, especially Bluetooth Low Energy, has an advantage in terms of power efficiency, our project prioritizes data transmission speed and range over power consumption. The Raspberry Pi, generally connected to a power source, can accommodate the increased power demands associated with Wi-Fi, which is a reasonable trade-off given the benefits of high-speed and reliable data transmission. Thus, while Bluetooth could work well for simple, low-bandwidth applications, Wi-Fi more effectively meets our project's requirements for real-time, high-speed data transfer.

3.3.6 Real-Time Processing Considerations

Need For Low Latency

Our application is one that requires a real-time response time; latency must be minimized at all costs. Strictly defined, "real-time" refers to the mixer's ability to respond to input within a short, predictable time window. Latency is the time between a raw audio signal being inputted, performing the requisite processing on it, and outputting it.

The live performative nature of musical performances, where rhythm and time are an inherent part of the artform, makes it very sensitive to latency in the order of magnitude of milliseconds. The upper limit of acceptable latency for musical applications is approximately 10 milliseconds – latency that exceeds this are merely conspicuous at best, and at worst being highly disruptive to performers and causing the monitor signal to be unplayable against, rendering the mixer unusable.

Need for Stable Processing

The processing deadlines in real-time audio processing are rigid and cannot be missed under any circumstances. Unlike traditional software, where delays merely results degraded user experience or reduced performance, the failure to meet deadlines in real-time audio applications has a direct and immediate impact on the integrity of the audio output. These deadline overruns manifest as audible artifacts introduced into the outputted audio signal.

These artifacts range in severity depending on how many samples failed to be processed before the deadline. A few missed samples produces *popping* and *clicking*, onomatopoetic terms referring to artifacts in the output audio caused by sub-millisecond durations of erroneous silence. As more samples are missed, and longer portions of the output stream become silent, the artifacts become outright *dropouts*, causing sections of audio to be skipped or delayed. These are not minor glitches that can be overlooked; even a single errant sample produces a highly conspicuous pop, and the accumulation of enough of these artifacts can reduce the audio quality to the point that the mixer is rendered unusable.

Trade-off Between Latency and Stability

A fundamental trade-off exists between **stability** and **latency**, hinging on the size of the audio buffer. Buffering is used to mitigate the effect of variability in processing time and data availability. The size of the buffer determines how many samples, and therefore what length of audio, is processed at once in chunks, and it is this fact that causes the mixer's "buffer size" to be a crucial value in honing in the balance between low latency and high stability.

A larger buffer increases the stability of the audio stream by providing more leeway for the system to handle delays or spikes in processor load. Even if the average processing time per given length of audio is far less than that length, a long worst-case runtime can lead to deadline overruns were it not for the time headroom provided by a large buffer to absorb the delay and complete the buffer's worth of processing before the next deadline. With a smaller buffer size, every intermittent processing delay could result in audible artifacts.

However, the trade-off for the increased stability of larger buffers is an increase in latency. When a larger buffer is used, the system must accumulate larger amounts of data before it can be processed. Any amount of buffering must necessarily introduce a delay between the time the audio signal is captured and when it is heard, equal to the length of audio that the buffer's size represents.

Therefore, for the purposes of real-time processing, the optimal value for buffer size is the smallest possible setting such so that audio artifacts are non-existent, or at least very infrequent, in the output stream. This value results in the minimum possible latency without hindering stability, and is dependent upon on the optimization of the software and the power of the hardware it runs on.

Given the importance and salience of both latency and stability to users of audio software, most vendors of professional audio applications provide a setting to allow the manual adjustment of the application's buffer size, letting the user decide their own optimal balance between the two factors, depending on their hardware. Since the hardware of our mixer is a known factor, it is not a high priority for us to implement user control over buffer size; rather, we will likely hardcode it manually after empirical testing to determine its optimal value.

Principles for Designing for Real-Time

All audio processing applications intended for live use must handle the challenge of ensuring real-time response. We can draw from precedent established by the developers of existing audio processing software, and follow their prescriptions for designing our mixer's software architecture.

In a word, the key principle for ensuring real-time response guarantees is **determinism**. Every process, algorithm, and operation involved in the audio pipeline must be predictable. The time it takes for an audio signal to be processed must be predictable and bounded, regardless of system load or external factors. It must be known that every task is guaranteed complete within an allotted time frame.

It is not enough for a task's average runtime to be short. If any part of the system has an arbitrarily long worst-case runtime, then the triggering of that worst-case can cascade and cause delays that propagate throughout the system, eventually leading to missed deadlines and audio artifacts. As such, certain operations, such as locking mechanisms, dynamic memory allocations, and complex algorithmic processes, are strictly avoided within the real-time processing loop to minimize the risk of introducing variability.

3.3.7 Choices of Runtime Environment on the Raspberry Pi

Selecting a runtime environment for the Raspberry Pi requires weighing the trade-offs between computational overhead and ease of development. The decision affects how effectively the software interfaces with hardware, handles system tasks, and maintains real-time constraints. Below, we evaluate three major options: bare-metal programming, real-time operating systems (RTOS), and general-purpose operating systems (GPOS).

Bare-Metal Programming

In a bare-metal setup, software runs directly on the hardware without an operating system, offering unparalleled control over system resources. Frameworks like Circle or low-level programming in ARM assembly or C/C++ allow direct access to the Raspberry Pi's hardware.

This approach maximizes performance by eliminating OS overhead, enabling precise control over hardware for near-real-time behavior. It uses minimal system resources since no background processes are present. Bare-metal programming, however, demands deep hardware knowledge and significant effort to implement basic functionality such as multitasking, networking, or file systems. The lack of standard libraries and ecosystems further compounds development complexity.

While bare-metal programming offers determinism and resource efficiency, its steep development overhead and limited feature set make it unsuitable for projects like ours, where robust peripheral support and networking capabilities are essential.

Real-Time Operating Systems (RTOS)

RTOS options like FreeRTOS, RTLinux, and Xenomai strike a balance between bare-metal control and the flexibility of an operating system. Designed for deterministic timing guarantees, RTOS environments are well-suited for tasks with strict low-latency requirements, such as audio processing.

These systems provide predictable timing behavior, efficient resource usage, and reliability for mission-critical applications. However, they often lack comprehensive features like GUIs and networking stacks, limiting their utility for user-facing applications. RTOS environments also demand greater setup and programming effort, with fewer hardware drivers available compared to GPOS.

Although RTOS solutions offer the determinism required for real-time audio processing, their limited feature set and complexity render them less suitable for our project, where networking and a rich development ecosystem are critical.

General-Purpose Operating Systems (GPOS)

General-purpose operating systems, such as Raspberry Pi OS, Ubuntu Server, and Arch Linux ARM, provide a full-featured environment tailored for a wide range of use cases. These systems offer extensive hardware support, built-in networking stacks, and compatibility with a rich ecosystem of development tools.

GPOS environments excel in ease of development, with support for high-level programming languages like Python, Rust, and Flutter. Pre-built libraries and frameworks simplify tasks such as UI development, communication, and networking. Furthermore, the extensive community support for GPOS ensures robust troubleshooting resources.

The primary drawback of GPOS is its higher overhead, which can hinder applications requiring deterministic timing. However, in our project, where the Raspberry Pi manages state and user interface logic while delegating real-time audio processing to the xCore DSP, the overhead is inconsequential.

Final Choice: Raspberry Pi OS

Given the trade-offs, we selected Raspberry Pi OS as the runtime environment for our project. Its comprehensive ecosystem, hardware compatibility, and ease of use outweigh its resource overhead, which is negligible in our use case. By leveraging the rich development environment of Raspberry Pi OS, we can focus on building modular, maintainable software while ensuring seamless integration with the hardware and user interfaces.

3.3.8 Comparison of Protocols for Communicating Audio Parameter Changes

In designing a protocol for communicating audio parameter changes between software modules, we explored several standards with varying levels of complexity, flexibility, and overhead. Our goal was to select a protocol that balances efficiency,

interoperability, and ease of implementation while meeting the unique requirements of our mixer system. Below, we evaluate three candidate protocols: Open Sound Control (OSC), Virtual Studio Technology (VST), and a bespoke JSON-based protocol.

OSC (Open Sound Control)

Open Sound Control (OSC) is a protocol developed in the 1990s to facilitate communication between computers, sound synthesizers, and other multimedia devices. It is designed as an alternative to MIDI, offering higher resolution, flexible addressing, and extensibility. OSC messages are typically transmitted over UDP, which provides low-latency, connectionless communication suitable for real-time applications.

OSC defines messages in a hierarchical format, where each message consists of an address (e.g., `/mixer/volume`) and associated data (e.g., 0.75). The hierarchical addressing scheme allows for intuitive organization of parameters, making it well-suited for systems with complex configurations like audio mixers. Additionally, OSC supports various data types, such as integers, floats, strings, and arrays, providing the flexibility to encode a wide range of parameter changes.

However, OSC has notable drawbacks for our use case. First, its reliance on UDP means it lacks built-in mechanisms for reliability or acknowledgment, requiring additional effort to ensure delivery in systems where message loss is unacceptable. Second, the lack of standardized semantics for parameter changes means each implementation must define its own conventions, which can lead to inconsistencies across systems. Finally, while OSC excels in interoperability, this advantage is less relevant for our tightly integrated system, where all components are developed in-house.

VST (Virtual Studio Technology)

Virtual Studio Technology (VST), developed by Steinberg in 1996, is a comprehensive standard for audio plugins and their integration with digital audio workstations (DAWs). VST defines a robust protocol for transmitting parameter changes, supporting automation, and synchronizing audio processing with host timelines. These features make VST a cornerstone of modern audio production.

The VST parameter communication model is built around a tightly coupled host-plugin architecture. Plugins expose a list of parameters, each with a unique identifier, and the host sends parameter change messages to update their values. These messages support high precision and can be automated over time, making VST well-suited for complex audio workflows. Furthermore, VST includes support for MIDI, allowing seamless integration with external hardware.

Despite its strengths, VST introduces significant overhead that makes it impractical for our application. Implementing VST requires adherence to a complex API, including mechanisms for real-time audio processing, GUI integration, and preset management. This level of sophistication is unnecessary for our system, which

does not aim to be a plugin host or integrate with third-party tools. Additionally, VST's focus on host-plugin interactions limits its applicability as a general-purpose protocol for intermodule communication.

Bespoke JSON-Based Protocol

A bespoke JSON-based protocol provides a streamlined, purpose-built solution tailored to our specific needs. JSON (JavaScript Object Notation) is a lightweight, human-readable format for data exchange that is widely supported across programming languages. Its hierarchical structure aligns well with the requirements of a mixer system, where parameters are often organized into nested categories (e.g., channels, effects, routing).

Our JSON-based protocol encodes parameter changes as structured messages, specifying the target (e.g., a channel or effect) and the new value. For example, a message to update a channel's volume might take the form:

```
{  
  "command": "update",  
  "target": "channel_1.volume",  
  "value": 0.85  
}
```

This approach offers several advantages. It is language-agnostic, allowing us to use the same protocol across modules implemented in Rust, Flutter, or other technologies. JSON's flexibility enables us to encode complex data structures, such as effect chains or automation curves, without imposing unnecessary constraints. Moreover, the simplicity of JSON makes it easy to debug and extend as our system evolves.

Unlike OSC or VST, a bespoke protocol imposes no external dependencies, allowing us to optimize it for our specific requirements. While it lacks the interoperability of OSC and the feature set of VST, these trade-offs are acceptable given our focus on internal communication within a unified system.

Conclusion

After extensive research, we selected a bespoke JSON-based protocol as the most appropriate solution for our system. OSC's flexibility and interoperability are impressive but come with unnecessary overhead and a lack of reliability for our use case. VST's powerful feature set is tailored to plugin-host environments and is overly complex for our needs. By adopting a custom JSON-based protocol, we achieve a balance of simplicity, efficiency, and flexibility, ensuring effective communication between software modules while maintaining full control over implementation details.

3.3.9 User Interface Frameworks, Comparison and Selection

In order to develop both the mobile and onboard graphical user interfaces for our mixer, we must select a suitable framework. Each UI framework is designed to cater to specific use cases and offers strengths in different areas. To narrow down our search, our selection process was guided by the following priorities:

Cross-compilation The chosen framework must support compilation to both mobile (iOS and Android) and desktop (Linux distro on the Pi). This ensures that we can maintain a single, unified codebase for all UIs.

High performance with low overhead Given the resource-constrained nature of the Pi, which must concurrently run the backend manager process for our mixer, the chosen framework should consume minimal CPU and memory resources.

Low learning curve The framework should be easy to adopt and efficient for our team to work with, allowing more development time to go toward working on the backend and DSP software components.

Support for audio software interface components Ideally, the framework should provide prebuilt audio-specific components, such as level meters, equalizer controls, and frequency spectrums.

Given these requirements, we narrowed our search to a shortlist of three frameworks: **Flutter**, **Slint**, and **Qt**. As follows in table 3.15 is a brief description and comparison of each:

Justification for Selecting Flutter

After careful consideration, we chose Flutter as the framework for developing our mobile and onboard user interfaces.

The reasons for this decision are as follows:

Mobile-oriented Our mobile interface is intended to be the mixer's primary UI, with the onboard interface serving in an auxiliary capacity. In addition, our primary design objective for our UI is to provide a fluid, user-friendly interface. Flutter, as a mobile-first, modern, and responsive framework, is the best framework to meet these objectives.

Developer productivity Out of the shortlisted frameworks, Flutter has the lowest learning curve, as well as quality-of-life features such as hot reload and modular widget architecture. These factors will serve to significantly reduce development time, allowing for quicker iteration and refinement of our UI.

Expansive ecosystem Flutter has a large and active community, which offers a wealth of open-source third-party libraries and packages, making it easier to bootstrap the audio-specific UI components we require without the need to build entirely from scratch.

Feature	Flutter	Slint	Qt
Language	Dart	Rust, C++, JavaScript	C++, QML
Cross-Platform Support	Mobile, Web, Desktop (Linux, macOS, Windows)	Embedded, Desktop (Linux, macOS, Windows)	Desktop, Mobile, Embedded
Target Platform	Mobile, Web, Growing Desktop Support	Embedded Systems	Desktop, Mobile, Embedded Systems
Performance	Good (especially on mobile, not as great on desktop)	Excellent for embedded systems, very lightweight	Excellent, high performance, optimized for embedded
Ecosystem and Libraries	Large, many third-party libraries	Small, evolving	Very large and mature, rich set of modules
Learning Curve	Moderate (Dart)	Easy to moderate (depending on language)	Steep (especially with C++ and QML)
Use Case Fit	Best for mobile and web apps, some desktop apps	Best for resource-constrained embedded projects	Excellent for complex, cross-platform projects
License	BSD 3-Clause License	GPL-3.0 or Commercial	GPL/LGPL for open source, Commercial license for proprietary use

Table 3.15: Comparison of Flutter, Slint, and Qt

Performance Flutter's suboptimal performance is its largest pitfall; due to its bulky rendering engine, it is an order of magnitude less performant than Slint and Qt, which interface directly with the native APIs of platforms. Nonetheless, its performance is sufficient for our purposes, and the advantages Flutter offers outweighs

Ultimately, on a project with such a high degree of complexity, developer time is a highly scarce resource. Flutter strikes the best balance between ease of development, expansive built-in, and (sufficient) performance, out of the remaining frameworks, making it the most suitable choice for our project.

3.3.10 Automixing Methods, Comparison, and Selection

Automatic mixing has become an essential feature in modern studio systems, especially in live performances, broadcasts and general recording environments. It falls under this broader realm of intelligent music production (IMP), a field focused on applying automation, incorporating a level of artificial intelligence (AI) and machine learning (ML) into the field of music production to streamline the music creation process. Automixing involves dynamically balancing audio channels to achieve a cohesive and professional mix without the constant need for manual oversight. This makes it particularly advantageous for speeding up workflows and offering ease of access to amateurs or those without advanced audio engineering skills.

There have been a litany of approaches to automixing, offering distinct advantages depending on the context and requirements of the application. These methods can range from traditional, rule-based or knowledge engineered DSP methods to more recent advancements in machine and now deep learning. Knowledge-based systems are generally rules and guidelines created by audio engineering professionals or specifically selected by the designer and the respective algorithm to apply the domain knowledge for the task. Whereas, machine learning methods use statistical models and algorithms in a data-driven manner by applying a sort of optimization to accomplish the general function. [35]

A Historical Overview of Automatic Mixing

It is well documented in academia and industry practice of the history and direction of music production. Here we present a more historical overview exploring the progression and highlight the major developments in the field that show approaches changing from rule-based systems to the integration of machine learning due to the improving efficiency, creativity, and flexibility aided by general commercial excitement.

Efforts in automixing began with deterministic, rule-based systems like Dan Dugan's automatic microphone mixing in 1975 [36] performing basic tasks such as gain adjustment during live events by equalizing input channel levels. This design was simple yet effective and became the foundation for subsequent designs for automixing.

Fast forward to the early 2000s, where Pachet and Delerue's constraint optimization approach for multitrack audio mixing introduced the idea that there could be an optimal solution to balancing audio, changing the landscape of IMP. Here they proposed a system for real time mixing of audio sources with spatialization capabilities with a GUI to represent sound sources and their properties. When a user moves a source, the constraint system ensures the specified properties are the same. It allows users to choose different listening viewpoints while preserving the audio engineers intentions [37, 38].

The years between 2007 and 2010 saw a new wave of innovative research through the contribution of Enrique Perez Gonzalez's work, which broadened the

definition of automatic mixing by incorporating stereo panning, equalization, gain, etc into automated systems [36, 38]. Gonzalez's work on optimizing gain levels included calculating the ratio between each channel's loudness and the overall loudness of the stereo mix. He uses a cross-adaptive algorithm to map the ratios to channel gain values with the goal of optimizing gain levels of individual input channels in a live audio mixture [39]. To clarify, as shown in Figure 3.16, an adaptive audio effect is a sound transformation that adjusts its settings based on the characteristics of the audio signal it's processing through time-varying control. A cross-adaptive audio effects are adaptive audio effects where the parameters of one effect are modulated by features of another audio signal [40]. Another example was his work on equalization through the use of accumulative spectral decomposition techniques to analyze the frequency content of each channel. This approach applies a cross-adaptive audio effects to equalize the levels [41]. This work had strong contributions to the applications of live audio mixing and improving setup time. The work also marked the true start of automixing as its known today, with systems automating a more wider range of effects and parameters.

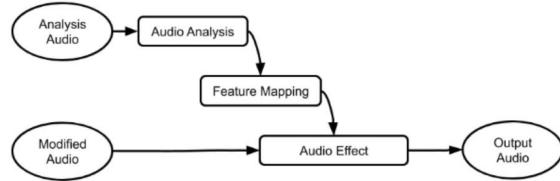


Figure 3.24: Flow diagram of an adaptive audio effect

By 2013, Brecht De Man and Joshua Reiss' introduction of a knowledge-engineered system highlighted the growing shift from simpler level and panning tasks as mentioned earlier, to more complex processing like dynamic range compression, equalization and high pass filtering, fading etc. Also, incorporating semantics like instrument and genre specific processing instead of just the low-level features of the raw audio signals [36]. By implementing these semantic mixing rules derived from audio engineering practices and textbooks, they ensured that the system made decisions that mirrored the expertise of professional audio engineers. This approach allowed automixing systems to handle more sophisticated tasks mentioned previously , which were previously difficult to automate without user input. By focusing on the semantics of the mixing process, De Man's system went beyond basic signal features and simple volume adjustments to deliver mixes that were more aligned with professional standards as well as leaving the door open to future research to include viable autonomous reverb/delay processors and reverberation rules [42].

Machine learning approaches rose to prominence between 2016 and 2017, bringing more advanced tools to the realm of automixing. The field transitioned from rule-based systems and knowledge-engineering techniques that dominated early automixing, towards more data-driven approaches. This shift was driven by the increasing complexity of tasks required in music production, such as dynamic range compression, reverb control, and audio source separation, which could no

longer be handled efficiently through deterministic, rules-based methods alone. Machine learning enabled more flexible and adaptive solutions. It also spurred the development of hybrid models and approaches that combined the strengths of both paradigms. For example, a system might use knowledge-based rules for initial tasks like instrument identification, but leverage machine learning for tasks like dynamic range compression and reverb which require finer tuning. This fusion allowed for more context-aware automation to handle technical mixing tasks but more subjective decisions based on training data [43]. The models are trained on large datasets of audio mixes to recognize the patterns and fine-grained nuances in the data that cannot be determined by predefined rules with music operating in such a complex multi-dimensional and rich parameter space. These predefined rules are now thought to be a departure from contemporary practice as its simply too restrictive. However, a common limitation of classical machine learning methods(methods as frequent late 2000s/early 2010s) that is proving to prevent further development is the lack of availability in data, with these models requiring large amounts of it to perform well. In recent years, deep learning has become widely accepted because of its ability to still generate meaningful results from smaller or out-of-domain datasets. These methods and architectures will be discussed in greater detail in the section below.

With mixing being a highly multidimensional in nature, the process requires audio engineers to make several simultaneous decisions. Is it the volume too quiet? The frequency range well-balanced? Is the panning fit withing the stereo field? Is the reverb set to appropriately to blend the instrument to the mix? Each of these elements can effect the overall mix and cannot be solved in isolation, attempting to solve one issue without considering how it interacts with other factors, can lead to new unresolved problems. Which is why mixing needs to be an all-in-one approach that is fully comprehensive and integrated [36, 44]. The transition to deep learning is what is needed to handle this interconnected and intricate task.

Classical Machine/Deep learning Methods and Models

Researchers in the 2010's and later wanted to explore machine learning by leveraging the large amounts of parametric data collected from professionals to train systems to do such tasks. Below is the in depth overview of some of these prominent classical machine learning approaches for general automixing and specific complex processing units. The drawbacks for each are also mentioned to provide the necessary evidence to move towards deep learning focused models.

Genetic Algorithms Here an automatic mixing was based on a genetic algorithm that extracted audio features for leveling tracks. It calculates the Euclidean distance between modified spectral histograms of a mix and a target sound. This defense serves as a measure of timbral similarity, which the system uses to adjust parameters during the mixing process. Timbre is the quality of a sound that allows a listener to differentiate between two sounds that have the same pitch, duration, and intensity. The genetic algorithm optimizes

the coefficients required for best possible recreation of the mix. While the system can recreate mixes, its limitation is in dealing with complex audio signals where timbral features may not capture nuances and often impossible in providing "objective validation". The genetic algorithm also struggled due to the lack of data available [45]

Linear Dynamical Systems The proposed system automates the mix-down process for multitrack audio files using parameters learned through supervised machine learning within a structured audio framework. The primary model employed here is based on linear dynamical systems, which provide a way to model the relationships between different audio tracks over time. The reliance on linear models poses limitations, particularly when handling non-linear dynamics of music production, such as compression and EQ adjustments. This simplification may prevent the model from fully capturing the intricate interplay between different musical elements that are often key to achieving a professional sound. It also was only proven to work on estimated parameters, mentioning the need for a larger collection of multitrack session files and only being limited to one particular genre, likely implying its inability to properly generalize mixes for other genres as well [46].

Hinge-Loss Markov Random Fields (HL-MRFs) The method described here was used for multitrack reverberation via hinge-loss Markov random fields to model the spatial and temporal dependencies between audio tracks. The model enables the intelligent application of reverberation by balancing the relationships between tracks and the desired spatial characteristics. HL-MRFs provide a probabilistic framework that allows for flexible reverberation effects. However, the challenge with this approach lies in the computational complexity of the model, which can limit its scalability to larger, real-time mixing environments. Additionally, the probabilistic nature of the model may sometimes lead to reverberation artifacts that detract from the overall mix quality, especially when handling highly dynamic sound sources. It was also shown not to be comparable to professional mixes in subjective evaluations with excessive reverberation leading to poor performance. It had a lack of parametric data here and also relied on estimations. [47]

General Statistical Feature-based Optimization Approaches Such systems have tried to capture the automation of a DRC, compression, and reverb. The dynamic range compression system employs a statistical approach by extracting real-time features from the audio signal—such as amplitude envelope and frequency content—to optimize the attack, release, and ratio parameters automatically [48]. Feature-based optimization is used for personalized compression, where the model statistically learns user preferences through feedback and adjusts future audio compression settings accordingly, tailoring the output to individual listeners [49]. Another system, similar to HL-MRFs, attempts to improve automation in reverberation with a feature-based approach. Its done by training a statistical model to learn mappings from

datasets of reverberant and non-reverberant signals, and analyzing the signal features(frequency spectrum and temporal dynamics) to apply appropriate reverberation effects [1]. The common drawbacks across these statistical feature-based optimization approaches for audio processing include a consistent mismatch between mathematically derived or automated parameters and perceptually accurate outcomes, leading to results that may sound unnatural or suboptimal. Additionally, the automation of key parameters limits user flexibility, reducing the ability to fine-tune settings for specific, complex audio scenarios. These methods also rely on predefined or static mappings between features and parameters, which fail to adapt dynamically to unexpected or rapidly changing audio inputs, further restricting their effectiveness in diverse real-world applications.

Overall, the classical methods over rely on predefined relationships between limited sets of audio features and parameters. It also only had limited and often manually curated datasets that included only specific audio features without comprehensive coverage of all relevant parameters (e.g., detailed processor settings, dry and processed audio comparisons) that often had to be estimated or excluded, impacting qualitative and at times quantitative performance. Such scarcity of these datasets prevented further pursuit and development of this field.

Deep Learning vs. Classical ML

Despite the availability of larger, more comprehensive datasets today, classical machine learning techniques for audio processing, and more specifically automixing aren't revisited or as accepted due to several reasons:

Superior Performance Deep learning architectures are made up of neural networks with many layers, optimized via loss functions enabled by gradient descent and backpropagation [50] to provide powerful generalization and representation learning capabilities. For example, models such as convolutional neural networks(CNNs), Long Short Term Memory (LSTMs), autoencoders, etc, have consistently outperformed classical machine learning techniques. These model architectures are capable of learning complex, non-linear relationships between high-dimensional audio features and output directly from the data, without requiring manually crafted features or predefined mappings.

End-to-End Learning Deep learning has the ability to perform what is known as end-to-end learning. Instead of separating feature extraction and model training as classical methods do (shown in Figure 3.17), deep learning models simultaneously learn both processes, optimizing them together. This integrated approach can capture richer, more nuanced relationships in the data and leads to better generalization, which classical ML models, even with new datasets, struggle to replicate effectively.

Transfer Learning and Pretrained Models In deep learning the concept of transfer learning is a very flexible and powerful tool. Transfer learning is when a

Method	Description
Genetic Algorithms	Uses genetic algorithms to optimize audio features, adjusting parameters to achieve timbral similarity between a mix and target sound, but struggles with complex signals and lacks objective validation due to limited data [45].
Linear Dynamical Systems	Automates multitrack mixing using learned parameters, but linear models are limited in capturing non-linear music production elements, resulting in poor generalization across genres [46].
Hinge-Loss Markov Random Fields (HL-MRFs)	Applies probabilistic modeling for reverberation, balancing spatial dependencies, but computationally intensive, and subjective evaluations show limitations in achieving professional mix quality [47].
Statistical Feature-based Optimization	Optimizes dynamic range compression and reverb using statistical features, but often lacks perceptual accuracy and flexibility, limiting adaptability to dynamic audio environments [1, 48, 49].

Table 3.16: Summary of Classical Machine Learning Methods for Automixing

model is trained on one task and is used as the starting point for a more domain-specific task. These pretrained models are originally trained on large datasets that can then be fine-tuned to specific tasks. This gives more adaptability to deep learning models for various inter and intra domain specific tasks related to audio processing. For example, a large network could be pretrained on large dataset of commercial music tracks to automate mixing parameters for new audio input across different genres or instruments. Once this model is pretrained, the model can be fine-tuned to a smaller, more limited dataset specific to a new task like mixing live concert audio or podcasts. The pretrained model uses its learned knowledge to quickly adjust and automate the mixing parameters, leading to faster, potentially more accurate

decisions for the new tasks even with less data and manual intervention.

Feature Engineering Limitations Classical ML methods heavily depend on manual feature engineering, which is labor-intensive and often limited in scope. While more parametric data is available now, classical methods would still require considerable effort to determine the right features and how to use them for specific audio tasks. On the other hand, deep learning automatically learns relevant features from raw data, significantly reducing the need for manual feature selection.

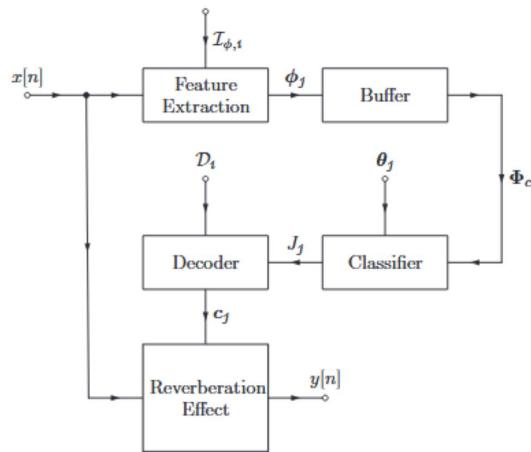


Figure 3.25: Architecture of [1] applying Reverb. Feature Extraction and Model training Training are Decoupled.

Computation Power and Infrastructure With modern computing infrastructure, we can efficiently handle large datasets and complex models, whereas revisiting classical ML techniques might not leverage computational resources as effectively. Classical models generally don't scale as well with larger datasets or benefit as much from parallelization on modern hardware like GPUs.

Research & Development Support/Focus Many areas of machine learning, has shifted focus toward neural networks and deep learning. As a result, there is more funding, tools, and community support for advancing deep learning approaches. This focus means less incentive to revisit classical methods, even with the availability of better datasets.

Challenges for Deep Learning Models in Automixing

The promise of deep learning is that it has the potential to learn mixes directly from tracks and mixes without any knowledge of the parameters used, alleviating a major constraint in data procurement and availability. However, in order to get there, it has to account for a number of challenges. **1)** There is *still* limited training data,

Characteristics	Deep Learning (DL)	Classical ML
Performance	Utilizes deep neural networks with multiple layers to capture complex, non-linear relationships and achieves superior generalization and representation [50].	Depends on handcrafted features and predefined mappings, which can limit model flexibility and performance.
End-to-End Learning	Integrates feature extraction and model training, optimizing both simultaneously for improved generalization.	Requires separate feature extraction and model training, limiting capture of nuanced data relationships.
Transfer Learning	Leverages pretrained models on large datasets, enabling quick adaptation and fine-tuning for domain-specific tasks with less data.	Lacks transfer learning capability, often requiring models to be built from scratch for new tasks.
Feature Eng.	Automatically learns features from raw data, reducing need for manual feature engineering.	Relies heavily on manual feature selection, which is labor-intensive and often requires extensive domain knowledge.
Computation Power and Infrastructure	Efficiently scales with modern hardware (e.g., GPUs), making it feasible to handle large datasets and complex models.	Less scalable on modern hardware, with limited benefits from parallel processing.
R&D Support	Receives significant funding, tools, and community support, with more resources focused on advancing DL.	Has limited support and fewer new developments, with less incentive for research focus.

Table 3.17: Comparison of Deep Learning and Classical Machine Learning for Automixing

needing paired inputs of original tracks and the good mixes. The demands on generalization make this task extremely data hungry, even as more datasets become available. **2)** Evaluating a mix is not necessarily as straightforward as it seems due to the subjective nature of music. There are not universal rules or guidelines to what makes a good mix. **3)** Mixing often has highly variable inputs with no consistent size and structure to inputs. Audio recordings can vary significantly, with

different mixes varying the number of input tracks (vocal, guitar, bass, drums, etc). This could prove to be an obstacle when it comes to generalizing across different mixes. The structure of the input could have different spectral, temporal and dynamic properties. Some tracks may be more high-frequency and some may be dominated by low frequencies. Mixes come in different genre's and styles with some being more vocal and others more instrumental. These subjective, artistic decisions can prove to be difficult to model. **4)** Ensuring a high-fidelity output is not easy. Such audio requires high sampling rates, often at 44.1 kHz, 48 kHz, or in professional contexts 96 kHz. The high sampling rate ensures the full spectrum of audio is captured in order to preserve all audio information. But doing this results in large volumes of data which can be very computationally demanding on both training and inferencing times. In real time automixing environments, the models need to process audio with minimal latency and the large inputs can generate artifacts when the system cannot keep up with the audio stream. Artifacts are unwanted alterations that distort the sound, like clicks, pops or unnatural frequency modulations. When deep learning applies various audio effects, it must account for this and reduces the margin of error in parameter estimation or direct transformations to prevent such distortions. The model has to have temporal and spectral consistency. There cannot be abrupt changes in amplitude over short time windows, this can create noticeable clicks. Thus, smooth transitions are necessary, which puts more stress on loss functions to minimize this. **5)** User interaction is something that should be incorporated into the deep learning process. For audio engineers, they often may need to tweak and adjust the mix output. While a model can fully automate the process, the creative and subjective nature of audio mixing requires human intervention to fine-tune the results. [51, 52].

Considerations for Automixing

The challenges outlined can then be distilled into several important considerations. The first being **interpretability**. It is important that the model can show what kind of manipulation is being applied to the input tracks in mix production as deep learning models tend to be *black-box* in nature. In the same vein, the **controllability** of the system can help understand not just what types of signal effects are being applied but allow user interaction so that adjustments to these transformations can be made. Extending from this is the idea of **context-aware systems**. Taking into the context of the genre, instrument, so that the final mix is more tailored and musically relevant.

The systems input recordings are unordered. Therefore, the order which the recordings are given to the system should not change the results. There are two main ways to go about doing this. One being a **fixed input taxonomy**, which trains the model using a dataset where the kind of and number of input sources are fixed in both training and inferences. E.g - In a vocal recording, the input could follow a set structure of lead vocal, backing vocal 1, backing vocal 2, and ambient mic. This allows for simplification in processing but can restrict flexibility and handicapping generalization capabilities where input tracks can vary unpredictability. The other

more permutation invariant approach is to through weight sharing. The same set of weights are applied to all input tracks, independent of its order, thus focusing on the track's intrinsic features rather the position of the track. Similarly is how the model handles the **number of inputs** given to it. A max number could be set, with silence being added when there is fewer than the maximum. If its greater than the maximum number of recordings, the shared weight implementation can handle this since it treats the tracks as set instead of a fixed vector. Another way is through possible aggregation or concatenation techniques and context sharing of the inputs [53]. Finally, the **expressivity** of the model should not be looked over. It defines how well the model can replicate to a wide range of mixing scenarios present in the dataset [51].

Deep Learning Model Applications and Overview

Recently, deep learning models have served a variety of nuanced techniques: data generation via source separation, creating differentiable audio effects, automatic upmixing, audio source manipulation and style transfer. But in general the goal is to reverse engineering an existing mix or create a new mix given a set of stems.

Early iterations applied deep neural and feed forward networks for specific processing of audio effects such as with dynamic range compression and for parameter prediction [54–56]. Implementations for specific effects also saw the use of convolutional neural networks, such as here with end-to-end equalization [57]. Models have shown how to deconstruct a multitrack mix and obtain the specific parameters from it given the original stems and the final mix [58]. A popular model known as Wave-U-Net was proposed for source separation directly in the time domain, eliminating the need for spectral transformations [59]. This model was then repurposed for specifically producing mixes for drum sets, with this variant being more data driven, swapping the input and output [53].

Furthering the development of automixing was the **Differentiable Mixing Console** developed by Steinmetz et al. They pushed boundaries by creating a differentiable pipeline for emulating the series connection of audio effects of EQ, compression, and reverb. The model enabled the full automation of multitrack mixing by predicting parameters directly from raw audio waveforms, introducing neural proxies for traditional effects processors, and incorporating a stereo-invariant loss function [60]. Despite these advances, the challenge of generalization remained but was significant step toward learning mixing conventions and paving the way for more sophisticated neural audio effects systems.

To address the limitations of datasets and generalization capabilities in models, Martínez-Ramírez et al. proposed an approach to reuse out-of-domain datasets like wet stems for training deep learning models to mix audio. The method involved normalizing wet stems based on audio features related to common effects, allowing the model to approximate professional-level mixes. They also designed a novel stereo-invariant, perceptually motivated loss functions to improve mix quality. The use of large datasets in conjunction with the new loss function showed promising results and reduced the gap between auto and human-made mixes [61].

Working around genre limitations brings about the idea of style transfer. Style transfer involves applying the characteristics of one mix or genre to another. In the context of music production, this could mean adapting the mixing decisions or audio effects typical of one genre to another. By leveraging deep learning models trained on large datasets, a system could automatically identify and adapt the appropriate mix based on the target style, making it more intention-driven to what a potential client wants [51].

Addressing the need for style transfer, the latest work of Vanka et al. (2024) introduced **Diff-MST**, a framework that automates the generation of a multitrack mix by inferring production attributes from a reference song. Their system builds of the differentiable mixing console with a transformer-based controller to predict control parameters for audio effects to generate mixes with high fidelity and can accommodate any number of input tracks. Diff-MST offers user control and allowing post-hoc adjustments, offering the practicality that previous models may not have had for real-world applications and tasks. Additionally, the audio production style loss function captures the global sound, dynamics and spatial attributes of the reference song, improve the transfers accuracy [62].

Model Architecture Comparison and Selection

Before delving into the specific architectures under consideration, its important to look at the frameworks or paradigm approach these models will fall under, why its relevant to our task, and a justification of one over the other. The two dominant schema for automixing model architecture is either direct transformation or parameter estimation.

Direct Transformation

Direct Transformation approaches aim to simplify the process by directly mapping input recordings to a final audio mix, bypassing the need for manual parameter adjustments required by traditional audio engineering methods. The model tries to generate a mix that is as close as possible to the reference mix, without requiring information about specific audio effects or the internal workings of digital audio workstations (DAWs). The training is guided by a loss function that minimizes the difference between the predicted and ground truth mixes in either the time, frequency, or combination of both.

Steinmetz et al. characterizes Direct Transformations as follows: A dataset consists of multiple examples, with each example containing a set of input recordings $\mathbf{x}_1^{(i)}, \mathbf{x}_2^{(i)}, \dots, \mathbf{x}_N^{(i)}$, along with a corresponding mixed output $\mathbf{Y}^{(i)}$ produced by an audio engineer. Formally, the dataset is represented \mathcal{D} as:

$$\mathcal{D} = \left\{ \left(\mathbf{x}_1^{(i)}, \mathbf{x}_2^{(i)}, \dots, \mathbf{x}_N^{(i)}, \mathbf{Y}^{(i)} \right) \right\}_{i=1}^E$$

where E is the number of examples. The model f_θ is trained to take the recordings as input and produce an estimate $\hat{\mathbf{Y}}^{(i)}$, aiming to minimize the difference between $\hat{\mathbf{Y}}^{(i)}$ and the ground truth mix $\mathbf{Y}^{(i)}$. This is done by optimizing a loss function

$$\mathcal{L}_a(\hat{\mathbf{Y}}^{(i)}, \mathbf{Y}^{(i)})$$
 [51].

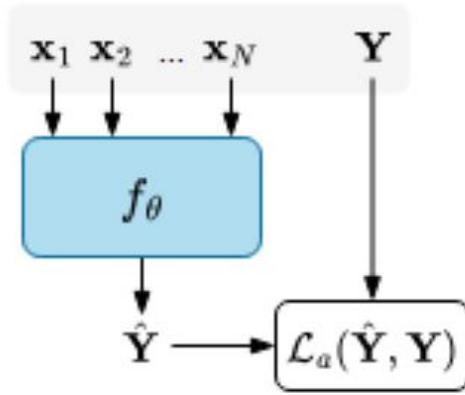


Figure 3.26: Visualization of Direct Transformation courtesy of Steinmetz et al.

Advantage Due to its minimal assumptions. It allows the model to implicitly learn and apply audio effects without needing explicit knowledge of the signal processing chain.

Disadvantage(s) A large dataset is required for a reasonable training performance and generalization since the model makes few assumptions. It tends to produce degrading signal transformations, aka artifacts. **Interpretability and controllability** is limited since the user cannot see the mixing parameters but only the final audio transformation.

Parameter Estimation and Justification

The limitations of direct transformations motivate this schema to focus on automatically adjusting audio processing parameters (gain, pan, EQ, etc) to optimize the overall mix. In deep learning, the process begins with the network observing the multi-track input audio, and learning how to predict the best set of parameters for different effects. The predictions are most often framed as regression problems where the network is trained to put specific values that would result in the desired sound [63].

Steinmetz. et al models it as follows: A mixing console modeled by a function h , which processes a set of N input recordings along with parameters for each recording, resulting in a final stereo mix [51]:

$$\mathbf{Y} = h(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N, p_1, p_2, \dots, p_N)$$

Advantages Modularity because the signal processing change assumes each input is processed through the same structures - a series connection of audio effects. **More Interpretability/Control** as this manipulation is handled by

predefined signal processing devices outlined in h . This in turn reduces the work that has to be done by the model.

Disadvantages Reduced expressivity limited by h . A loss computed in the parameter space measuring the distance between estimated parameters and ground truth parameters does not have many datasets that support this. It would require audio engineers to follow a fixed mixing chain to generate training data which would hurt practical use cases and data representativeness. Additionally, there is the idea of overparameterization, where different combinations of parameters could result in very stark differences in the audio domain. Sets of parameters could generate similar mixes but be far apart in parameter loss and be penalized [51].

However, there is a work around to the computing parameter loss where the audio loss is simply used. This would require a differentiable signal processing chain of audio effects so that gradients could be automatically calculated and so the parameters could be optimized in the audio domain, thus circumventing the limitations of the parameter space, avoiding challenges of overparameterization and data scarcity. With preexisting methods for automatic differentiation in PyTorch, TensorFlow available, a parameter estimation approach can be leveraged within the deep learning framework. And given the nature of our project as a whole, where real-time performance and user control are cornerstones, it is the best approach. The only concern with automatic differentiation being computation as signal processing has recursive operations that, in the time domain, can be problematic for gradient computation because the gradient will need to be calculated sequentially at each time step during back-propagation [51]. This will be taken into account in the design phase.

With the problem formulated, the architectures can now be outlined, compared and one method ultimately selected.

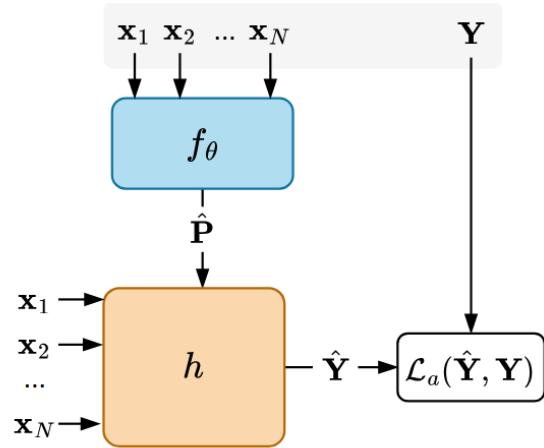


Figure 3.27: Visualization of Parameter Estimation with Audio Loss courtesy of Steinmetz et al.

Architecture Breakdown and Justification By breaking down the various models in Table 3.18, we can see that while Mix-Wave-U-Net offer promising capabilities due to its simplicity and generalization in the context of wet stems, it leaves much to be desired due to the output being a final stereo mix as a result of direct transformations. The Diff-MST overall structure and need for a reference track does not align with our necessary application where in our setting we would only be giving the raw audio supplied by the performers using the mixer. Although, there are aspects of its architecture that are desirable such as how it implemented the differentiable effects fairly straightforward without proxies via the **dasp-pytorch** library. This leaves the Differentiable Mixing Console as the ideal candidate for the automixing algorithm to be implemented because of its flexibility and control over parameters. Also the input being in mel-Spectograms allows for faster processing compared to raw audio recordings. The complex gradient computation will be something to be weary of as its put into design, noting potential drawbacks for real-time inferencing.

Time aligned pairs: Stems and the final mix should be time-aligned so that temporal relationships and transformations can be learned.

Diverse Instruments and Genres: Better generalizing capabilities during inferencing, exposed to more diverse data

Dry multitrack stems: Wet stems have effects already applied, limiting the model's ability to learn how to apply mixing techniques independently and its hard to remove these effects during training. Dry stems are unprocessed and are "ready to be learned".

It is clear from Table 3.19 that the best dataset to use for training is **MedleyDB** given the sheer size, and variety compared to all other datasets, and what seems to be no drawbacks at face value. It is also the dataset that the originators of the DMC recommend and use. Cambridge Multitrack is second choice and perhaps could be trained alongside MedleyDB if preprocessing is implemented to account for the recordings not being time-aligned or as a comparative dataset to evaluate performance.

Method	Input	Output	Architecture	Adv.	Disadv.
Mix-Wave-U-Net	Multitrack audio recordings (typically 8 channels) in the time domain	Final stereo mix (direct transformation)	Encoder-decoder with 1-D convolutions, skip connections inspired by U-Net architecture	Efficient at mimicking professional mixes (EQ, compression, reverb) without requiring knowledge of individual effects	No user control, artifacts with out-of-distribution inputs, fixed number of input channels
DMC	Mel-Spectrograms of multitrack audio recordings with pre-processed context embeddings for each track	Predicted mixing parameters (EQ, compression, gain, panning, etc)	CNN-based VGGish encoder for feature extraction, shared-weight networks, context embedding, MLP for parameter prediction via Post-processor, Transformation Network for DSP Chain,	Flexibility in input track numbers, user control over mixing parameters, adapts to different scenarios	Complex gradient computation, requires differentiable effects, computationally expensive without proxy networks pretrained beforehand
Fx-Normalization (CRAFx2)	Wet multitrack recordings (processed tracks with applied effects)	Predicted mix with adjusted loudness, panning, EQ, compression	Convolutional and recurrent neural networks (CRAFx) with stacked TCNs and BLSTMs for normalization and mix prediction	Utilizes real-world wet data, adaptable to real scenarios, produces mixes comparable to professional mixes	Requires extensive preprocessing (Fx-normalization), complex architecture and large architecture, requires large datasets for training
Diff-MST	Multitrack audio recordings and a reference track for style transfer	Predicted mixing parameters (EQ, compression, gain, panning)	Transformer-based controller, two encoders (one for input tracks, one for reference), outputs parameters to DMC for final mix	Effective style transfer, adaptable to variable input tracks, allows further control and fine-tuning by the user	Relies on reference mixes, potential for artifacts with mismatched inputs, requires large and high-quality reference datasets

Table 3.18: Comparison of Automatic Mixing Methods

Dataset	Size (Hrs)	no. of Songs	no. of Instrument Category	no. of Tracks	Type	Usage Per-mis-sions	Other info	Remarks
MedleyDB	7.2	122	82	1-26	Multitrack, Wav	Open	44.1KHz, 16 bit, stereo	-
ENST Drums	1.25	-	1	8	Drums, Wav/AVI	Limited	44.1KHz, 16 bit, stereo	Drums only dataset
Cambridge Multitrack	>3	>50	>5	5-70	Multitrack, Wav	Open	44.1KHz, 16/24 bit, stereo	Not time-aligned, recordings for all songs are not uniform
MUSDB	10	150	4	4	Multitrack, Wav	Open	44.1KHz, Stereo	Used mainly for source separation, wet stems
Slakh	145	2100	34	4-48	Synthesised Flac	Open	44.1KHz, 16 bit, stereo	Used mainly for source separation; sometimes wet stems
Shaking Through	4.5	68	>30	>40	Multitrack, Wav	User only	44.1/88.2KHz, 16/24 bit, stereo	-
BitMIDI	-	>1M	>5	>5	Multitrack, MIDI	Open	MIDI data	MIDI data submitted by users across world

Table 3.19: Comparison of Various Multitrack Audio Datasets for Automixing Courtesy of Steinmetz et al.

Chapter 4

Standards and Design Constraints

This chapter is split into two overarching sections, standards and design constraints. It is imperative to have a grasp on both. Standards are critical to engineering projects in general; they provide consistency across the product and are particularly important in interfacing systems. They are particularly pertinent to the plethora of inter-device communications present in this project as well as the outward-facing connectors for audio input and output.

Design constraints, economic constraints, and other practical considerations will guide the project and limit its scope to that which is feasible within the time-frame of Senior Design 1 and Senior Design 2. Without detailing these constraints, certain aspects of the project may never reach completion.

The two sections are outlined below.

4.1 Standards

4.1.1 Audio Hardware

Phantom Power

Title	Standards for Phantom Power: IEC 61938:2018
Authors	International Electrotechnical Committee
Year Created	2018
Description	IEC 61938:2018 gives guidance on the best practices for multimedia analog interfaces to achieve interoperability. Among these is a specification for microphones and the use of phantom power. Three formal standards are specified, P48, P24, and P12. These specifications include the nominal voltages, operating currents, coupling resistor values, and more.
Application	For the mixer, phantom power must be available to power the use of certain microphones. The actual phantom power design will actually not adhere to the phantom power specifications, but it will extrapolate from the specifications to create a design that provides the same amount of power as the formal specifications. This is due to the difficulty of implementing a voltage regulator that matches the phantom power specification.

4.1.2 Communication Protocols

Universal Serial Bus (USB)

Title	USB Implementers Forum
Authors	Apple - Dave Conroy, HP Inc. - Alan Berkema, Intel Corporation - Brad Saunders, Microsoft Corporation - Toby Nixon, Renesas Electronics - Philip Leung, STMicroelectronics - Gerard Mas, Texas Instruments - Anwar Sadat

Year Created	1996
Description	This standard outlines the guidelines and technical specifications for cables, connectors, and communication protocols used to enable seamless interaction between devices such as computers, peripherals, and other hardware. The USB standards, which are developed and overseen by the USB Implementers Forum (USB-IF), were established to create universal compatibility and simplify the connection process across devices. The USB-C connector, our chosen type for this application offers both speed and versatility. Its reversible oval design allows for insertion in either orientation. In addition to offering faster data transfer rates and enhanced power delivery, USB-C supports a broad range of functionalities, including audio streaming through USB audio interfaces. Furthermore, USB-C maintains backward compatibility with earlier USB versions, including USB 1.0 through USB 3.2.
Application	For the mixer, a USB interface will be implemented to facilitate the recording of audio into a host computer. The USB interface will conform to the physical standards of USB-C and the operability of USB 2.0 to ensure maximum compatibility. USB drivers and software are handled by first party libraries for the microcontroller platform, ensuring compliance.

Serial Peripheral Interface (SPI)

Title	KeyStone Architecture Serial Peripheral Interface (SPI)
Authors	Texas Instruments
Year Created	March 2012
Description	The Serial Peripheral Interface (SPI) is synchronous, meaning it uses a clock signal shared between the master and slave devices to ensure data integrity during transmission. SPI uses separate lines for sending and receiving data (MOSI and MISO), allowing full-duplex communication, which enables simultaneous data transfer in both directions.

Application	The mixer will utilize the Serial Peripheral Interface protocol to communicate in between devices on the PCB as well as between the xcore microcontroller and Raspberry Pi 5. Care will be taken to implement all the data lines necessary to utilize SPI for full-duplex communication. In terms of the software implementation, both the xcore and Raspberry Pi 5 support SPI with robust, well-documented libraries. This ensures compliance with the standard and should allow for smooth communications between the connected systems.
--------------------	---

Inter-Integrated Circuit (I²C)

Title	I ² C-Bus Specification and User Manual
Authors	Philips Semiconductors
Year Created	1982
Description	The Inter-Integrated Circuit (I ² C) protocol is a synchronous, multi-master, multi-slave communication standard. It uses only two wires: a data line (SDA) and a clock line (SCL), minimizing the number of connections required. I ² C is byte-oriented and supports bidirectional data transfer. Each device on the bus is identified by a unique address, allowing multiple devices to communicate on the same lines.
Application	The mixer will utilize the I ² C protocol for low-speed communication between the Raspberry Pi 5 and the XMOS xcore. The Raspberry Pi 5 will act as the master, initiating communication and managing the connection with the xcore device. Software libraries provided by the Raspberry Pi ecosystem, such as i2c-tools and smbus, will streamline integration and ensure compliance with the standard.

WebSockets

Title	The WebSocket Protocol
Authors	Ian Fette, Google, and Alexey Melnikov, Isode Ltd.

Year Created	December 2011
Description	The WebSocket protocol is a communication protocol that provides full-duplex, bidirectional communication over a single TCP connection. It begins with an HTTP handshake to establish the connection, then switches to a lightweight WebSocket frame format for continuous data exchange. WebSockets are ideal for real-time applications, allowing servers to push updates to clients without the overhead of frequent HTTP requests.
Application	The mixer will utilize the WebSocket protocol to facilitate communication between the Raspberry Pi 5 backend processes and client modules. Two WebSocket servers will run on the Raspberry Pi: one for state data (e.g., mixer settings) and one for audio data (for visual monitoring and AI inferencing). Each server will be distinguished by a unique port number. The WebSocket protocol was chosen for its low-latency, persistent connection capabilities, which are crucial for real-time synchronization between the user interface and the mixer's backend systems, as well as the relative ease with which connections can be established and utilized

Wi-Fi

Title	Evolution of the IEEE 802.11 Standard: Current Trends and Future Perspectives.
Authors	D'Andrea, A., Silva, A.
Year Created	July 2020

Description	<p>Wi-Fi, short for Wireless Fidelity, is a technology based on the IEEE 802.11 standards that facilitates wireless communication between devices by transmitting data over radio frequency waves. Operating in the physical and data link layers of the OSI model, Wi-Fi enables devices to exchange data without requiring physical connections, providing the convenience of wireless connectivity. The technology supports varying speeds, ranges, and frequencies depending on the specific version of the IEEE 802.11 standard implemented, such as 802.11n (Wi-Fi 4), 802.11ac (Wi-Fi 5), or 802.11ax (Wi-Fi 6). Each iteration improves bandwidth, network capacity, and efficiency. Wi-Fi operates in the 2.4 GHz and 5 GHz frequency bands, with newer standards also supporting the 6 GHz band (Wi-Fi 6E). These frequencies enable a balance between range and speed. It uses technologies such as Orthogonal Frequency Division Multiplexing and Multiple Input, Multiple Output to improve data throughput and handle multiple simultaneous connections effectively.</p>
--------------------	--

Application	Wi-Fi is integral to the communication architecture of our project, serving as the physical and data link layer for wireless data exchange between the Raspberry Pi and the mobile application. This choice is driven by Wi-Fi's ability to provide high-speed, reliable, and long-range connectivity, which is crucial for our system's functionality. The primary role of Wi-Fi in our project is to establish a wireless connection that supports real-time transmission of FFT data from the Raspberry Pi to the mobile app. This connection enables seamless updates to the app's user interface, where the frequency spectrum is visualized in real time. By leveraging the IEEE 802.11 standard, the system benefits from the high data throughput necessary for streaming large amounts of FFT data continuously and efficiently. Wi-Fi's extended range also plays a significant role in our project, allowing the Raspberry Pi and mobile app to remain connected over larger distances compared to other wireless protocols like Bluetooth. This flexibility ensures that the user can interact with the system without being physically constrained to close proximity. Another important aspect is the ability of Wi-Fi to support application layer protocols like WebSockets, which are essential for real-time, bidirectional communication between the Raspberry Pi and the mobile app. WebSockets enable low-latency, full-duplex communication over a persistent connection, ensuring that FFT data is transmitted and displayed without delays.
--------------------	---

4.1.3 JSON Extensions

JavaScript Object Notation Schema (JSON Schema)

Title	JSON Schema: A Media Type for Describing JSON Documents
Authors	IETF JSON Schema Working Group
Year Created	2013

Description	JSON Schema is a standard for defining the structure, constraints, and validation rules of JSON documents. It enables precise specification of the data model, including required fields, data types, default values, and more. JSON Schema ensures interoperability and facilitates robust data validation by providing a contract for JSON-based communication.
Application	The mixer will use JSON Schema to define the structure of its central state, which will be managed by the Raspberry Pi 5. This schema will ensure consistent encoding of state information across the system, including the central state manager process and the Flutter UI. The JSON Schema will specify properties such as channel configurations, effect parameters, and system status, enabling data validation and error checking during state synchronization. Libraries such as jsonschema (for Rust) and json_serializable (for Dart/Flutter) will be used to enforce schema compliance.

JavaScript Object Notation Patch (JSON Patch)

Title	RFC 6902: JavaScript Object Notation (JSON) Patch
Authors	Paul C. Bryan, Kris Zyp, and Mark Nottingham
Year Created	March 2013
Description	JSON Patch is a standard for describing modifications to a JSON document. It defines a list of operations (add, remove, replace, move, copy, test) that can be applied to a JSON object to transform its state. Each operation specifies a target location (using a JSON Pointer) and the intended modification, allowing efficient and precise updates to complex JSON structures. JSON Patch is particularly well-suited for systems requiring frequent updates to shared states, as it transmits only the changes rather than the entire document.

Application	The mixer will utilize JSON Patch for communicating modifications to the central state between the UI, AI, and DSP modules, following the publisher-subscriber design pattern. The central state manager process will mediate these updates, ensuring that changes are applied atomically and propagated to all subscribed modules. JSON Patch allows for the efficient transmission of only the relevant updates, reducing bandwidth usage and computational overhead. Robust libraries for JSON Patch are available for all relevant modules, allowing for easy integration into our system.
--------------------	--

4.1.4 Programming Languages

Rust (programming language)

Title	Rust Programming Language
Authors	Graydon Hoare, Mozilla, The Rust Foundation
Year Created	May 2015
Description	Rust is a multi-paradigm, systems-level language with a focus on safety, concurrency, and performance. It offers strong memory safety guarantees via its unique ownership model, which statically eliminates data races and memory leaks without requiring garbage collection. Its ecosystem includes extensive libraries for safe systems programming, audio processing, and more.
Application	Rust will be used to implement the state manager module on the Raspberry Pi 5. It is the natural choice due to the need for low-level control over hardware interfaces between the Pi and other devices in conjunction with high-level constructs for abstracting our custom data types.

Dart (programming language)

Title	Dart Programming Language
Authors	Lars Bak and Kasper Lund, Google
Year Created	October 2011
Description	Dart is a client-optimized programming language designed for building web, mobile, and desktop applications. It features a sound type system, a rich standard library, and asynchronous programming support. Dart compiles to efficient native code or JavaScript, making it highly versatile. It is the primary language used in the Flutter framework, which simplifies UI development with a declarative, reactive programming model.
Application	Dart will be used to develop the user interface for the mixer, leveraging the Flutter framework. Flutter's compatibility with Dart allows for efficient development of visually responsive, cross-platform UIs for both the Raspberry Pi and tablet devices. Dart's asynchronous programming features, such as Future and Stream, will be particularly useful for managing WebSocket communications between the UI and the backend processes, ensuring smooth and responsive user interactions.

4.2 Design Constraints

Below is a summary table of the main constraints upon the mixer project. Their details and impact on the project are detailed in the sections below.

Constraint	Description
Cost	The ideal budget for this project should be under \$400, though \$500 may be required. There is no definitive maximum, but the intent is for it to be competitive with mixers on the market.
Time	The current semester of SD1 gives roughly 3 months (or 12 weeks) of time to properly plan and develop the software and hardware prototype for this device. Once in SD2, we will have another 3 months to finalize our project before the final deadline. Using our time wisely during SD1 is the most crucial part for the success of the project.
Hardware Constraints	Certain unavoidable hardware details place constraints on the design of the mixer. Part of the design process is finding suitable workarounds/solutions to such constraints.
Social	The mixer lies in the realm of music and the creative arts, and as such, the developers bear a level of responsibility to ensure the mixer is not used for offensive content or other negative purposes.

Table 4.11: Identified Constraints

4.2.1 Economic Constraints

In regards to economic constraints, the mixer project is self-funded, and as such, significant budgetary constraints are placed upon the project due to the limited resources of the developers. The most straightforward way in which the budget impacts the project is in the price of the selected components. Where possible, the cheaper option is used.

An example is in the operational amplifier selection. The NE5532 operational amplifier is quite old and does not perform like the best operational amplifiers on the market. However, audio op-amps are already quite good across the board. In order to reduce the cost of the hardware, the NE5532 was chosen due to its tremendous value.

Another example is in the selection of the single-board computer. In many ways, the Beaglebone AI-64 was the best option. It boasted vastly better audio IO compared to the Raspberry Pi 5 as well as a plethora of accelerator hardware components on the SOC which would have befitting our auto-mix algorithm. However, the cost was almost triple that of the Raspberry Pi 5. The extra features could not make up for the huge difference in cost.

There are several less obvious ways in which the economic constraints of the

project impact the design process. Among these is the way in which budget impacts the iterative nature of the prototyping process. While it may be ideal to test a variety of parts for each type of component, the budget constraints make it so that for most components, it is most viable to simply construct and test a single part for any given hardware design. Additionally, it is imperative that great care be taken care of to minimize the amount of hardware iterations necessary to create a functioning prototype. Beyond the repeated additions to the overall cost from the price of parts and board re-fabrication, another consideration is the shipping cost. Most electronic parts are unavailable from online shopping platforms that offer free shipping like Amazon. Instead, the predominant marketplaces like Digikey and Mouser charge money for shipping. As such, it is imperative to ensure the inclusion of all necessary parts in each order as well as the inclusion of enough spare parts to minimize the number of repeat orders.

The restriction on repeated parts reordering and board re-fabrication also factor into the constraints on time. For one, parts shipping is often a lengthy process. Any lack of parts will set back prototyping by several days. Additionally, board manufacturing and shipping is a very hefty time bottleneck. It can take upwards of several weeks to receive an order. Board orders must be thoroughly inspected and re-checked before they are sent in order to expedite the design process.

4.2.2 Time Constraints

Time is a significant constraint for our project due to several hardware details. For one, the fine pitch of the XMOS xcore microcontroller pins means that board assembly cannot be done by hand. This introduces extra delays in the production of a board and also prevents quick rework. The necessity of a board that is easy to iterate on and rework has directed the PCB level design choices to emphasize modularity and reworkability through the use of expansion headers and breakout boards.

Time constraints also significantly influenced our software development process, shaping nearly every decision we made regarding architectural design and choice of technologies/frameworks. Throughout the project, we consistently encountered a trade-off between performance and development time. While more performant solutions often yield higher efficiency, they also require significantly more effort to develop, test, and integrate. Given our use of the Raspberry Pi and xcore boards, whose processing capacity exceeds the computational requirements of our system, we consistently prioritized minimizing development time over achieving maximum performance.

Many of our software design choices, as outlined during our research and comparisons in Chapter 3, were made explicitly with these time constraints in mind. By favoring technologies and frameworks that enable rapid iteration and implementation, we can accelerate development progress without outright sacrificing function.

For instance, we opted to establish communication between modules hosted on the Pi via WebSockets rather than more performant inter-process communication (IPC) methods. While the use of WebSockets for internal communication within a

device inherently introduces overhead, it is nonetheless easier to implement and use, reducing both development and debugging effort. Similarly, we chose to use a general-purpose operating system (Raspberry Pi OS) over a bare-metal implementation. This decision allowed us to leverage well-documented libraries and tools, eliminating the need to manage low-level hardware interactions, which would have demanded considerable time for comparatively little gain.

Furthermore, the adoption of Flutter, a high-level, modern UI framework, exemplifies our commitment to rapid development. Flutter's expressive API and pre-built widgets enables us to rapidly build a functional user interface with minimal time investment, compared to the complexity of implementing a lower-level solution, like Qt. Flutter's hot-reload functionality is also a great boon, given the heavy amounts of trial-and-error required in prototyping the integration of our software modules together.

The time constraints on our project has weighed heavily throughout the design of our hardware and software, and as a result we have consistently made the design decisions that enable us to develop and test according to our project timeline, while still delivering a fully functional and robust system.

4.2.3 Hardware Constraints

As touched on above, the actual details of the hardware used for the mixer project impose certain constraints on the design process as well as the design itself. While it may be argued that the entire point of engineering is to work around and figure out the hardware that is available, there are some ways in which these hardware constraints are so unworkable such that they influence the part choice and other critical design decisions.

One example is the way in which the Raspberry Pi 5's IO situation affected the design of the mixer. While the Raspberry Pi 5 was indeed the best choice for the project due to its ability to cover the AI auto-mix functionality and wireless communications functionality, its limited audio IO complicated the design of the system. Because the Raspberry Pi 5 only has 2 channels of I2S, the full 8 channels of audio input cannot be streamed in real-time to the Pi. This restriction is particularly challenging to the design of the reactive audio GUI and the auto-mix feature, both of which require the input audio to operate. This hardware restriction necessitated some design compromises to be made, such as the onboard transformation of the audio on the microcontroller to be sent in small packets over SPI instead of the live audio being transmitted.

4.2.4 Accessibility of Datasets

For production use, MedleyDB, MUSDB, Slakh, Cambridge Multitrack, and BitMIDI are generally suitable due to their open licenses, though some may require attribution or have minor restrictions on commercial use—reviewing specific license terms. MedleyDB, MUSDB, and Slakh are especially favorable, as they offer open access with minimal constraints. In contrast, ENST Drums and Shaking Through

have restrictive or limited permissions, making them unsuitable for commercial production without additional licensing agreements or explicit permission but can still be used for educational purposes.

4.2.5 Social Constraints

The social aspect of the mixer is largely constrained and governed by its existence in the sphere of music and art. In as much as the mixer can be a helpful tool for artists who want to spread their legitimate artistic messages, bad actors can also use the mixer for purposes contrary to what is helpful or ethical. While the burden does not directly lie on the developers to mitigate such issues, it may be worthwhile to consider how the design of the mixer can be thought out so that it is unlikely to be used for purposes outside of what is intended.

Chapter 5

Comparison of ChatGPT with Other Similar Platforms

For better or for worse, ChatGPT and other LLMs have established their presence in our academic and professional lives. In this chapter, the ramifications of such a technology will be discussed. The topic is split into several overarching themes. One is a brief comparison between ChatGPT and other LLMs available to the students as well as a personal assessment of the usefulness and limitations of LLMs in application to our project. The next section will present examples of how ChatGPT has impacted the learning and research experience of the students.

5.1 Limitations, Pros, and Cons

A key advantage that ChatGPT has over its competing platforms is the immense financial endowment that it has compared to other LLMs. Research and development, training, and inferencing are all incredibly expensive activities, computationally and financially. The increased funding that OpenAI possesses translates to faster generation, more accurate and helpful inferences, more advanced models – all around, an enhanced user experience.

However, it is possible for LLMs to specialize for performing certain types of tasks. Microsoft's *Github Copilot* is a fine-tuned GPT-4o model, developed to assist with programming. It is unique in that it has been trained on snippets of code, can directly draw from public GitHub repositories, and operates within a text editor/IDE, offering real-time code suggestions to the programmer. Github Copilot is an excellent example of an alternative LLM implementation that offers increased utility to a developer compared to the chatbot interaction format of ChatGPT.

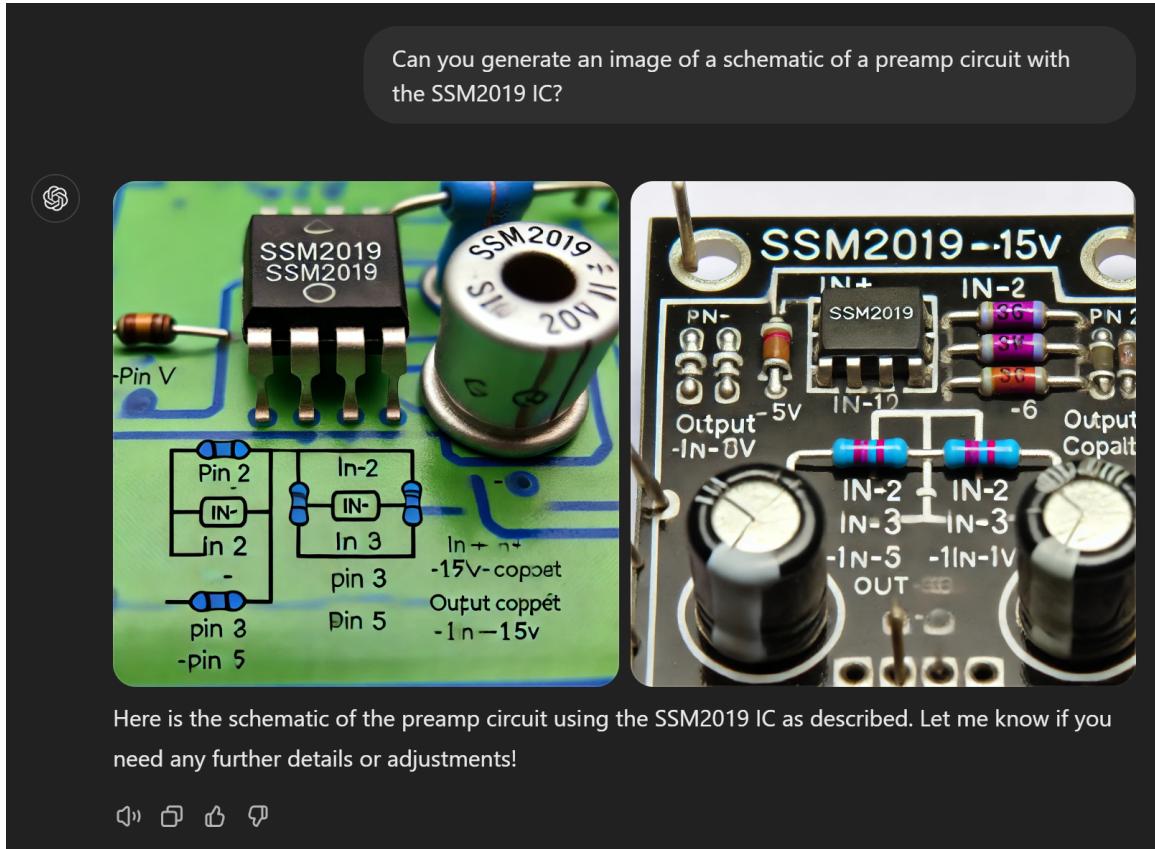
Another example of an LLM with specialized utility is Google's Gemini, specifically the Gemini instance embedded within Google search functionality. Gemini provides helpful summaries and explanations that directly answer the question you have typed into Google search. This is in contrast with the normal Google search functionality, which simply gives you hyperlinks to websites; you still have to look on the website yourself and infer the correct information. This gives Gemini a similar functionality to ChatGPT. However, it is superior in terms of answer reliability. Gemini, in addition to directly performing inference off of the search data, offers reference links showing where it got its answers from. That way, you can verify the legitimacy of its answers. This makes it particularly useful for research, where LLMs may tend to hallucinate and give you untrue or inaccurate answers to questions. Such is another example of a specialized LLM which extends the utility of LLMs beyond that which is offered by ChatGPT alone.

In terms of its overall utility, ChatGPT and other LLMs have proven quite useful to the students in their research for the mixer project as well as their studies in other fields. It functions as an enhanced web browser search which presents information in an organized and readable way. If accuracy is paramount, results can simply be verified through a browser search after a ChatGPT prompt.

However, there are certain limitations to ChatGPT and other LLM platforms. For one, LLMs offer drastically varying amounts of utility depending on the field of study. One may note that ChatGPT, in popular culture, has particularly affected the plight of computer science majors and software developers. LLMs are particularly useful in coding because of their text-based chat-bot format. One can simply ask some sort of prompt and they will receive a batch of ready-made code. Of course, the code may still need refactoring to work as intended. For one, LLMs are not omniscient in their coding abilities. Another factor is that LLMs typically cannot see the context of larger coding projects. This is even true of IDE-integrated LLMs like GitHub Copilot, although certain "software engineering" LLMs have emerged (e.g. Devin) with limited success. As such, LLMs have limited utility in large codebases. Code snippets still necessitate the developer to understand how to integrate them into the code base.

However, with that said, LLMs offer drastically more utility to software engineers than they do to electrical engineers or other disciplines. In the case of electrical engineering, LLMs cannot produce an output such that the output can immediately be utilized or incorporated into a larger work. LLMs cannot output things useful to electrical engineers such as schematics, footprints, and PCB layouts. At best, they can describe what needs to be done in a verbal manner; however, such output still requires a knowledgeable engineer in order to incorporate such verbal sugges-

tions into a real engineering design. One may suggest that there are generative image AI platforms incorporated into LLM platforms which can theoretically generate schematics and designs. However, such generative AI technologies are so laughably inaccurate such that they are entirely useless for engineering work. Below is an example of an attempt to generate a schematic of a preamp circuit with the SSM2019 IC? The prompt used was "Can you generate an image of a schematic of a preamp circuit with the SSM2019 IC?" The prompt was fed into ChatGPT, specifically the 4o model.



One may clearly see that the output is totally incoherent. Not only is the output itself not a real schematic, but the connections and text presented in the image output are complete gibberish. Generative image AI is particularly deficient when it comes to text, and the results are clear.

5.2 Examples

There are many ways in which LLMs have affected the learning experience of the students. Below are several examples. Most of them are positive, albeit some are negative.

Model	Specialization	Key Features and Limitations
ChatGPT	General-purpose LLM	Provides natural language responses and assists with a wide range of topics. Useful in research and as an enhanced search tool, but prone to inaccuracies (hallucinations) and less specialized for technical fields like electrical engineering.
GitHub Copilot	Coding and IDE integration	Fine-tuned on code, operates within IDEs for real-time coding suggestions. Useful for software engineers, but limited in context-awareness for large projects and requires user knowledge for integration.
Google Gemini	Search integration and information reliability	Offers direct answers with reference links for verification, making it more reliable for factual queries and research tasks. Limited to search-related applications and less interactive than a chatbot.

Table 5.1: Comparison of ChatGPT, GitHub Copilot, and Google Gemini

5.2.1 Automation of Boilerplate Tasks

One incredibly useful manner in which LLMs have been useful in the process of this project is the automatic generation of LaTeX tables. This document is typeset in LaTeX. For the uninitiated, LaTeX is a typesetting system for technical papers that offers incredible flexibility and aesthetically pleasing documents. However, there is some increased complexity when it comes to the creation and formatting of tables. Tables are not visual elements that can be modified with mouse clicks or simple dragging actions. They are typeset with code, and as such, it can be cumbersome to create them. However, LLMs have created an incredibly simple workflow booster. With this new workflow, I can create a table in Microsoft Word or some other document editor, take a screenshot, and paste it into an LLM. Upon prompting, the LLM will immediately generate the necessary code to create the table in LaTeX, reducing overhead with LaTeX typesetting significantly.

5.2.2 Information Retrieval

LLMs can be seen in the context of the evolution of information discovery and retrieval mechanisms. Every advancement in this process involves the automation of tedious, repetitive tasks involved in procuring knowledge relevant to users. For instance, consider how search engines automate the process of finding relevant information given a query; the alternative would have been manual trawling of indexes and consciously identifying relevant entries, diverting time and attention away from the ultimate purpose of research – the synthesizing of information to solve a novel problem. In this light, LLMs can be viewed as a supercharged search engine, capable of taking complex, natural-language queries and producing an overview of relevant information for a highly specific problem, delivered in a narrative format.

In this fashion, LLMs have proved invaluable in their ability to illuminate "unknown unknowns" while researching complex, expansive topics. These are concepts and sources that one would not even know to solicit more information about, as opposed to "known unknowns". The uncovering of "unknown unknowns" is a task that search engines are particularly weak at – their strength of returning highly relevant information given a specific query is also a weakness, as search engines are of little help in navigating concept networks laterally, necessitating, again, the manual trawling of search results and conscious identification of gaps in one's existing knowledge.

5.3 Conclusion

For better or for worse, LLM-based AI is a technological innovation that is here to stay. The titanic investments being made by the US government [64] and private investors [65] into AI companies and research go to show the important role they will play in tomorrow's economy. As such, all professionals employed in knowledge work must learn to adapt to their presence, and understand the implications that it has for their work. For such professionals, the most realistic and beneficial view they can take regarding LLMs is to conceive of them as *tools*. Like any other tool, LLMs are best used as a means of performing work more efficiently, by making certain problems require less manual labor to overcome.

Inherent to the nature of tools is that a given tool will be helpful with performing some types of tasks, and useless with others. They are *not* a panacea; LLMs are *not* AGI, and the path to AGI will most likely *not* consist of an exponential runaway of the intelligence of LLMs [66]. The overestimation of the abilities of LLMs are a mistake made by many, especially by the professional-managerial class, consisting almost entirely of knowledge workers who generally feel as if their livelihoods are existentially endangered by advances in AI [67]. The fear of AI, and automation in general, eliminating one's highly specialized position in society is a valid one to harbor. However, LLMs specifically are unlikely to be the harbinger of great upheaval. Like any other technological innovation, LLMs will in fact cause transformation in

the economy, contributing to the loss, change, and addition of jobs. But ultimately, the maximum capabilities of LLMs are bottlenecked by their fundamental nature – they are simply memorization machines, no more, no less.

Chapter 6

Hardware Design

The following chapter details the hardware design of the digital mixer.

6.0.1 Preamplification Circuitry

The input topology of the mixer includes 4 microphone/line combo inputs and 4 dedicated high-impedance (Hi-Z) instrument inputs. As such, this represents eight discrete preamplifier circuits split into two designs. The designs are shown below.

Microphone and Line Preamp

Below is the schematic for the microphone and line combo preamplifier circuit. As mentioned in the research section on input levels, the total input signal chain must offer at least 40 dB of gain for microphone inputs and a small gain (<10 dB) for line inputs. Additionally, the preamplifier must be able to supply phantom power to the XLR input for microphones. The requirements are satisfied through the design presented.

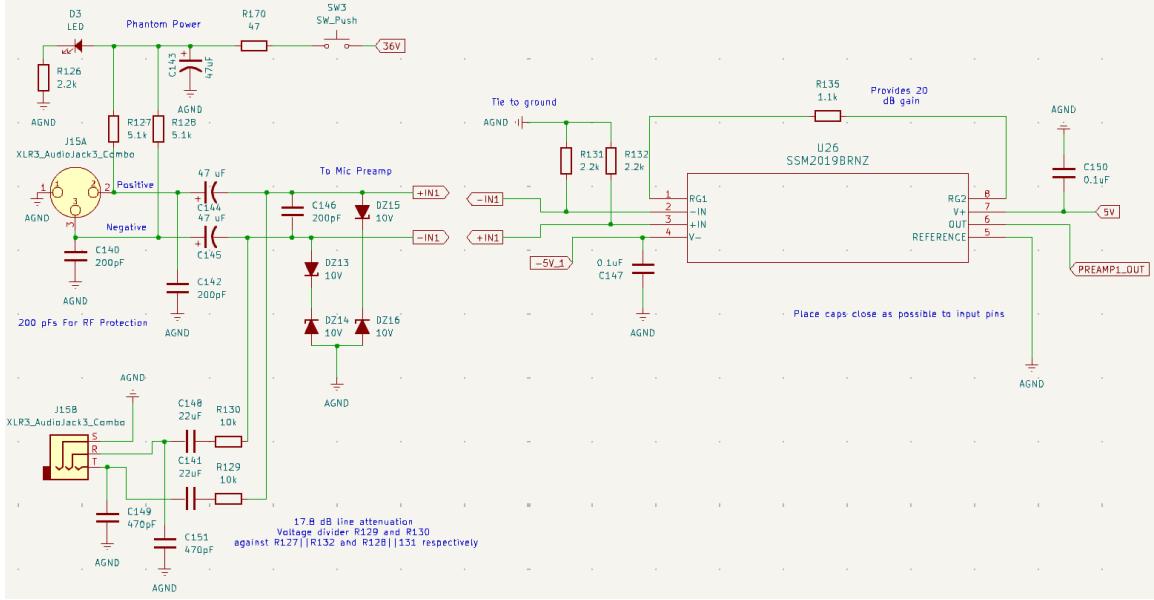


Figure 6.1: Schematic of Microphone/Line Preamplifier

There are quite a few design details of note. First are the various EMI protection capacitors strewn throughout the circuit. There are capacitors attached to the XLR and TRS inputs with values chosen based on THAT Corporation's design note [68]. Additionally, there is a 200 pF capacitor attached between the differential inputs into the preamplifier, with a value chosen based on the Analog Devices SSM2019 datasheet [69]. Similarly, there are 0.1 uF decoupling capacitors placed as close as possible to the supply pins of the SSM2019 IC.

Also of note is the phantom power implementation. The phantom power implementation uses 36 volts. This design is not adherent to the IEEE phantom power specification and is brought into the circuit with custom resistor values of 5.1k Ohms. This design is based on a reference design where the author utilizes custom phantom power voltages [12].

For the line input, the voltage is attenuated by 17.5 dB. This attenuation will allow for extra headroom for the line-input to operate without clipping the operational amplifier or the ADC. The attenuation is achieved by a voltage divider. The T and R line signals are attenuated by a voltage divider between R129 and R127—R132, and R130 and R128—R131, respectively. This yields a voltage attenuation of

$$\text{Attenuation}_V = \left(\frac{5.1k \parallel 2.2k}{10k + 5.1k \parallel 2.2k} \right) = 0.133$$

which translates to a decibel attenuation of

$$\text{Attenuation}_{dB} = 20 \times \log(0.133) = -17.5 \text{ dB}$$

Once the microphone or line signal is brought into the preamplifier IC, a gain of 20 dB is applied. This gain is set by wiring a 1.1kOhm resistor between pins

1 and 8 of the SSM2019 IC as seen in the diagram. This is straight from the datasheet [69].

Because of the capability of the ADC to amplify the signal by over 30 dB, the preamplifier design above meets the gain requirements.

High Impedance Instrument Preamp

Below is the schematic for the high-impedance instrument preamplification circuit. As mentioned in the research section on input levels, the total input signal chain must offer at least 20 dB of gain for instrument inputs. The requirements are satisfied through the design presented.

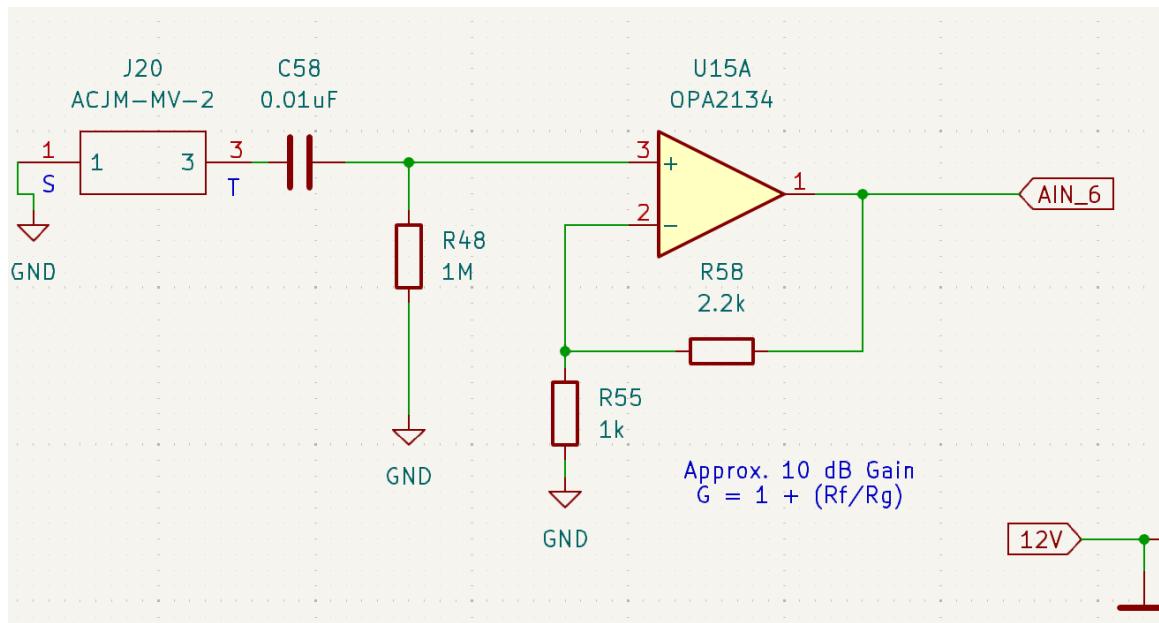


Figure 6.2: Schematic of Instrument Preamplifier

There are several design features of note. One is the high-impedance RC filter created by the 0.01 uF capacitor and 1 MΩ resistor. Another is the configuration of the op-amp. The op-amp is configured in a non-inverting configuration to preserve the phase response of the input. The gain is configured with the feedback resistors.

$$Gain_V = \left(\frac{2.2k}{1k} + 1 \right) = 3.2$$

This translates to a decibel gain of:

$$Gain_{dB} = 20 \times \log(3.2) = 10.1 dB$$

Because of the capability of the ADC to amplify the signal by over 30 dB, the preamplifier design above meets the gain requirements.

6.0.2 LDO Circuitry

As mentioned in previous sections, LDOs offer a near-DC output voltage that is able to be done with very minimal circuitry. In addition, LDOs have significantly less noise than Buck converters at the cost of being more energy-inefficient. Because we are working with analog devices that need noise-free power supplies, using an LDO is a must for those analog power pins. Because we are only using the LDO voltage output for specific analog power pins, we do not need high current output, allowing the low-cost NCP1117 to be used. Similarly, we can also use a TPS73601 for lower voltage outputs. Shown below is the very simple circuit needed to operate the LDO, with only a few capacitors/resistors required.

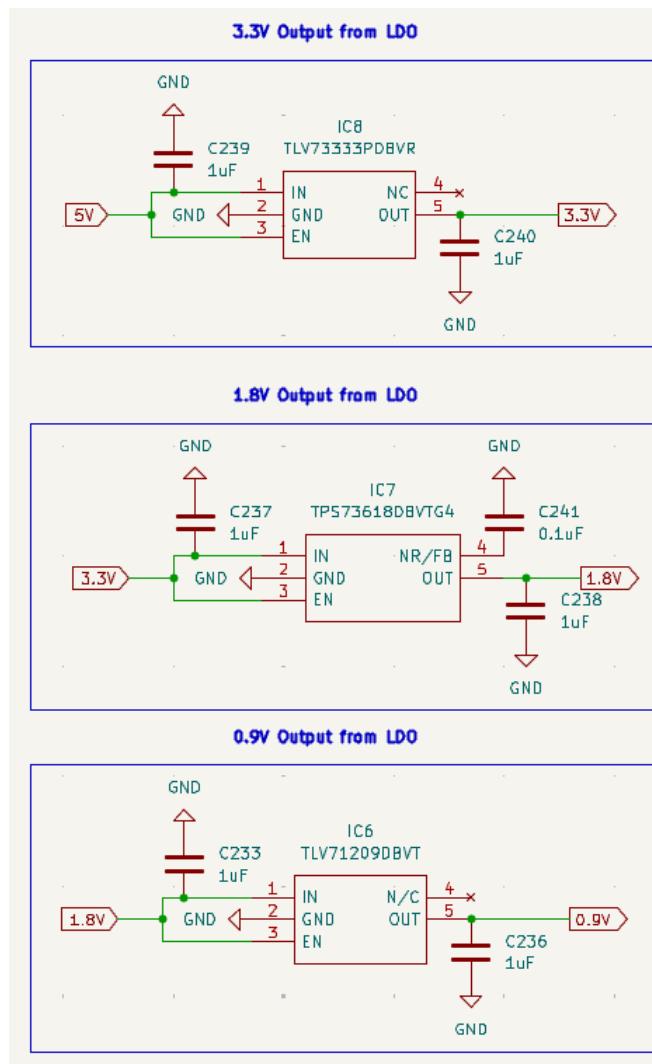


Figure 6.3: Schematic of the three LDOs used in power circuitry

6.0.3 Phantom Power Circuitry

While the standard for phantom power is officially 48V, the vast majority of microphones (besides extremely expensive condenser mics) do not require the full 48V, rather stepping the 48V down to a lower voltage. There is an excellent reference design online that utilizes 15V for phantom power [12] but a better design for the mixer's design is 36V, which would mitigate the overall current required. To do this, we are using the LT8338 as a Boost Converter to boost our original 15VDC input into 36VDC to power all 4 of our phantom power enabled inputs. As the voltage is relatively high for phantom power, the overall current used would be much less than the maximum 100mA that the LT8338 can output.

Shown in Figure 6.4, we are also using the basic phantom power circuitry of having two equivalent resistors on the positive and negative pins (6.8K for 48V) in addition to electrolytic capacitors for smoothing and zener diodes for voltage clipping. In this case, we used 10V zener diodes to ensure that the maximum voltage that the microphones get is around 10V, which is more than sufficient for phantom power operation and protects circuitry. There is also a simple enable switch with LED to indicate the use of phantom power for the mixer. The one switch pictured would indicate that phantom power is enabled for all four channels of the TRS/XLR combo-jack. After the XLR channels receive phantom power and are smoothed and clipped, the signals are then sent to the preamplification circuitry to then be processed.

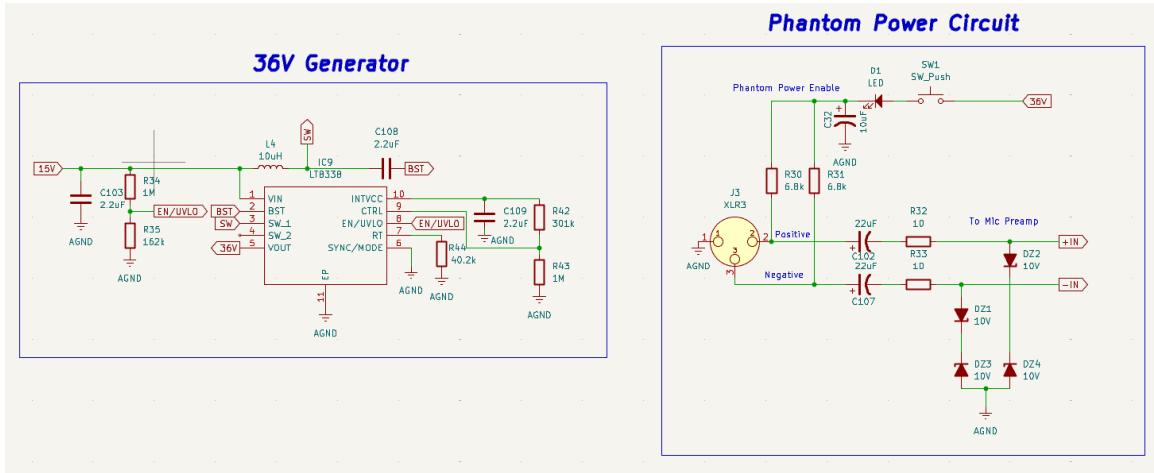


Figure 6.4: 36V Generator and Phantom Power enabled XLR circuit

6.0.4 ADC Circuitry

After the input signals are amplified, they must be converted to digital representations to be processed by the microcontroller and DSP chip. To do this, one must use an ADC, and since we are working with audio, we need to ensure that the signal quality is not degraded through conversion. To do this, we require a high SNR (ideally above 80) coupled with a sampling size of over 16 bits. While our mixer

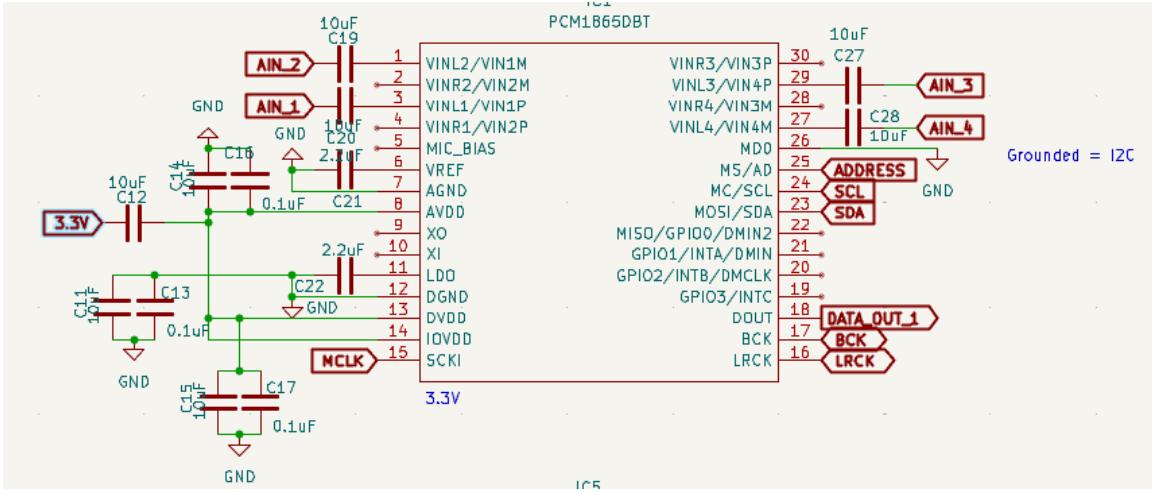


Figure 6.5: Schematic of one ADC

will require decently high quality conversion, due to budgetary and practical limitations, a 44.1KHz sample rate is sufficient as we are not working with specifically Hi-Fi audio technology. This limitation ultimately allows the processing of the data to be significantly easier on the microprocessor itself, lowering the required RAM by multiple factors.

Due to the fact that we are going to be mixing many audio inputs, we chose an ADC that allowed for I2S communication while having the highest number of inputs possible. Using the PCM1865, we are able to have four channels be output per ADC and thus two will be required. This adds a slight level of complication to the routing of signals, as proper separation between the analog and digital pins is very important to mitigate noise.

6.0.5 DAC Circuitry

After processing the signals, they must be converted back to analog to be heard by the musicians and audience. Just as bitrate mattered for the input conversion, it also matters for the output conversion with how accurate the final signal will be in comparison to the digitized version. Inside the mixer, there only needs to be a singular DAC as the two desired outputs (line-out and headphone-out) can be obtained from the same signal after splitting it. If there were more desired audio buses, more DACs would need to be utilized. The DAC will be communicated to via the xcore DSP chip, similar to the ADC, and will be clocked in unison with the other microcontroller chips. There are very few external parts required in this specific DAC circuit, with only some resistors and electrolytic capacitors being sufficient for proper regulation. Although most ADCs and DACs can be powered by either 3.3V or 5V, choosing to operate with 5V allows the SNR to be higher and reduces sampling errors due to a larger reference voltage.

The chosen DAC (PCM1780) requires filtering after conversion to ensure audio integrity, which requires a second-order low-pass filter to prevent any high-

frequency noise (e.g., switching) from being introduced. After going through the low-pass filter, the two audio signals going out would be the left and right audio channels. These signals are going to split off: one copy goes to the Headphone Amplifier for proper amplification and another goes to the line-out TRS connection for use as an input in an amplifier or PA system.

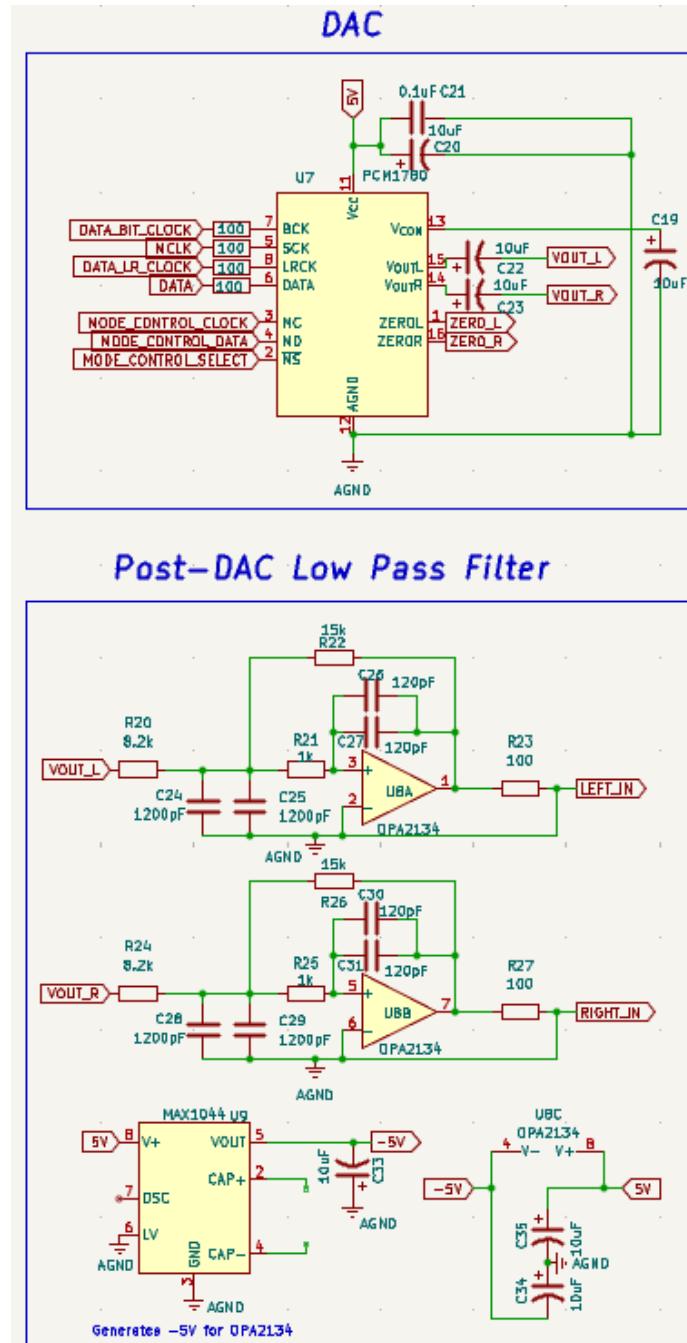


Figure 6.6: Schematic of DAC circuit and Post-DAC Low Pass Filter

Because the low pass filter is essential in maintaining the quality of the output

signal, simulating the frequency response in a program such as LTSpice is helpful to confirm the proper cutoff frequency is obtained from the given resistor and capacitor values. The schematic placed into LTspice is shown by Figure 6.7.

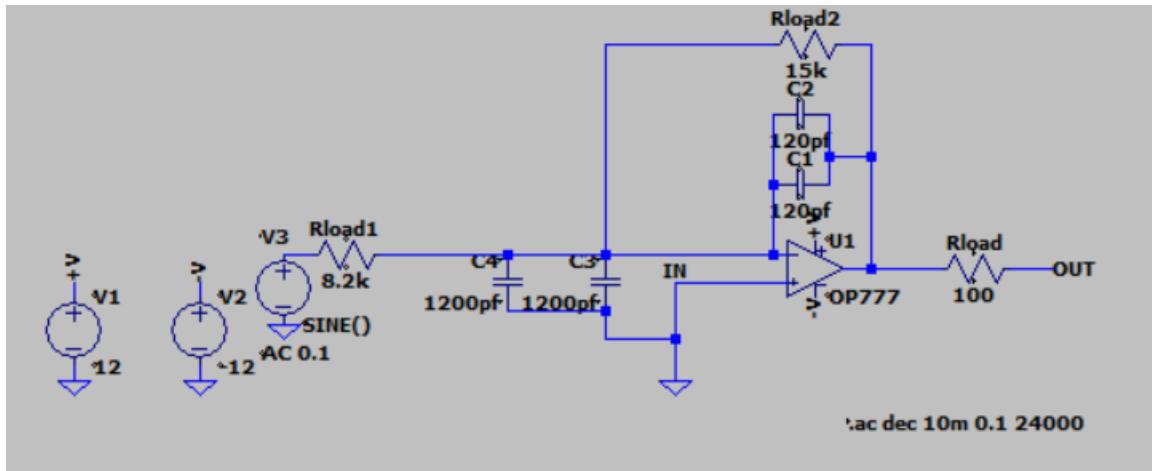


Figure 6.7: Simulation Schematic used to find frequency response

After running the frequency sweep up to 100KHz (over 4x the maximum auditory frequency), the results of the low pass filtering are shown below in figure 6.8:

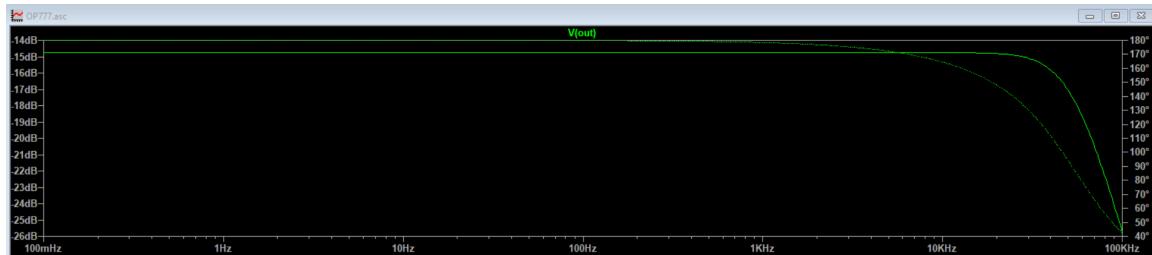


Figure 6.8: Frequency response of Post-DAC Low Pass Filter

It is evident that the cutoff (3dB) frequency is around 50KHz, well over the Nyquist sampling frequency for audio. The main purpose of this low pass filter is to filter out extremely high frequency noise signals, to which it will succeed for anything above 100KHz.

6.0.6 Differential Outputs

While having a single-ended output is fine for most audio applications, creating a differential output allows for a higher SNR and thus less degradation of the output signal. With this in mind, it is in the best interest to split the output signals from the DAC into a differential output using a simple Op Amp circuit, shown below in Figure 6.9. This specific circuit turns a single ended input (for both left and right

channels) into a + and - version to be used in the output for both line out channels. This is done by using an inverting op amp setup with a gain of -1. As used earlier, the gain equation for an inverting op-amp is $G = -\frac{R_2}{R_1}$, and in this case our resistors are both the same value, thus giving us unity gain. This is very practical as stereo line out signals do not really get used in the real-world—speakers will be both to the left and right of the players, after all.

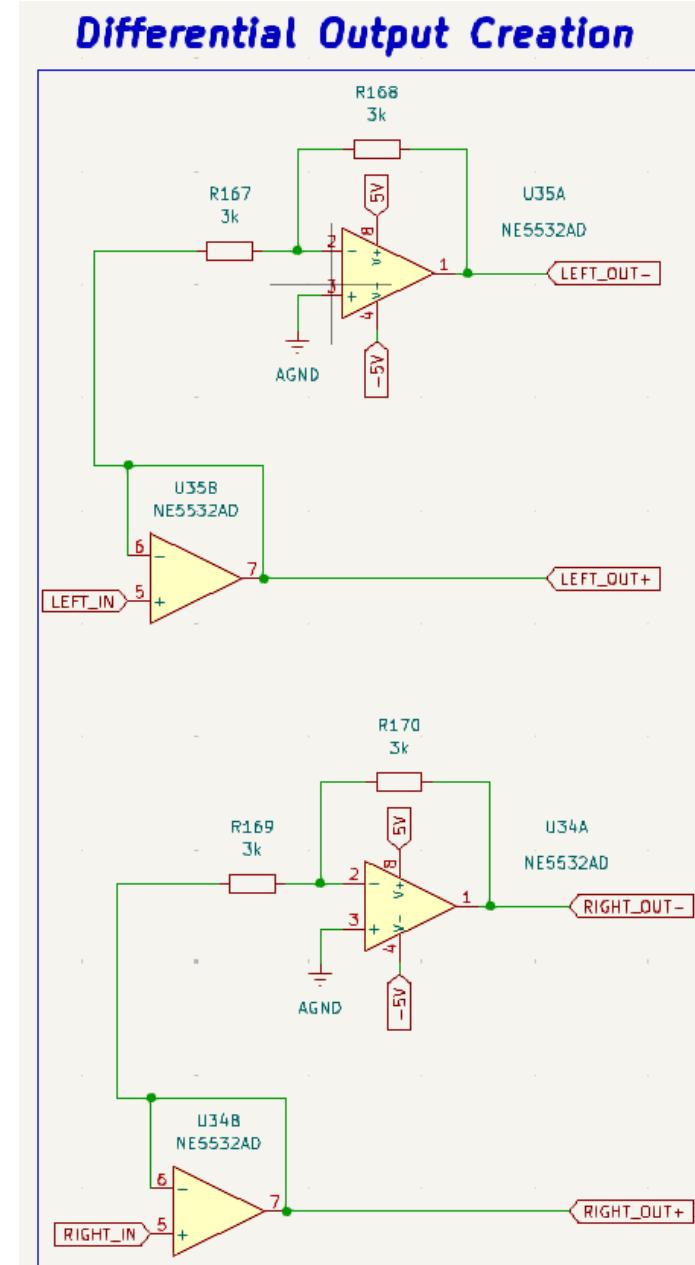


Figure 6.9: Schematic of Synthesized Differential Output

6.0.7 Headphone Amplifier

After being converted back to an analog signal, the end user must be able to hear it. With a headphone output being standard for an audio mixer, it is only natural that our mixer also include one for monitoring purposes. To properly transmit a signal to headphones, the line-out signal (which is the signal obtained from the DAC and then sent through a buffer) must be amplified to a level that headphones are able to drive. This technology is similar to a preamplifier (by way of being a post-amplifier) and there are many ICs on the market that are able to do this without introducing noise into the system.

In our case, we are using the TPA6110A2DGN, which offers a high SNR and requires minimal external circuitry (capacitors and resistors) to properly operate. Once output, the lines are then sent to a TRS connector for output, where the T connection corresponds to the left audio channel and R corresponds to the right audio channel (and the S connector always being the reference ground connection).

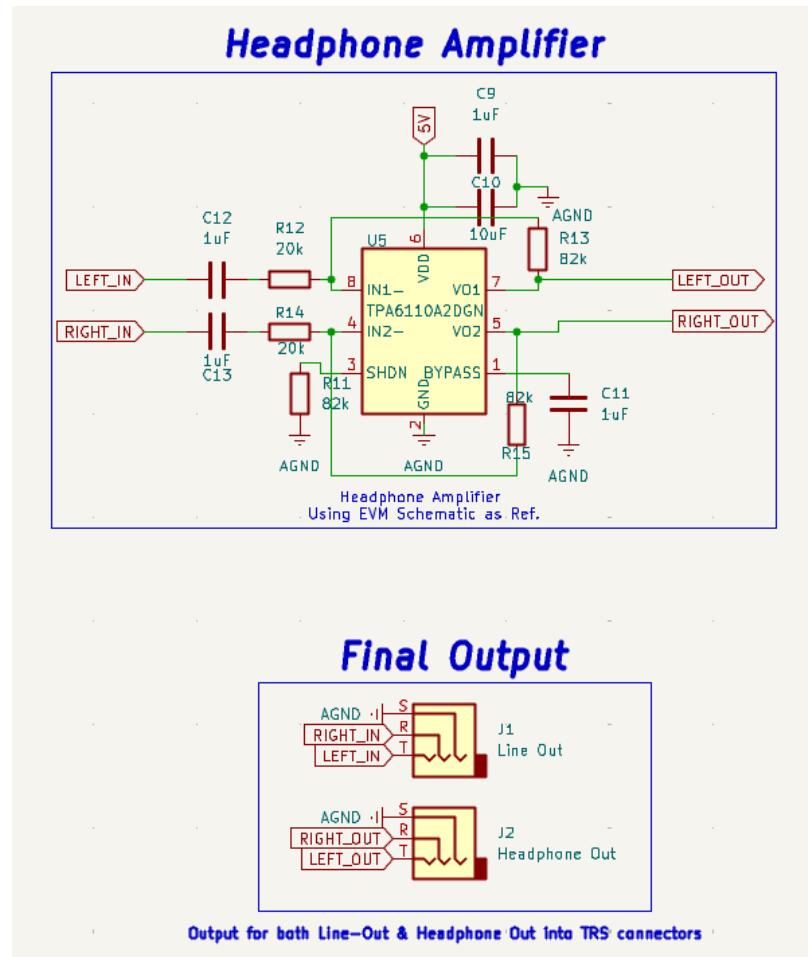


Figure 6.10: Schematic of Headphone Amplifier and Line/Headphone Output

6.0.8 Proper Grounding

The circuitry within the mixing device incorporates multiple grounding systems, specifically for analog and digital components. The distinction between these grounds is critical due to the varying levels of noise present in the signal path. Analog circuits, which typically have low current requirements, are highly sensitive to noise, making proper grounding essential for maintaining audio quality [70]. Therefore, the analog ground must be isolated from the digital ground on the PCB to prevent interference from devices with high switching frequencies, which can introduce noise into the signal path.

In contrast, digital grounds are designed to handle noisy signals, as digital devices are less affected by noise. This is particularly important in power regulation components like buck converters, which operate at high switching frequencies. Ensuring that analog and digital signals are routed to separate ground planes is crucial for the reliable operation of the circuits.

Although analog and digital grounds must be separated to function correctly, they still need to be connected at a single point, known as the "star ground." This narrow junction allows for a shared reference voltage between the two planes while minimizing the switching noise that could affect the analog ground. Proper placement of the star ground is essential for maintaining signal integrity. As illustrated in Figure 6.11, keeping analog and digital components as isolated as possible is key to ensuring a properly functioning circuit.

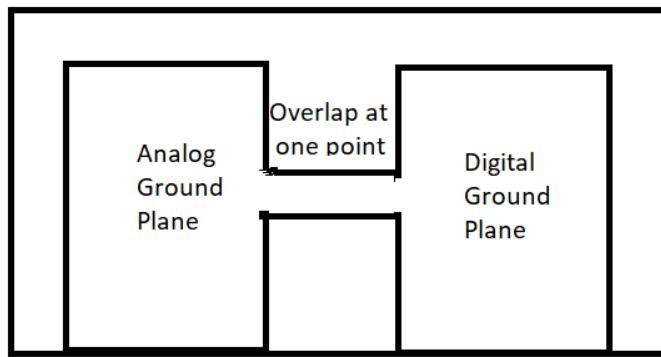


Figure 6.11: Example Organization of Mixed Signal Circuit

6.1 Overall Hardware Design Schematic

As shown in Figure 6.12, the overall hardware design is very interconnected and has many moving parts surrounding it. The connections between specific pieces of hardware are shown, but specific voltage connections are omitted due to spacing concerns. The required voltages for each IC are shown in the sides of the schematics.

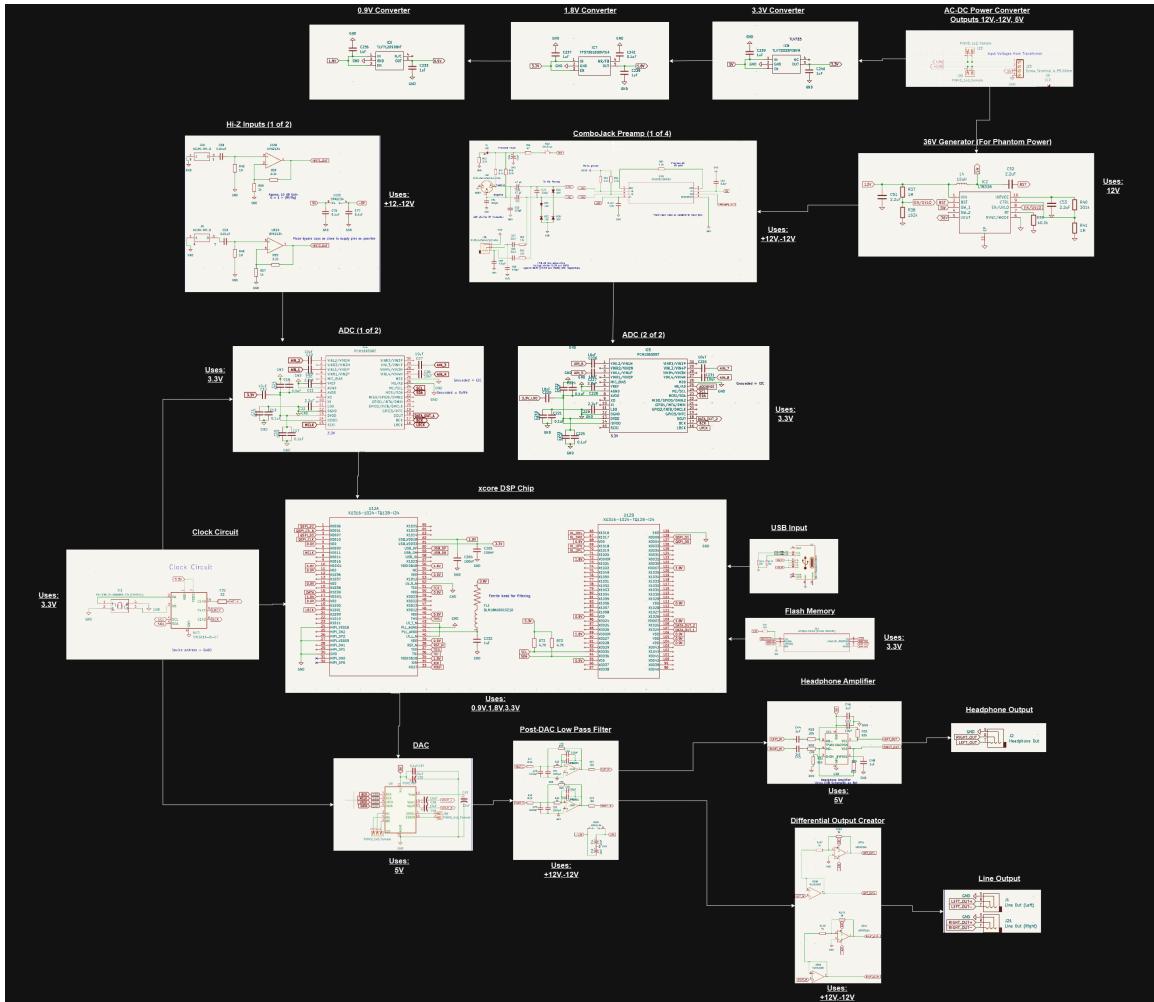


Figure 6.12: Overall Schematic of Entire Circuit

Chapter 7

Software Design

7.1 Software Requirements and Constraints

Our project aims to be a fully-fledged digital mixer, and we aim to deliver a product that will meet professional constraints for audio quality and allowing for diverse

functionality. The software that is required to meet our project's objectives requires a substantial amount of engineering effort.

To this end, we will adopt modern software development methodologies to guide the development of our software architecture – namely, the *Agile* methodology. The starting point for project development under Agile is to formally and exhaustively define the requirements of our software, from the perspective of the end users.

While we have outlined the overall goals of our project earlier in Section 2.3 of our report, here we will specifically focus on the requirements of our mixer's *software*, as well as other relevant information that have implications for our software design, and discuss how it will inform our architecture's design.

7.1.1 User Stories and Acceptance Criteria

A foundational component of the Agile methodology is the use of *user stories* – natural language descriptions of desired functionality from the first-person perspective of users – to guide development and ensure that engineering effort is focused on directly implementing and servicing the requirements of end-users.

We will take these user stories into consideration when designing our software architecture, by ensuring that our design is as parsimonious as possible while still meeting the acceptance criteria for every following story point.

1. BASIC AUDIO MIXER CAPABILITIES

- (a) **As a user**, I want to connect up to 8 analog input channels to the mixer so that I can route my instruments through the system.
 - **Requirement:** Support for 8 analog inputs with ADC.
 - **Acceptance Criteria:** The system can recognize and process up to 8 analog input channels without signal degradation.
- (b) **As a user**, I want to be able to control volume and panning for each input channel so that I can balance the audio mix.
 - **Requirement:** Implement volume and panning control for all input channels.
 - **Acceptance Criteria:** Each input channel has adjustable volume and panning via the touchscreen interface.
- (c) **As a user**, I want the mixer to sum audio from multiple channels into a stereo output, so that I can send a combined signal to the speakers.
 - **Requirement:** Implement summing functionality for stereo output.
 - **Acceptance Criteria:** The mixer sums multiple input channels and outputs a stereo signal to line out and headphone out.

2. CREATIVE AUDIO EFFECTS

- (a) **As a user**, I want to apply creative audio effects (reverb, delay, compression, EQ, chorus/flanger/phaser) to my audio inputs so that I can enhance the performance.
 - **Requirement:** Implement DSP effects: reverb, delay, compression, EQ, chorus/flanger/phaser.

- **Acceptance Criteria:** Each input can have multiple effects applied, with adjustable parameters, processed in real-time.
- (b) **As a user**, I want to adjust effect parameters through the touchscreen interface so that I can fine-tune the sound during performances.
- **Requirement:** Intuitive effect parameter control via touchscreen.
 - **Acceptance Criteria:** Each effect's parameters can be modified through the touchscreen, with visual feedback.
3. **INTUITIVE TOUCHSCREEN INTERFACE**
- (a) **As a user**, I want to use an intuitive touchscreen interface to control the mixer so that I can easily navigate between effects and input routing.
- **Requirement:** Develop a digital touchscreen interface with user-friendly navigation.
 - **Acceptance Criteria:** The touchscreen interface is responsive, easy to navigate, and allows quick access to mixer functions such as routing and effect control.
4. **WIRELESS CONTROL VIA MOBILE APP**
- (a) **As a user**, I want to control the mixer wirelessly from my mobile device so that I can adjust settings remotely during performances.
- **Requirement:** Implement wireless control via a mobile app (WiFi/Bluetooth).
 - **Acceptance Criteria:** Users can control key mixer functions (volume, effects, routing) via a mobile app on both Android and iOS platforms.
5. **AUTO-MIX FEATURE**
- (a) **As a user**, I want the system to automatically adjust the volume levels based on the input dynamics so that the mix is balanced without manual intervention.
- **Requirement:** Implement automatic volume balancing based on input levels.
 - **Acceptance Criteria:** The auto-mix algorithm adjusts volume in real-time, maintaining a balanced mix.
- (b) **As a user**, I want the system to automatically apply effects based on the genre and input type (e.g., vocals, guitar) so that I can quickly set up the mixer.
- **Requirement:** Implement automatic effects selection based on genre and input type.
 - **Acceptance Criteria:** The system selects and applies appropriate effects without manual configuration based on predefined genre profiles.
- (c) **As a user**, I want the system to automatically apply EQ to unmask instruments in the mix, so that all instruments can be heard clearly.
- **Requirement:** Implement automatic EQ adjustments to reduce masking between instruments.
 - **Acceptance Criteria:** The auto-mix feature applies EQ automatically, ensuring all instruments remain distinct in the mix.
6. **USB COMPUTER CONNECTIVITY**

- (a) **As a user**, I want to connect the mixer to my computer via USB so that I can record and transfer audio data directly.
- **Requirement:** Implement USB connectivity for recording and data transfer.
 - **Acceptance Criteria:** The mixer can be connected to a computer, allowing audio to be recorded and data to be transferred via USB.

7.1.2 United Modeling Language (UML) Diagrams

Throughout our software documentation, we will utilize diagrams from the United Modeling Language (UML), a technical standard for the visual modeling of the design, implementation, and behavior of software systems. Due to the clarity and conciseness offered by UML diagrams in the communication of the design choices we make regarding our architecture, as well as their ubiquity in industry, we will prefer to use UML standards when producing visual models of our software design. We will provide expository detail about diagram types as they are used.

7.1.3 Use-Case Diagram

The use-case UML diagram depicts a high-level overview of the possible interactions that users can have with a system. It is a type of *behavioral* diagram, which are a class of diagrams that display the dynamic behavior of systems (as opposed to *structural* diagrams, which display the static structure of systems). The use-case diagram represents the overall functional requirements of the system, displaying the system's inputs produced by users and outputs consumed by users, as well as the high-level actions that the system must take in response to inputs.

7.1.4 User Roles

As shown in figure 7.1, there are four distinct roles that a user can have in relation to the mixer, as determined by their primary form of interaction with it:

Performers generate input signals, most of the time using instruments or microphones, to be converted to digital signals, processed, and mixed.

Audio engineers determine what digital processing to perform on the input signals, in order to achieve the best overall mix, both on a technical and aesthetic level. It is this action that is often simply referred to as "mixing".

Audience members consume the output signal, being the summation of the processed input signals, and converted back to analog to be fed to a transducer, such as a loudspeaker or PA system.

External computers utilize the mixer as a USB audio interface, treating the mixer's inputs and outputs as an extension of the computer's audio recording and routing capabilities.

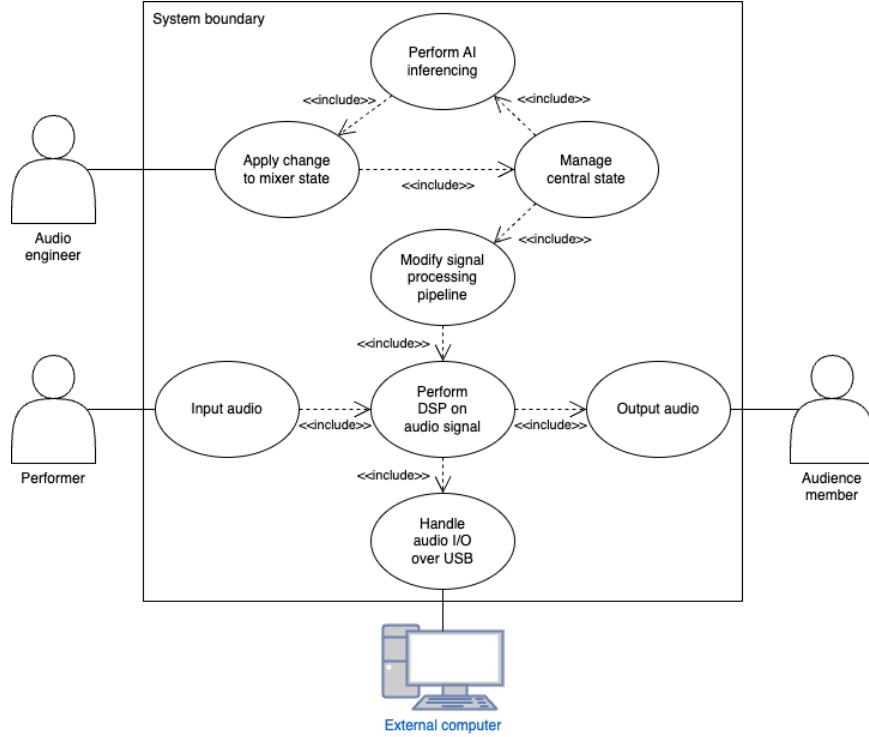


Figure 7.1: Use-case diagram for mixer software

7.1.5 Use-Cases Correspond to Software Modules

An essential task in designing complex software systems is to decompose the system into self-contained modules, connected by simple, stable interfaces. These modules should be constructed such so that they encapsulate groups of closely linked data and actions, and have clearly defined purposes.

The individual use-cases in the diagram are natural-language definitions of major actions that the system should take in response to actions performed by users. This lends the use-case diagram as being a convenient blueprint for determining the scope and interactions of formal modules, as will be expanded upon in the upcoming section 7.2.

7.2 Development of Software Architecture

In this section, we aim to provide readers with a detailed overview of the considerations that went into designing software architecture, building up to it from first principles. We will not only describe the "what" – the actual structure and behavior of our system, but also the "why" – details regarding design decisions, alternatives considered, and trade-offs involved in our chosen approach. As for the "how" – implementation and development details, we discuss this the next section 7.3.

Given the diverse set of functionality that we intend on implementing, a certain degree of complexity in our design is inevitable. Complexity in any software project

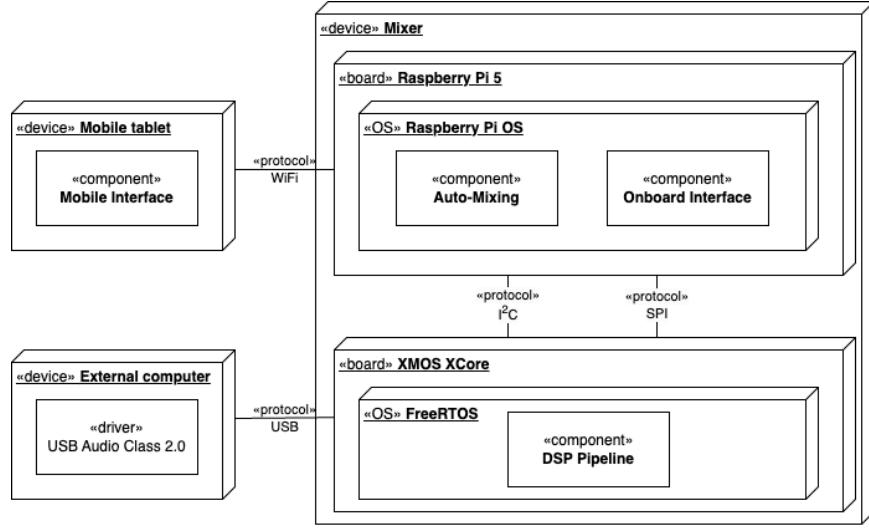


Figure 7.2: Deployment diagram of computing hardware

is undesirable, but it is especially hazardous in our project, with the extent of the work that is required to achieve our minimal viable product goals; unmanaged complexity may result in failure to deliver a functional product all together. This makes good software engineering practices all the more important; special attention will be directed toward the decision of how to **modularize** our system.

7.2.1 System Topology and Deployment

First, we must address a major constraint on our software design – the physical topology of the system’s computing hardware. The organization of our software must conform to the nature of the hardware it runs on. We will display the major hardware components, as well as their primary software responsibilities with a UML deployment diagram.

Onboard SoCs

To reiterate, the onboard computing hardware on our mixer is comprised of two system-on-a-chips: the **Raspberry Pi 5** and the **XMOS xcōre**. They are connected to each other via the **I²C** and **SPI** communication protocols, the purpose of each will be discussed later in section 7.3.3. As for the SoCs, their primary responsibilities are as follows:

Raspberry Pi 5 General purpose computer, to be used for all non-DSP tasks – AI inferencing, onboard touchscreen UI, and connection with mobile device.

XMOS xcōre Specialized crossover platform for handling real-time DSP, including effects processing and handling audio I/O over USB.

The relationship between the Raspberry Pi and the xCore in our system is structured as a **manager-worker paradigm**. The Raspberry Pi, acting as the manager,

is responsible for high-level control, coordination, and state management. It assigns audio processing workloads to the xCore, which operates as a dedicated worker, executing audio processing tasks more or less autonomously based on the commands it receives.

The impetus for opting for this software design pattern lies in the distinct strengths of the two hardware platforms. The Raspberry Pi, as a general-purpose single-board computer, is well-suited for tasks requiring versatility, such as managing user interfaces, networking, and system coordination. In contrast, the xCore is specifically designed for DSP tasks, offering deterministic execution and low-latency performance, critical for real-time audio applications.

Our architecture leverages the specialization of each of the boards. The xCore's deterministic guarantees is central to achieving real-time safety; all code that we execute on the xCore will be designed to be strictly deterministic, avoiding non-deterministic behavior that could compromise audio processing. Likewise, any code or logic that could potentially introduce non-determinism, such as user interactions, network operations, or dynamic state changes, is confined to the Raspberry Pi.

External Devices

The strict boundary of our mixer's hardware is indicated by the *Mixer* device node in figure 7.2. However, the functionality of our system extends to two external connected computing devices:

Mobile tablet Runs the mobile remote-control user interface, connected to the mixer over a **WiFi** connection.

External computer Utilizes the mixer as an audio interface over **USB**, as explained in section 2.2.1. All major operating systems provide built-in drivers for USB audio interface support; these generic drivers are called *class drivers*, and should enable plug-and-play functionality for our mixer's interface functionality.

These devices are technically ancillary to the central functionality of the mixer; the essential mixing functions can be operated without the use of the mobile UI, nor does the connection of an external computer in any way impact the internal flow of audio within the mixer. Nonetheless, the design and integration of these external devices require careful consideration to ensure they operate with their own encapsulated environments, and adhere our defined interfaces to avoid unintended dependencies and maintain the integrity of our system's architecture.

7.2.2 Managing Data Flows

The primary engineering challenge that will inform the design of our architecture is the management of the flow of data between various locations in the system. Especially considering the complex topology of our system, with data consumers and

producers scattered between disparate locations in hardware, a solid understanding of the connections that each module requires is crucial to designing a robust architecture.

Data Needs of System-Level Modules

We will begin by taking into consideration the data needs of the four preliminary software components from the deployment diagram (figure 7.2). The way these components were obtained was by grouping related user stories (from section 7.1.1) under umbrella terms, for the purpose of indicating the primary responsibilities of each hardware component in our system.

The components obtained by this loose categorization of desired functionality lends itself as a blueprint for formal modularization. These components, as they are, make for effective system-level subdivisions; they encapsulate tightly linked behavior and state, and are loosely coupled with each other. Here, we will describe what coupling is entailed between these components:

User interfaces Produces commands in response to user actions, such as adjusting mixer settings. The effect of these commands should be to mutate the current state of the mixer. Since both interfaces (mobile and onboard) must remain synchronized to each other, they must also consume this state information to update their local state. In addition, audio data must be consumed by the interfaces for the visual monitoring of audio levels and frequency spectra.

AI auto-mixer Consumes audio data in the form of spectrograms; this is the input data to the model, which then produces parameters to be applied to the mixer's state.

DSP pipeline Consumes state commands, which direct it to modify the processing pipeline in some fashion. Produces audio data, which is extracted from the processing stream and forwarded to other modules that require real-time audio data.

Types of Data

As evidenced by the data needs of our modules, there are two primary types of data that are used for inter-module communication in our system: **state data** and **audio data**. These distinct types of data are diametrically opposite of each other in nature, and require different handling and mechanisms to ensure correctness in communication.

State data refers to *discrete*, trigger-based changes in system state, such as adjusting a fader, enabling/disabling effects, or changing a parameter like panning. These events are instantaneous and represent individual updates to the system's state.

State data must be synchronized between subscribing (consuming) modules, including the DSP, AI auto-mixer, and both user interfaces, via a central mechanism. Each module will react to state changes as needed. For example, if the onboard interface triggers a volume change, the DSP will apply this change to the audio, while the mobile interface will update its UI to reflect the new state. We discuss the centralized mechanism for this in the upcoming section 7.2.2.

Audio data refers to the *continuous* stream of audio samples that flows between the DSP, the AI auto-mixer, and the visual monitoring components of the system. In contrast to the discrete nature of event-based data, audio data is continuous and must be handled in real-time for effective audio processing and visualization.

Beyond actual mixing in and of itself, audio data is central to the auxiliary functions of the mixer. The AI auto-mixer requires it to generate inferences, and the user interfaces require it for visual monitoring (e.g., displaying audio levels in the meters of channel strips, frequency spectrum visualizer for equalizer effect, etc.). Without mechanisms to ensure reliable and continuous audio data streaming from the DSP module into the modules on the Pi, these functionalities would be compromised.

Centralized Manager Modules

Due to the need for synchronization in our state data as well as real-time distribution of our audio data, we will employ centralized managers to be intermediate hubs for inter-module communication. These managers will run on the Raspberry Pi 5 and be responsible for the routing, management, and transformation of data between system modules.

State manager Handles all event-based, discrete state data. It will follow a **publisher-subscriber design pattern** to ensure that all relevant modules are synchronized in real-time. When a state change occurs (e.g., volume adjustment, effect toggling), the State Manager will publish this update, and all subscribing modules—such as the DSP, AI auto-mixer, and both user interfaces—will receive and apply the change accordingly. The State Manager ensures consistency across all modules, maintaining a single source of truth for mixer parameters, thus avoiding conflicts or discrepancies in system state.

Audio manager Handles the continuous flow of audio data within the system. Its primary role is to manage the streaming audio data coming from the DSP and ensure it is routed appropriately. It will forward high-fidelity audio streams to the **AI auto-mixer** for real-time inferencing and processing. Simultaneously, it will downsample the audio for the **user interfaces** (both onboard and mobile) to provide real-time visual feedback, such as audio levels and frequency spectra. The Audio Manager ensures efficient handling of high-throughput

audio data while maintaining the low-latency requirements necessary for real-time processing and visualization.

By separating state management and audio data handling, and centralizing the intermediary logic between the various processing and interface modules, we create a modular system architecture that allows for easier testing and integration, as well as facilitating distribution of development work between team members.

7.2.3 Complete Description of Architecture

At this point, we have now provided enough exposition to fully explain our software architecture block diagram shown earlier in Chapter 2, providing an overview of our computing hardware components and their physical connections, and our primary software modules plus their data dependencies. With the overview of our architecture established, we will move from discussing high-level design, onto concrete implementation details.

7.3 Concrete Design and Implementation Details

We will now discuss the concrete implementation of our abstract, high-level software architecture. We begin by examining the state and audio manager modules, the linchpins of our system. Their roles as mediators and coordinators between the remainder of our software components make them a good starting point for exploring both the internal implementation of individual modules and the design of their interfaces.

7.3.1 State Manager Process

The state manager module will be implemented as a Rust background service process running on the Raspberry Pi. It will be initiated during the Pi's boot sequence and operate indefinitely while the mixer is powered on. At its core, the state manager acts as a WebSockets server, employing an event-based runtime model powered by the `tokio` crate. Its primary function is to respond to incoming state commands delivered by publishers (such as the UIs or the AI model), reconcile these commands into a single consistent state, and subsequently broadcast the updated state to all relevant subscribers.

Publisher-Subscriber Pattern

As described earlier in section 7.2.2, the state manager follows the publisher-subscriber pattern to synchronize the mixer state between our various software modules. The Onboard UI, Remote UI, and AI Auto-mixing modules are **publishers**, meaning they can mutate the central state by sending (publishing) state-modifying commands to the state manager. The state manager reconciles these

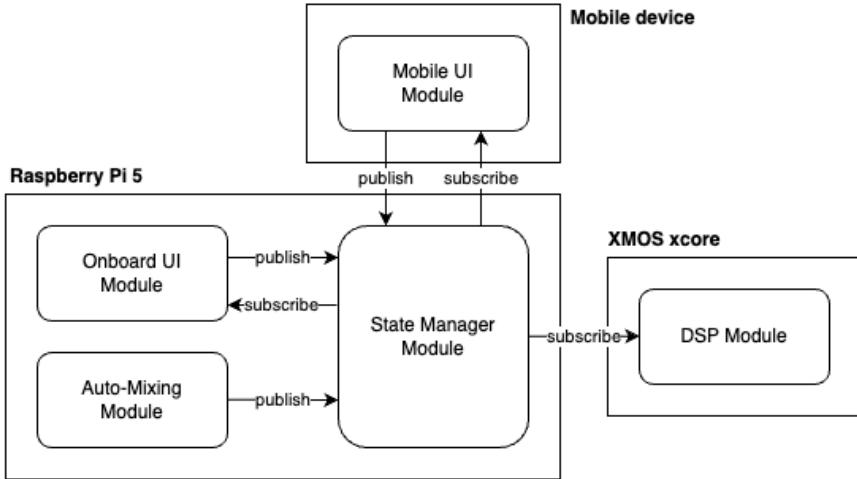


Figure 7.3: Modules Linked Via Publisher-Subscriber Pattern

commands into a single, consistent central state in a transactional and atomic manner. Once updated, the central state is forwarded to the DSP module and broadcasted back to the UI modules; these modules are the **subscribers**.

The purpose of using the publisher-subscriber design pattern is to *centralize* state management. The state manager now acts as the definitive arbiter of the mixer's state, isolating changes made by publishers and ensuring that state updates are propagated to all subscribers, effectively synchronizing them to the single source of truth dictated by the state manager.

Event-Based Runtime Loop

The state manager's runtime environment is powered by `tokio`, a Rust library that provides an asynchronous event-driven runtime. This framework enables concurrency without the overhead typically associated with traditional multi-threading, which is critical given the need for efficient and non-blocking communication with multiple modules.

`tokio` is particularly well-suited for the state manager due to its ability to handle asynchronous I/O operations, such as listening for WebSocket connections and managing I²C/SPI communication. It provides high-level extensions like `tokio-tungstenite` for WebSocket support and `serde` for data serialization, simplifying the implementation of the state manager's runtime logic.

WebSockets Communication

The state manager operates as a WebSocket server, presenting a unified and consistent interface to the state-mutating modules, which act as WebSocket clients; this means that the onus is on these modules (the UIs and the AI modules) to establish their own connection to the manager process, which otherwise has no way of knowing of the existence of the publisher modules.

The choice of using WebSockets to connect each publisher module to the state manager, even the ones that run alongside the manager process on the Pi, was done in pursuit of the simplicity afforded by being able to maintain a single module-agnostic interface, that can be used directly by publishers without requiring helper code to establish the connection with the manager process.

Communication With DSP Module Over I²C

Although the DSP module is conceptualized as being a subscriber module, the implementation of its connection with the state manager process differs from the remainder of the state-handling modules, in that communication occurs over I²C, not WebSockets.

Nonetheless, we ensure that the I²C communication is tightly integrated with the operation of the publisher-subscriber model, leveraging `tokio` for non-blocking I/O handling. This ensures that the state manager can handle transmitting high-frequency state updates without impacting the responsiveness of the WebSockets listeners.

7.3.2 State Command Standard Definition

In light of the need for a platform-agnostic method to communicate mixer state changes between various software modules, we have established a standardized, hierarchical JSON-based command scheme. This internal standard ensures that state data is handled by software modules created by different team members, and as it flows to and from each layer of our software architecture.

As discussed in section 3.3.8, the design of our internal standard was informed by existing audio plugin standards, which allow for arbitrary host automation of controls. Similarly, our JSON-based protocol abstracts the specifics of user interaction, enabling any of the three distinct user interfaces—onboard touchscreen, mobile app, or automated AI module—to effect changes in the mixer state in a unified and predictable manner.

Our software will *enforce invariants* for parameters with pre-defined ranges. Mechanisms will exist on each intermediary software layer that handles state data to validate and enforce constraints on parameters, ensuring that commands with invalid or out-of-bounds values do not propagate beyond their point of origin.

To simplify operation and development, our mixer architecture maintains a *fixed set of effects for each channel*, as opposed to allowing for the dynamic addition and removal of effect inserts. The user can toggle effects on or off and reorder them as needed, but the set of available effects remains constant.

Channel Parameters

These parameters define the attributes and controls associated with each channel in the mixer. These parameters govern fundamental aspects of channel behavior,

such as gain, panning, and muting, and they form the top-level controls available for configuring the audio mix.

Parameter	Type	Range	Description
input_gain	float	-48.0 to 24.0	Controls the input gain for the channel.
output_gain	float	-48.0 to 24.0	Controls the output gain for the channel.
panning	float	-1.0 to +1.0	Adjusts the stereo balance of the channel.
mute	boolean	True/False	Mutes the channel when true.
solo	boolean	True/False	Solos the channel when true, muting others.
order	array of string	N/A	Specifies the order of effects for the channel.
effects	object	N/A	Contains effect-specific parameters, as will follow.

Table 7.1: Channel Parameters

Effect Parameters

These parameters are specific to the various effects, associated with each channel. Each effect has its own set of adjustable properties, which determine its behavior and processing characteristics.

Equalizer

Parameter	Type	Range	Description
enabled	boolean	True/False	Whether the equalizer is active.
low_shelf_gain	float	-20.0 to 20.0	Low shelf gain.
low_shelf_cutoff	float	20.0 to 1000.0	Low shelf cutoff frequency.
band0_gain	float	-20.0 to 20.0	Band 0 gain.
band0_cutoff	float	60.0 to 1000.0	Band 0 cutoff frequency.
band0_q	float	0.1 to 6.0	Band 0 Q factor.
band1_gain	float	-20.0 to 20.0	Band 1 gain.

Parameter	Type	Range	Description
band1_cutoff	float	1000.0 to 7000.0	Band 1 cutoff frequency.
band1_q	float	0.1 to 6.0	Band 1 Q factor.
high_shelf_gain	float	-20.0 to 20.0	High shelf gain.
high_shelf_cutoff	float	4000.0 to 21050.0	High shelf cutoff frequency.
high_shelf_q	float	0.1 to 6.0	High shelf Q factor.

Table 7.2: Equalizer Parameters

Compressor

Parameter	Type	Range	Description
enabled	boolean	True/False	Whether the compressor is active.
comp_threshold	float	-60.0 to 0.0	Threshold.
comp_ratio	float	1.0 to 20.0	Compression ratio.
comp_attack	float	5.0 to 100.0	Attack time in milliseconds.
comp_release	float	5.0 to 100.0	Release time in milliseconds.
comp_makeup_gain	float	0.0 to 12.0	Makeup gain.

Table 7.3: Compressor Parameters

Reverb

Parameter	Type	Range	Description
enabled	boolean	True/False	Whether the reverb is active.
reverb_time	float	0.1 to 10.0	Reverb time in seconds.

Table 7.4: Reverb Parameters

7.3.3 Audio Manager Process

Data transfer from xcore to Raspberry Pi

As discussed in 3.3.2, our choice of method for data visualization will be Fast Fourier Transform signals. Due to their popularity in industry, as well as their ease of use, they were the obvious choice as a means to visualize the audio information the xcore sends to the Raspberry Pi. This section provides a detailed and structured explanation of the process involved in transferring FFT signals from the xcore to the Raspberry Pi using SPI.

FFT Data Preparation The first step in this process involves the transformation of raw audio into FFT data on the xcore. This transformation is performed using the Fast Fourier Transform algorithm, allowing us to extract key frequency components and their magnitudes. Upon performing the FFT, the xcore generates complex numbers for each frequency bin, consisting of both real and imaginary components. These numbers represent the amplitude and phase of the signal at each frequency, essential for audio analysis and visualization. In our system, it is important to structure this data efficiently for transmission via the SPI interface. Each FFT bin will be represented as a 32 bit floating point. The reason for this is because floating-point values can represent a much wider range of amplitudes, making them ideal for audio signals with a high dynamic range, such as in our project. The complex numbers are stored in a buffer on the xcore in preparation for transmission.

Packetization of FFT Data Given the nature of SPI as a serial communication protocol, it is essential to packetize the FFT data before transmission. The goal here is to divide the large FFT data into smaller, manageable packets for efficient transfer. For our project, a 1024-point FFT will be performed. This results in 1024 real values and 1024 imaginary values, which will be split into smaller packets of 64 data points. Packetization ensures that the data can be transmitted in chunks, reducing the risk of data overflow and enhancing error handling. Each packet will consist of a header that identifies its sequence in the data stream, followed by the actual FFT data. The header will also include error-detection mechanisms to ensure the integrity of the data. This structure allows the Raspberry Pi to verify the correctness of the received data and request retransmission if necessary.

SPI Configuration The xcore acts as the SPI master, controlling the data transfer to the Raspberry Pi. On the xcore, the clock speed, mode configuration, and data word size must be adjusted appropriately. The SPI clock speed will be 8 MHz, as this is the required data rate of the Raspberry Pi as the SPI slave while considering a 1024-point FFT, sampling rate of 44.1 kHz, and 32-bit floating point values. The SPI mode will be configured to ensure that the clock polarity and phase are synchronized between the xcore and the Raspberry Pi by operating in mode 0. This mode came from the fact that the

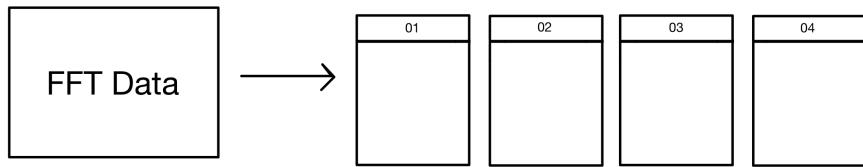


Figure 7.4: Packetization of FFT Signals

most commonly used configuration is SPI mode 0. In addition, mode 0 means the clock polarity is at 0 (clock idling at low) and clock phase 0 (data sampled at rising edge). For the data word size, we will be using 32-bit words. This word size was selected to optimize data transmission. Once the SPI interface is configured, the xcore sends out the FFT data packet-by-packet. Each bit is shifted out on the MOSI (Master Out Slave In) line, synchronized with the clock signal. The CS (Chip Select) line then is used to select the Raspberry Pi as the active recipient.

Processing and Visualization After the complete FFT frame is successfully received by the Raspberry Pi, it proceeds to process the data for visualization purposes. The FFT signals are used to generate a visual representation of the audio spectrum, which is then displayed on the on-board LCD of the mixer via UNIX domain socket. Furthermore, the processed FFT data is then transmitted to the mobile tablet via WebSockets for remote visualization. This process is detailed in ?? and 7.3.3, respectively.

Data transfer from Raspberry Pi to mobile tablet

Once the Fast Fourier Transform signals are successfully transferred from the xcore to the Raspberry Pi, the subsequent phase of the project involves transmitting these signals to a mobile application for real-time visualization and processing. The primary challenge here is to ensure efficient, low-latency communication between the Raspberry Pi and the mobile application while maintaining the integrity of the FFT data during the transfer. This section details the process for transferring FFT data from the Raspberry Pi to a mobile application, exploring the communication protocols involved and the steps required to facilitate real-time interaction with the audio data.

FFT Data Preparation After receiving the FFT data from the xcore, the Raspberry Pi stores this data in memory. Each FFT frame consists of a series of complex numbers, representing the real and imaginary components of the frequency

spectrum of the audio signal. These FFT values, which reveal frequency and phase information, are crucial for visualizing the audio's spectral content in real time. The first step in preparing the FFT data for transmission is serialization. Serialization refers to converting the in-memory data structure (the FFT array) into a format that can be transmitted over a network. The Raspberry Pi can serialize the FFT data into commonly used formats. Although JSON would be the easiest to work with due to previous experience as well as human-readability, we will be using a binary format as FFT signals generate a lot of data. For these large datasets, binary formats are more efficient and reduce overhead during transmission, especially in real-time applications. (Protobuf or CBOR)

Establishing Communication In order to send the FFT data from the Raspberry Pi to the mobile application, a communication protocol must be utilized. For our project, we will be using WebSockets, as it allows for real-time data transfer, as opposed to traditional HTTP/REST APIs which are not real-time and must be used by periodically sending requests to the server. By using WebSockets, we simplify our data communication pipeline with one active connection that allows for full-duplex connection, as opposed to HTTP/REST APIs, where we would need to continuously establish connections to transfer data.

Processing of FFT Data on Mobile Tablet Once the FFT data has been transferred to the mobile application via WebSockets, the next task is to process and visualize the data. The mobile application will display the FFT data as a frequency spectrum, with the x-axis representing frequency bins and the y-axis representing the amplitude of each frequency. This provides the user with a real-time view of the frequency components of the audio signal. As a result of using WebSocket-based communication, the mobile app continuously updates the visualization as new FFT data is received, offering us a dynamic representation of the audio spectrum.

7.3.4 Automix Model Design, Training, & Deployment

Overview

The AI model is trained on a development machine using large datasets mentioned previously of audio signals and their corresponding mixing parameters with PyTorch as the framework. Once trained, the model can be exported to a more lightweight format that can be deployed on the pi. The Raspberry PI OS will have to be optimized for real-time performance by reducing background tasks and enabling GPU acceleration. As an application it will listen for audio signals or FFTs from the DSP pipeline. Then it will pre-process the signals into mel-spectrograms. The spectrograms are then run through the model and the mixing parameters are inferred, which are then sent back to the DSP pipeline.

Data Preprocessing

The primary goal is to ensure the dataset is consistent, normalized, and free from artifacts such as silent audio segments or insufficiently long tracks. This step is essential for optimizing the performance of the Differentiable Mixing Console (DMC), which relies on precise input data to infer accurate mixing parameters.

The first step in preprocessing involves discarding audio tracks that are shorter than the required input length. Multitrack datasets often contain mixes and stems of varying durations. Feeding audio files shorter than the prescribed length into the model can lead to inconsistencies, such as errors during batch generation or model training. To address this, a filtering mechanism is employed to exclude such examples. For efficiency, this is achieved using the `torchaudio.info` function, which retrieves metadata about each audio file—specifically the number of frames—without loading the entire file into memory. By iterating through the dataset and retaining only those files with a frame count exceeding the required length, we ensure compatibility and uniformity across training samples.

Once the dataset has been filtered for length, the next step is loudness normalization. This step scales the amplitude of each audio sample to a consistent level. Normalization is performed by dividing the audio waveform by its maximum absolute value, ensuring a peak amplitude of 1. A small constant (e.g., `1e-8`) is added as a lower bound to prevent division by zero, which can occur if the audio contains silence. Normalizing both individual stems and their corresponding mixes is crucial, as it reduces the impact of varying recording levels and allows the model to focus on the relative differences between tracks.

Handling silence in the dataset represents another important aspect of preprocessing. Silent audio chunks, often found at the beginning or end of tracks, provide no meaningful information and can introduce errors during training. Specifically, normalizing silence results in `NaN` values due to division by zero, which disrupts the training process. To avoid this, a mechanism is implemented to detect and skip silent chunks. The preprocessing pipeline generates a random starting offset for loading frames from the audio, ensuring variability in training data. Each chunk is evaluated for energy, calculated as the mean of the squared amplitude. If the energy falls below a predefined threshold (e.g., `1e-8`), the chunk is classified as silent, and a new offset is generated. This process repeats until a non-silent chunk is identified. Afterward, the audio is normalized using the same peak normalization technique described earlier.

DMC Design

The Differentiable Mixing Console (DMC) is designed to model the traditional signal processing chain found in professional mixing consoles while enabling end-to-end learning directly from audio waveforms. The architecture is comprised of three primary components: the controller network, the transformation network, and the loss function. Together, these components allow the DMC to process multitrack audio, predict human-readable mixing parameters, and generate perceptually coherent

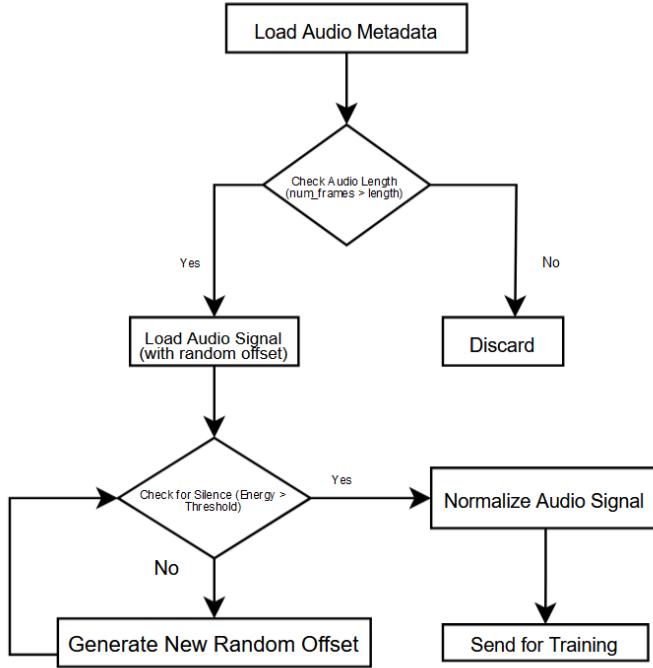


Figure 7.5: Preprocessing Pipeline

mixes.

Controller Network. The controller network is responsible for analyzing the input audio tracks and predicting mixing parameters for each channel. This network employs shared encoders and post-processors for all input tracks, ensuring permutation invariance. The encoder processes each track independently, extracting features relevant for mixing using a pre-trained spectrogram-based VGGish model. These features are fine-tuned during training to adapt to the mixing task. The post-processor predicts the mixing parameters for each track, such as gain, panning, compression, and equalization settings. To incorporate global context across all tracks, the post-processor also receives a context embedding, computed by averaging the embeddings of all input tracks. This allows the DMC to balance individual track processing with overall mix cohesion. See figure 7.6

Transformation Network The transformation network emulates the signal processing operations of the mixing console, including equalization, compression, and reverb. These effects are modeled using a Temporal Convolutional Network (TCN), which processes audio in the time domain. The TCN architecture is characterized by stacked dilated convolutions, enabling it to capture long-range dependencies in the audio signal. To condition the TCN on the predicted parameters, feature-wise linear modulation (FiLM) layers are incorporated. These layers inject parameter-specific information into the network, allowing it to adapt dynamically to the desired

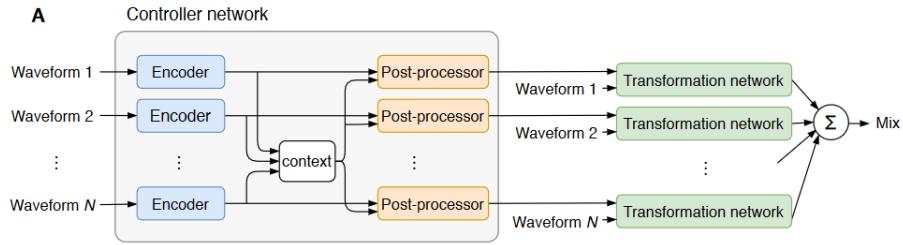


Figure 7.6: Controller Network - DMC

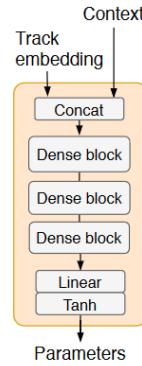


Figure 7.7: Post Processor - DMC

effect settings. The transformation network is trained independently before being integrated into the DMC, ensuring accurate emulation of complex signal processing chains. See figure 7.8.

Post Processor The post-processor in the Differentiable Mixing Console (DMC) serves as the final step in the controller network, converting the extracted track features and contextual information into actionable mixing parameters. This component is implemented as a lightweight, three-layer multi-layer perceptron (MLP) with PReLU activations and dropout to improve generalization during training. The post-processor takes two key inputs for each track: the track embedding, which encodes features specific to the track, and a context embedding, which represents the average of all track embeddings in the mix. This dual-input design allows the post-processor to balance individual track adjustments with the overall requirements of the mix. By incorporating global context, the post-processor can account for interactions between tracks, such as ensuring complementary frequency ranges or balancing relative loudness. The predicted parameters, such as gain, panning, and EQ settings, are human-readable, making them interpretable and adjustable in downstream applications. This design ensures that the post-processor contributes to both the precision and flexibility of the DMC's mixing capabilities. See figure 7.7.

Training Training the DMC involves two stages. First, the transformation network is trained to emulate audio effects using examples generated from digital signal processors. Once the transformation network is pre-trained, the entire DMC is trained end-to-end using the stereo loss function. The controller network learns to predict mixing parameters by minimizing the perceptual difference between the generated mix and the ground truth. The predicted parameters are interpretable and human-readable, allowing manual adjustments to the mix if needed. When a piece of music is input into the DMC, each track is passed through the encoder, which extracts its salient features. These features are combined with global context information to predict mixing parameters for each track. The parameters are passed to the transformation network, which applies the corresponding audio effects, such as equalization or compression. The processed tracks are then mixed together by summing the stereo outputs, producing a cohesive final mix. At every stage, the system leverages differentiable operations to enable gradient-based learning, ensuring that the DMC can refine its predictions during training.

Additional Properties The DMC is designed with several properties that make it robust and versatile:

- **Permutation Invariance:** The shared encoders and post-processors ensure that the order of input tracks does not affect the mix.
- **Human-Readable Parameters:** The mixing parameters predicted by the controller network are interpretable, enabling fine-tuning by audio engineers.
- **Scalability:** The architecture can handle varying numbers of input tracks, making it suitable for diverse mixing scenarios.
- **High-Fidelity Processing:** By leveraging the TCN for audio effects, the DMC produces mixes with minimal artifacts, even at high sample rates.

In summary, the DMC combines traditional signal processing principles with modern deep learning techniques to create a flexible and powerful tool for automatic multitrack mixing. Its architecture, training approach, and nuanced design elements make it uniquely suited to learn mixing conventions directly from raw audio data.

Implementing Non-Differentiable Effects

To extend the Differentiable Mixing Console (DMC) with additional audio effects beyond basic gain and panning, we integrate non-differentiable effects such as compression and parametric equalization directly using signal processing functions from the `dasp.pytorch` library. Unlike the original implementation, as shown in Figure 7.8 which relies on a neural network proxy (TCN) to emulate these effects, this approach allows the use of pre-built, accurate audio processing functions, avoiding the need for neural proxies.

The integration begins by incorporating specific processors such as `compressor` and `parametric_eq`. Each effect is parameterized to accept a range of values, nor-

Transformation Network

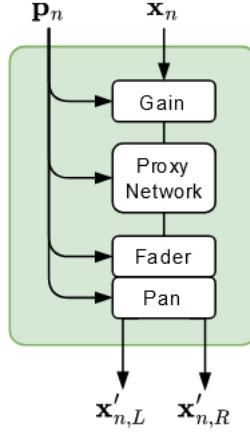


Figure 7.8: Transformation Network - DMC

malized between 0 and 1 during training. These normalized values are denormalized at runtime to their actual parameter ranges (e.g., threshold, ratio, and gain for compression or gain and cutoff frequencies for equalization). The parameter transformation ensures compatibility with the differentiable training pipeline while maintaining accurate and realistic audio processing.

The processing pipeline executes in two stages:

- Parameter Mapping:** The model predicts normalized parameters for each effect. These are then denormalized using predefined parameter ranges. For instance, normalized threshold values for a compressor are mapped to dB ranges suitable for audio processing.
- Audio Processing:** The denormalized parameters are passed to the corresponding signal processing function (e.g., compressor or parametric_eq). These functions operate directly on the input audio tensors, transforming them into their processed counterparts while applying the desired effects.

Compression. Compression is applied to reduce the dynamic range of the input audio signal, ensuring consistency and balance across the mix. It is composed of parameters mentioned earlier. These parameters are predicted by the model as normalized values in the range [0, 1]. During runtime, they are denormalized to their real-world ranges (e.g., thresholds in the range -60 dB to 0 dB). The compressor function applies these parameters to process the audio signal, attenuating loud sections while maintaining quieter parts, ensuring that the dynamic range aligns with the desired output characteristics.

Parametric Equalization. Parametric equalization allows frequency-specific manipulation of the audio signal to emphasize or attenuate specific bands. The

`parametric_eq` function in `dasp.pytorch` supports multi-band processing with parameters also mentioned earlier.

As with compression, parameters for equalization are predicted as normalized values and denormalized during runtime. For instance, cutoff frequencies may be mapped from a normalized range to physical frequencies such as 20 Hz to 20 kHz. The `parametric_eq` function then processes the audio by applying gain adjustments to the specified frequency bands, shaping the tonal balance of the mix.

Processing Pipeline. The integration of compression and parametric equalization occurs as part of the audio processing pipeline:

1. The model predicts normalized parameters for each effect based on the input audio.
2. These parameters are denormalized to their real-world ranges (e.g., attack time for compression or cutoff frequency for equalization).
3. The processed audio is passed through the `compressor` and `parametric_eq` functions sequentially, transforming the signal.

During the forward pass, the mixer processes each track in the batch by applying the operations in the following order: gain, equalization, compression and then panning. The predicted parameters for all effects are concatenated into the parameter tensor, enabling a comprehensive and interpretable representation of the mix settings. Each effect transformation contributes to the final output waveform, which is then summed across tracks to produce the stereo mix.

This integration enables the DMC to apply nuanced audio effects with high fidelity, directly leveraging `dasp.pytorch`'s pre-built functionality. By incorporating compression and equalization, the model can effectively shape dynamics and tone while maintaining differentiability for end-to-end training.

Loss Functions

The Differentiable Mixing Console (DMC) utilizes a variety of loss functions to optimize the model's performance and ensure that the resulting audio mixes are perceptually coherent and technically accurate. These losses capture different aspects of audio quality, ranging from waveform fidelity to spectral accuracy, and are carefully weighted to achieve a balanced training process.

Multi-Resolution STFT Loss (MR-STFT). The Multi-Resolution Short-Time Fourier Transform (MR-STFT) loss is a key component in training the DMC. This loss compares the spectral content of the predicted mix with the target mix at multiple time and frequency resolutions. By using various FFT sizes, hop sizes, and window lengths, the MR-STFT loss ensures that both fine-grained details and broad spectral features are accurately reproduced. The MR-STFT loss includes terms for

linear and logarithmic magnitude spectrums, each weighted to emphasize perceptual relevance. This loss is particularly effective for preserving the tonal balance and temporal structure of audio signals.

L1 Loss. The L1 loss measures the absolute difference between the predicted and target waveforms. It promotes smooth and continuous outputs by penalizing large deviations equally across all signal values. This loss is simple yet effective for waveform reconstruction tasks and provides a robust baseline for optimizing signal fidelity.

SI-SDR Loss. The Scale-Invariant Signal-to-Distortion Ratio (SI-SDR) loss evaluates the quality of the predicted mix by quantifying the distortion, interference, and noise relative to the target mix. This loss is computed directly in the time domain, making it sensitive to phase and amplitude errors in the waveform. SI-SDR loss ensures that the DMC produces clean and artifact-free mixes while maintaining the relative levels of input tracks.

Sum-and-Difference STFT Loss (SD-STFT). The Sum-and-Difference STFT loss is designed for stereo audio mixing, where maintaining spatial balance is critical. This loss operates on the sum and difference signals of the left and right stereo channels, ensuring that the model learns to preserve the intended stereo image. By incorporating STFT-based spectral comparisons, the SD-STFT loss balances left-right spatial coherence with spectral fidelity.

Loss Weighting and Combination. During training, the individual losses are combined with specific weights to form the total loss. These weights are hyperparameters that control the relative importance of each loss function. For instance, MR-STFT loss may be weighted more heavily to emphasize spectral accuracy, while L1 loss ensures waveform smoothness. The loss weights are adjusted empirically to achieve the desired trade-off between perceptual quality and technical accuracy.

Hyperparameters for Loss Computation. Several hyperparameters govern the behavior of the loss functions:

- **FFT Sizes, Hop Sizes, and Window Lengths:** For MR-STFT and SD-STFT losses, these parameters determine the resolution at which the spectral comparisons are performed.
- **Learning Rate Scheduler:** A cosine annealing or multi-step scheduler is used to adapt the learning rate over training epochs, ensuring stable convergence.

The batch size, which defines the number of audio tracks processed simultaneously, is set to **32** to balance computational efficiency with stable gradient estimation. A larger batch size ensures smoother updates to the model’s parameters but

requires sufficient memory, particularly when processing multi-channel audio data. For 8 input channels, this configuration provides a manageable trade-off between computational demands and training stability.

The learning rate is set to 3×10^{-4} , reflecting a standard choice for Adam optimizers in audio-related deep learning tasks. This value ensures steady progress toward convergence while avoiding the risk of overshooting the loss minimum. Given the absence of neural proxies and the direct use of dasp-pytorch effects, this learning rate provides sufficient granularity to fine-tune the mixing parameters.

The parameter ranges for audio effects are also key hyperparameters. Since effects like compression and parametric equalization are directly applied using dasp-pytorch, the parameter ranges dictate how the model interacts with these functions. For compression, the threshold is set between -60 dB and 0 dB, the ratio spans from 1:1 to 20:1, and attack/release times range from 5 ms to 100 ms. Similarly, for parametric EQ, gain adjustments range from -20 dB to +20 dB, cutoff frequencies are mapped between 20 Hz and 20 kHz, and Q-factors span from 0.1 to 6. These ranges ensure compatibility with professional audio standards while enabling the model to explore a diverse range of effect configurations.

The learning rate schedule is another important consideration. A cosine annealing scheduler with a maximum epoch count is employed to gradually reduce the learning rate, promoting fine-tuning of the model during later stages of training. This schedule is particularly useful in stabilizing the optimization process when working with high-dimensional audio data.

Finally, the maximum number of tracks is set to 8 for the DMC when handling 8 input channels. This hyperparameter ensures that the shared encoder and post-processor layers are equipped to handle the specific channel configuration, maintaining permutation invariance and scalability. Together, these hyperparameters are fine-tuned to ensure the DMC is optimized for processing multitrack audio with high fidelity and interpretability.

During each training step, the DMC computes the predicted stereo mix and compares it to the ground truth mix using the defined loss functions. The gradients are backpropagated through the model, enabling parameter updates that optimize both the spectral and time-domain characteristics of the mixes.

Model Deployment and Hardware Acceleration on Raspberry Pi

Deploying the Differentiable Mixing Console (DMC) model onto a Raspberry Pi involves selecting deployment methods and leveraging hardware acceleration to ensure real-time performance. The Raspberry Pi's constrained resources necessitate careful optimization, especially when processing multichannel audio data. Below, we evaluate several deployment methods and hardware acceleration techniques, emphasizing their compatibility with the DMC and their trade-offs in terms of efficiency, ease of implementation, and scalability.

Deployment Methods. The DMC model can be deployed on the Raspberry Pi using various approaches, including PyTorch with ONNX Runtime, TensorFlow Lite

(TFLite), and native PyTorch execution. Each method has specific benefits and limitations:

- **PyTorch with ONNX Runtime:** This method converts the model to the ONNX format, optimizing it for efficient execution on edge devices. ONNX Runtime supports hardware acceleration through libraries like OpenVINO or TensorRT, making it a versatile choice.
- **TensorFlow Lite (TFLite):** By converting the PyTorch model to ONNX and then to TFLite, the model can leverage TFLite's optimized interpreter, designed for ARM-based processors like the Raspberry Pi. TFLite also supports quantization, reducing model size and improving inference speed.
- **Native PyTorch Execution:** Running the DMC model directly in PyTorch is the most straightforward method but lacks the hardware optimizations offered by ONNX or TFLite.

Hardware Acceleration Techniques. To meet real-time processing requirements, hardware acceleration techniques such as GPU acceleration, the Raspberry Pi's NEON SIMD instructions, and external accelerators are essential:

- **NEON SIMD Instructions:** These are supported natively on the Raspberry Pi's ARM processor and provide fast parallel computation for matrix operations and signal processing tasks.
- **GPU Acceleration:** The Raspberry Pi supports OpenCL and Vulkan APIs for general-purpose GPU acceleration, enabling faster inference of computationally heavy models.
- **External Accelerators:** Devices like the Coral USB Accelerator (based on Google's Edge TPU) or NVIDIA Jetson Nano can offload processing, significantly boosting performance while reducing load on the Raspberry Pi.

Comparison of Deployment Methods and Hardware Acceleration. The following table compares the deployment methods and hardware acceleration options, highlighting their relative advantages and limitations.

We will select PyTorch with ONNX Runtime for deploying the DMC due to its balance of performance and flexibility. ONNX optimizations enable efficient execution on the Raspberry Pi, leveraging NEON SIMD instructions or external accelerators like TensorRT. This approach retains PyTorch's development versatility while ensuring real-time multichannel audio processing.

Method	Advantages	Key Features and Limitations
PyTorch with ONNX Runtime	Optimized for edge devices; supports hardware acceleration libraries like OpenVINO and TensorRT; versatile for various hardware configurations.	Requires additional conversion steps; performance gains depend on compatibility with hardware acceleration libraries.
TensorFlow Lite (TFLite)	Lightweight interpreter designed for ARM processors; supports quantization to reduce model size and latency; optimized for real-time applications.	Conversion from PyTorch to TFLite can introduce compatibility issues; limited support for complex PyTorch layers.
Native PyTorch Execution	Simplifies deployment by retaining PyTorch functionality; minimal overhead during conversion and integration.	Slower inference; does not leverage hardware acceleration effectively; less optimized for real-time performance.
GPU Acceleration	Enables faster inference for computationally heavy tasks; utilizes OpenCL or Vulkan APIs available on Raspberry Pi GPUs.	Limited GPU resources on Raspberry Pi; implementation complexity may require specialized knowledge.

Table 7.5: Comparison of Deployment Methods and Hardware Acceleration Techniques for the DMC Model

Chapter 8

System Fabrication/Prototype Construction

8.1 PCB Design Standardization

Due to guidance about power regulation ICs being very inconsistent, our group opted to split the overall circuitry into two distinct PCBs: one specifically for power

(using pin-headers) and another for the entire signal chain (input/output). Due to the structure of the Raspberry Pi, it can only be interfaced with jumper wires (as there are pin headers for every pin). These necessitates three separate zones of the overall enclosure, with the goal to maintain these separately. With these two circuits, we require a standard between them to mitigate any risk with the designs. The overall PCBs will be 4 layers with a stack-up of Signal-Ground-Ground-Signal (with the two ground layers being joined together using vias). The PCB itself will be 1.6 Oz using FR4 as the material. We have ultimately decided against using a separate analog and digital ground as a majority of the lines are relatively low speed and the analog section will be well separated from the digital section. This affords us a ground plane that should be able to dissipate a large amount of power in addition to offering less noise.

While component size constraints would normally be an issue when hand-soldering, we have opted to use part placement services (e.g. JLCPCB) to place components due to many of the ICs being used being extremely small. Due to this, we do not have constraints on the size of the components, and will utilize the optimal component size (0402, 0603, 0805) as per the datasheet of each specific IC. Generally, if the IC requires a high level of power dissipation, larger components will be used, but having larger capacitors increases parasitic inductance and will be mitigated for sensitive devices. As such, there cannot be a standard size of component to be used for this project.

As for the aforementioned pin headers, we will be utilizing 2.54 mm female sockets coupled with jumper wires to connect the three parts of the overall circuit. For sturdiness, we will also use male-to-male Dupont connectors to ensure the wires stay in place.

The two PCBs will be rectangular with mounting holes on each corner measuring 3mm in diameter with 5mm of spacing from the edges. Because both PCBs will be very different in scope, the overall sizes will differ. Adhering to the fabrication limitations from JLCPCB, the minimum trace width will be 0.3mm for signal lines and 0.5mm for power lines. Filled zones will additionally be used for areas that require higher power. The main via size to be used is 0.8mm in diameter (with 0.4mm in diameter for the hole itself) and pads should be 0.3mm apart at the very least. To ensure this was adhered to, these standards were input into KiCAD and used when executing a Design Rules Check (DRC).

8.2 Power Supply PCB

The power supply PCB takes in $\pm 12V$ and $5V$ inputs from the aforementioned commercial AC-DC transformer. The positive 12V will be stepped up to 36V to be able to supply phantom power and additionally serve as the positive rail of the many op amps used in the Mixed Signal circuitry (preamp, post-DAC LPF). The negative 12V supply will exclusively be used to power the negative rail of the op amps used in the circuitry. This transformer will be connected to our PCB using wire connectors that are capable of supplying relatively high current (<2 Amps),

most likely 18 or 20 gauge wires, ending at a terminal block for tight connection.

The main signal path is starting from the 5V DC input, which is then stepped down to $3.3V \rightarrow 1.8V \rightarrow 0.9V$. The 1.8V and 0.9V rails being exclusively used for the Xcore DSP chip. Cascading the power rails in series allows the LDOs to be more efficient and allow for proper startup of the Xcore chip. With higher efficiency of LDOs comes lower power loss, meaning less thermal management required.

Overall, this board should output the following voltages to be used in the main Mixed Signal circuit (and Raspberry Pi): -12V, 0.9V, 1.8V, 3.3V, 5V, 12V, and 36V. All of these power rails will be accessible via female pin headers as mentioned before. To do this, we will have at least two pin headers per supply (a column in a 2x10 layout). This will allow redundancy for the lower voltages as well as having a centralized location for the power. For supplies like 5V, due to many devices needing to utilize this voltage we will use at least 4 pins to allow this capability. The extra headers allow for us to utilize as many as necessary to supply all the pins on the external boards. There will also be output LEDs for each source of power (sans -12V) to assist in debugging the regulators.

As shown in Figure 8.1, the overall footprint of the Power Supply PCB is only 35 x 60 mm, which is quite compact to fit in the specified enclosure. Because four layers are being used for the PCB, the two ground planes need to be joined together using via stitching, hence the large amount of vias being used.

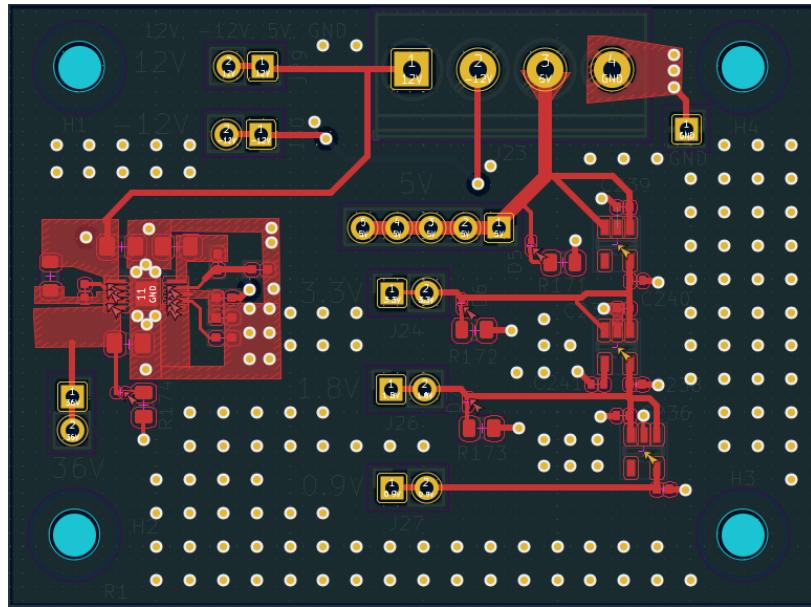


Figure 8.1: Schematic View of Power Circuitry PCB

Additionally shown is the 3D view of the Power Supply PCB in Figure 8.2.

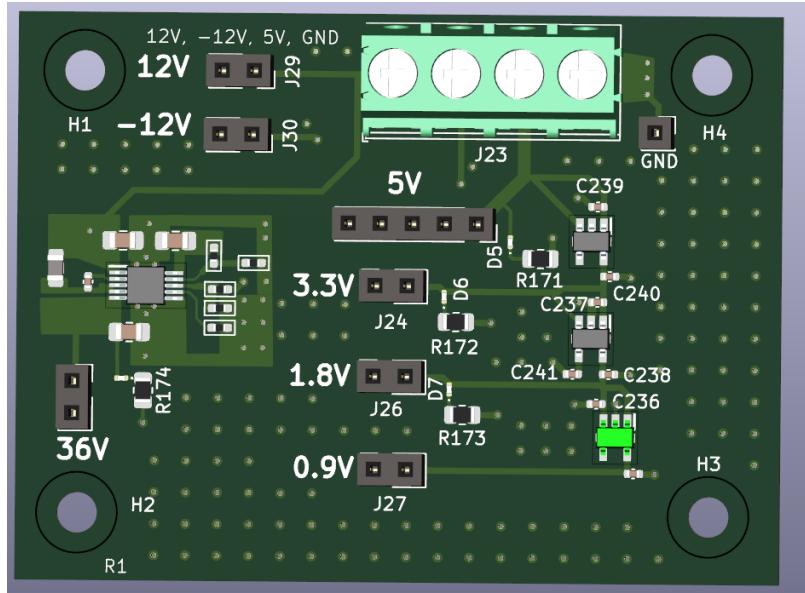


Figure 8.2: 3D View of Power Circuitry PCB

8.3 Mixed Signal PCB

Our project, due to being a mixed signal circuit, requires rather stringent adherence to standard practice to ensure signal quality is not degraded. This means that noisy digital signals should be separated from the analog signals, especially power rails. We are also specifically using the voltage output from the LDOs to power the analog supply rails as they are much less noisy from both switching and digital noise. The switching regulated voltages are to power the digital components. All of the input voltages will be obtained via connecting the pin headers on those from the Power Supply PCB to those on the Mixed Signal PCB.

The main components on this board following the signal path are Preamplifiers → ADC → Xcore Chip → DAC → Low Pass Filter → Headphone Amplifier & Line Output. The Xcore chip will be in constant communication with the Raspberry Pi for both input and output of data. Because of the fact that 1/4" connectors will be plugged into the board for the audio inputs, the overall dimensions of the board will be dictated by the feasible use of them. Pin headers will also be used to connect output pins from the xcore to those of the Raspberry Pi, as well as a set of pin headers exclusively for programming the xcore chip itself.

As shown in Figure 8.3, the main PCB integrates the seven separate voltages obtained from the Power PCB to power all the individual circuitry, with the xcore chip as the focal point of the PCB.

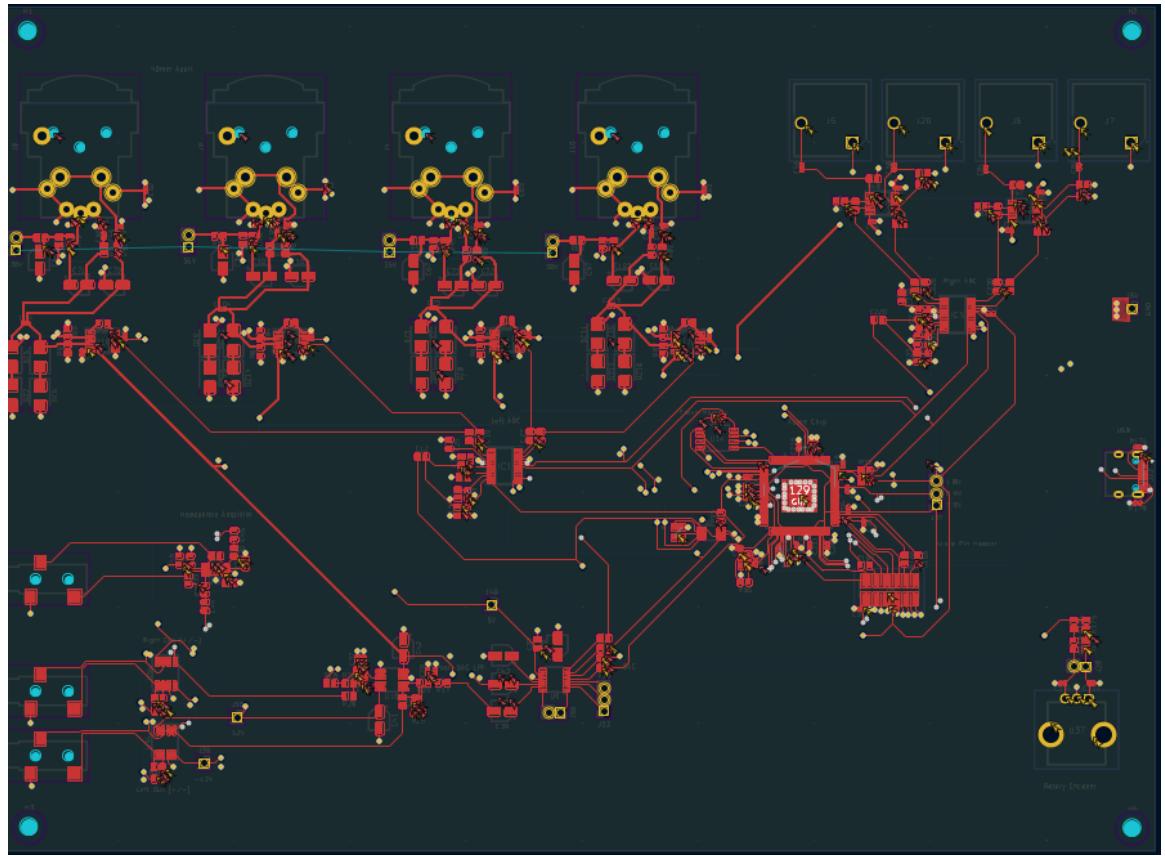


Figure 8.3: Schematic View of Main Audio PCB

Additionally shown is the 3D render of the PCB, with some parts being missing due to library limitations. The dimensions of the PCB come out to be 247.75 x 181.5 mm.

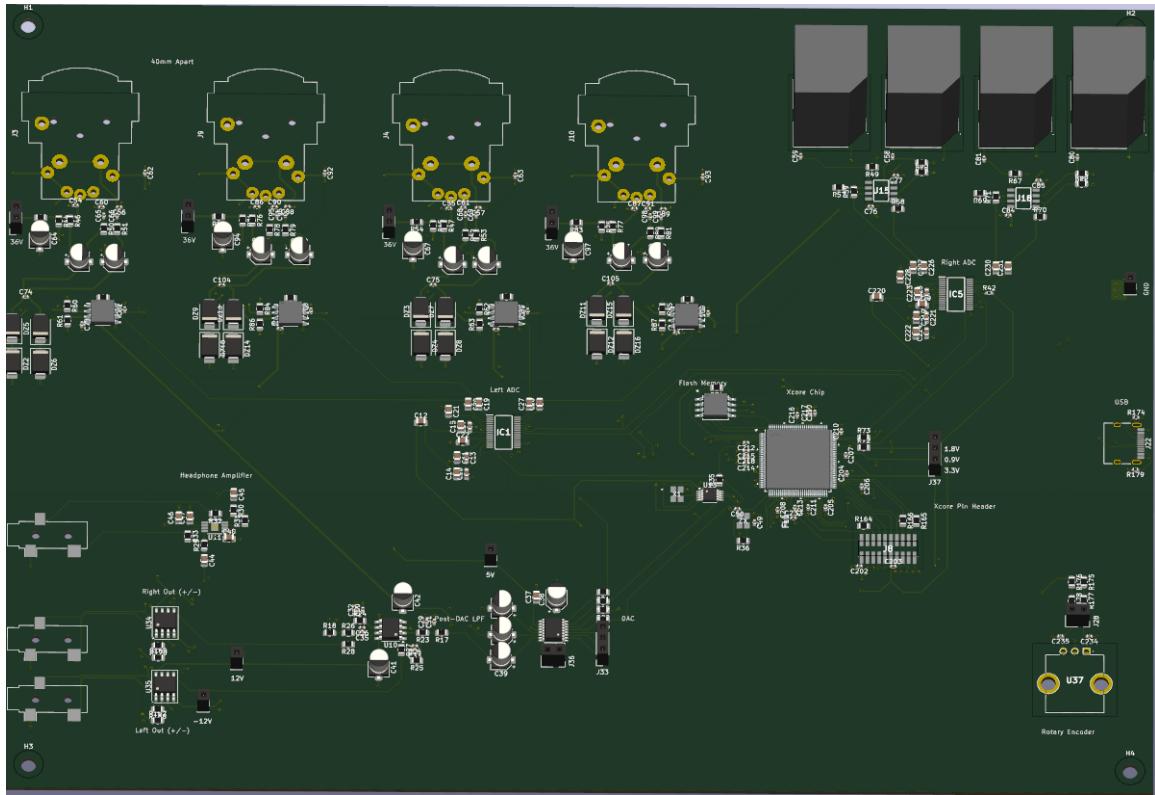


Figure 8.4: 3D view of Main Audio PCB

Chapter 9

System Testing

9.1 Hardware Testing

Hardware testing is an important component of the design process. This section will outline the test equipment and environment. The sections below will detail the testing results of individual components.

The hardware testing that has been conducted so far has been carried out in the UCF Senior Design Lab. Some of the equipment provided there includes:

- Oscilloscope: Rohde & Schwarz RTM3004
- Waveform Generator: Keysight EDU33212A

Testing the Preamps

One major component of the hardware design is the preamp design. There are two preamp designs, and it is important that they provide the specified amount of gain and that they handle their inputs properly, whether they are single-ended or differential.

The differential microphone preamplifier was tested with a differential input generated by the Keysight EDU33212A. The input signals were sine waves set at 1 kHz with a peak-to-peak value of 200 mV. One of them was 180 degrees out of phase. The circuit and oscilloscope reading are shown below.

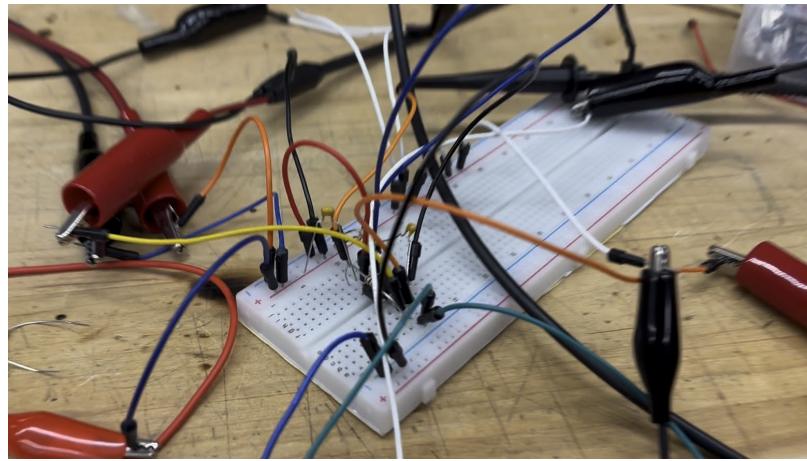


Figure 9.1: SSM2019 Preamplifier Circuit

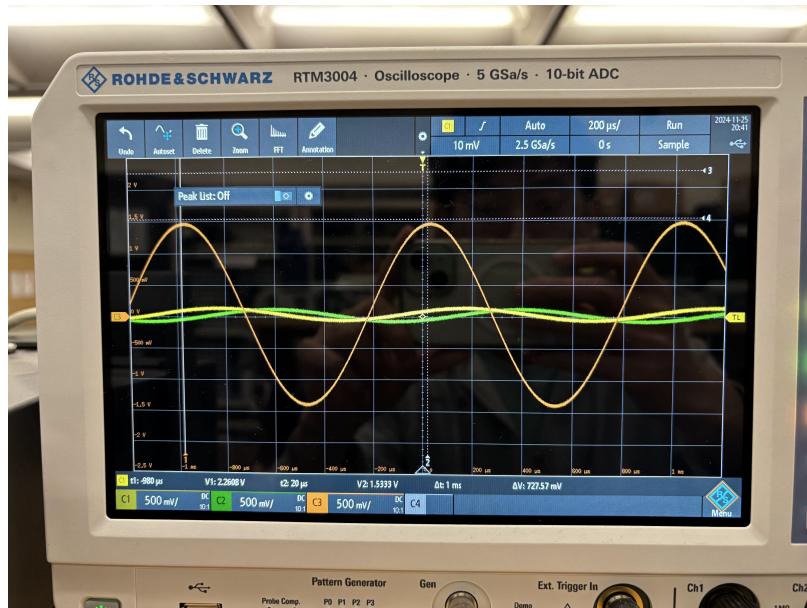


Figure 9.2: SSM2019 Preamplifier Circuit Output

The oscilloscope readings indicate a peak output of 1.5 volts which translates to a peak-to-peak rating of 3 volts. Compared to the 200 mV input, the output has been effected by a gain of approximately 15. This translates to a decibel gain of 23.5, which is close to the expected gain of 20 dB. The slight discrepancy is unaccounted for and remains to be explored. However, the output is clean and indicates a successful operation of the preamplifier circuit.

The high-impedance instrument amplifier was tested with the same waveform generator and same input signal, sans the second out-of-phase input signal. The circuit and oscilloscope reading are shown below.

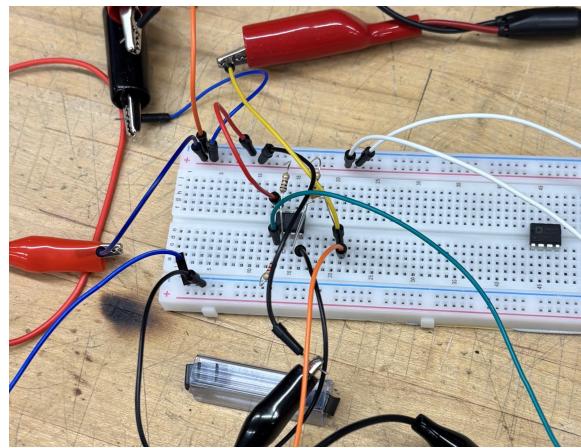


Figure 9.3: OPA2134 Preamplifier Circuit

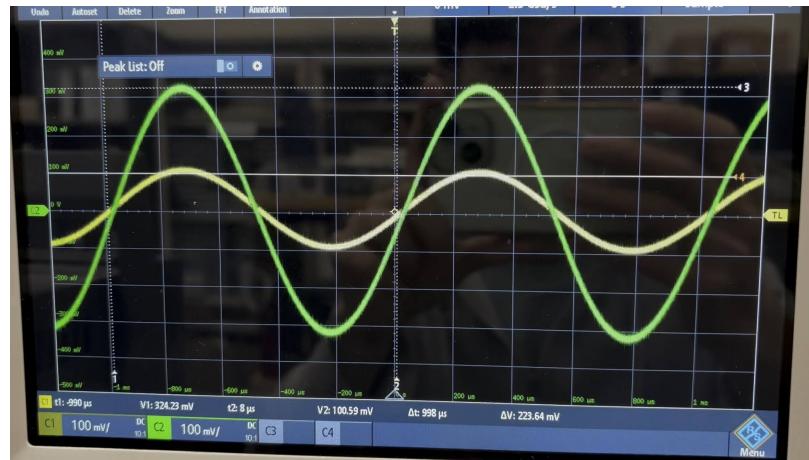


Figure 9.4: OPA2134 Preamplifier Circuit Output

The oscilloscope readings indicate a peak of approximately 320 mV which translates to a peak-to-peak rating of 640 mV. Compared to the 200 mV input, the output has been effected by a gain of approximately 3.2, which is exactly inline with the expected gain of the preamplifier setup. Overall, the output is clean and correct and indicates a successful operation of the preamplifier circuit.

9.2 Software Testing

Testing is a fundamental component of building any complex software system. It ensures the system functions as intended, especially in projects like ours, which features multiple software modules distributed across a complex hardware topology, along with stringent real-time requirements. Rigorous testing is critical for our being able to guarantee reliability, verify that the contracts of our modules are upheld, and save valuable development time.

9.2.1 Testing Methodology

Test-Driven Development

Our testing philosophy is informed by principles of extreme programming and test-driven development (TDD). We prioritize writing tests *before* implementing functionality, to ensure that each software module fulfills its design requirements. This philosophy aligns with the modular nature of our application: each developer is responsible for implementing a software process with a rigidly defined interface, allowing modules to be tested both independently and in integration.

By emphasizing the encapsulation of each developers' modules, we both ensure that implementation errors within modules are localized and quickly identifiable during development, and that the exposure of each modules' internals via their interfaces (and therefore the amount of testing required to verify each module) remains constrained. In this way, we create a foundation for confidence in the overall system's correctness upon the full integration of every module.

Automated Testing

In line with TDD, we aim to automate the validation of tests wherever possible. Automated testing accelerates development cycles by providing **consistent and rapid feedback** to developers. This is particularly necessary given the complexity of our system – the continuous validation of the behavior of individual modules and their interactions under a variety of simulated edge-conditions will be invaluable to ensuring correctness. Automation also ensures that our testing pipeline remains scalable as the project evolves, as our modules' internals and interfaces grow in complexity.

Code Coverage

We hope to achieve comprehensive code coverage across all testing layers. This means ensuring that all, or at least most, code paths are executed during testing, minimizing the risk of untested edge cases. Particular attention is paid to critical components such as the state manager, communication protocols, and real-time processing workflows, where undetected errors could have significant downstream effects.

The Testing Pyramid

Our testing strategy is structured around the ubiquitous testing pyramid strategy, which organizes tests into three layers, each layer encompassing tests which are increasingly larger in scope:

Unit tests Validate the smallest components of our system in isolation, ensuring that individual functions and methods behave correctly.

Integration tests Ensure that interactions between modules, such as state synchronization and communication protocols, function according to their design contract.

System tests Evaluate the complete system under realistic and edge conditions, testing end-to-end functionality, including hardware-software integration and real-time behavior.

Unit tests will likely comprise the majority of the tests we write, owing to their ease of implementation and execution. With a strong foundation of unit tests validating the behavior of individual modules, we reduce the risk of errors propagating to higher levels of testing.

Need for Testing of Real-Time Reliability

The real-time, distributed nature of our software is such that a high degree of correctness is required to ensure that we can make guarantees about the reliability and correctness of our system. In particular, this means that special attention must be given to testing our system for **non-deterministic behavior**.

We must specifically test sections of our codebase where concurrency or asynchrony introduces the possibility of race conditions. This means validating the thread safety and proper synchronization in shared resources, particularly in the state manager and communication modules, as well as testing for timeouts and communication errors, especially over physical-layer protocols like SPI and WiFi. Every module should be tested in a variety of simulated workloads and timing variations to ensure predictable behavior under all conditions.

Through the rigorous testing of non-deterministic scenarios, we can be assured that our system will remain robust and correct, even under highly dynamic conditions, and that the risk of unexpected exceptional states causing a failure of our system remains minimal.

9.2.2 Concrete Implementation of Testing

Whereas the previous section provided an overview of our testing methodology, the following sections will delve into the tools, frameworks, and specific techniques we will use to implement and execute tests at various levels on the testing pyramid. Our tests will span the full spectrum of our system's levels of complexity, from individual unit tests to full system stress tests on our physical hardware. We

will cover our intended testing strategies for each software module, and then the holistic behavior of the system.

State Manager

In testing our state manager module, we want to ensure that it meets its core design contracts of providing atomic, transactional modification to the central state, properly mediating the publisher-subscriber relationship between the remainder of our software modules, and providing robust communication with the xcore unit over I²C. In addition to unit and integration testing, we will use hardware interface validation to thoroughly vet the module's functionality.

To implement unit tests, we can use the testing tools provided by Rust's ecosystem. These range from the stock cargo test functionality, supplemented by libraries such as `mockall` for dependency mocking, and `tokio::test` for asynchronous testing. These tools will allow us to simulate concurrent operations and validate the system's ability to handle them safely even under high contention and exceptional circumstances.

The `loom` library will be used to model multithreaded behavior and exhaustively explore execution orderings. By injecting artificial delays and reordering events during tests, we can simulate real-world timing variances and ensure that the system gracefully handles out-of-order operations. This will ensure that the state manager consistently mediates the publisher-subscriber design pattern.

Hardware interface testing will focus on ensuring that the state manager's I²C connection with the Xcore board operates reliably. For this, we will employ the `rppal` library to interact with the Raspberry Pi's SPI and GPIO interfaces during automated tests. We will also perform manual SPI validation using `python-spi` to cross-check signal integrity and data correctness. These tests will verify the successful transmission of mixer state data to the DSP module, as well as the ability to handle communication interruptions or hardware-level anomalies gracefully.

Flutter User Interface

The testing strategy for our Flutter UI is designed to ensure its reliability, efficiency, and compliance with its functional requirements as a module within the larger system. Given the critical role the UI plays as the primary interaction point for the user, our tests will focus on validating both its visual behavior and its ability to communicate effectively with the backend manager processes. To achieve this, we will employ a combination of unit tests, widget tests, and integration tests, taking advantage of Flutter's mature testing ecosystem.

For unit testing, we will use Flutter's built-in `test` package to validate the behavior of individual UI components, such as state models, data transformations, and utility functions. These tests will ensure that the internal logic of the UI operates as intended and remains consistent across various input scenarios. State-driven UI elements will be tested to confirm that updates to the state manager are correctly reflected in the interface.

Widget testing will verify the functionality of individual interface elements, such as buttons, sliders, and visual indicators. Using Flutter's `flutter_test` package, we will simulate user interactions with these widgets to ensure that they respond appropriately and trigger the correct state changes. These tests will also confirm that the widgets render correctly across different screen sizes and orientations, maintaining the responsiveness required for use on any tablet device.

Integration tests will provide end-to-end validation of the entire UI, ensuring that it interacts correctly with the backend processes. Using Flutter's `integration_test` package, we will be able to simulate longer sequences of user interactions. These tests will also evaluate the UI's responsiveness and performance under typical usage conditions, ensuring that latency and resource consumption remain within acceptable limits, and that there are no components of the UI with unexpectedly high computational complexity.

Automixing Evaluation

A structured listening test will focus on evaluating the mixes generated by the Automix model compared to baseline approaches and a ground truth reference. The evaluation will be conducted within our group of four participants, ensuring consistency and alignment in subjective assessments.

Each participant will listen to a set of audio passages mixed using three different methods: the DMC implementation, a baseline approach (e.g., mono mix or random gain and panning), and a manual ground truth mix. The audio passages will be presented in a randomized order to minimize bias, and participants will rate each mix on a numerical scale (e.g., 1 to 5) based on the specified criteria. To ensure meaningful results, the test will use unseen samples from MedleyDB, featuring a variety of instruments and styles.

The evaluation will include scenarios ranging from simple gain and panning adjustments to more complex mixes involving EQ and compression. Each scenario will be repeated for multiple audio tracks to account for consistency in ratings. The results will be averaged across participants. This structured test will provide an informed, qualitative evaluation of the DMC's mixing capabilities relative to established benchmarks.

9.3 Overall Integration

9.3.1 Hardware Integration

Hardware integration in Senior Design 2 will involve the integration of several key components. Integration will involve the power PCB, overall PCB with MCU and peripherals, hardware enclosure, as well as the Raspberry Pi and its LCD Touchscreen. The plan is to test all the components individually and then integrate. For the most part, the physical integration of the components is facilitated through pin headers and jumpers. One large point of integration is that integration between the

xcore MCU and the Raspberry Pi. Given that there is not a designated hardware connection between the two devices, it is necessary to ensure a secure connection even with just jumper cables. Part of the challenge of integration is the design of the hardware enclosure and layout such that it facilitates solid connections between the components.

9.3.2 Software Integration

Our integration strategy for the software side of the mixer is fundamentally informed by the interfaces and design contracts between the various software modules. To streamline integration, each team member responsible for a software module will first produce a minimum viable product of the module they are responsible for. These MVPs will focus on establishing the basic communication patterns between modules. Establishing these communication protocols early provides a foundation for further development, ensuring that subsequent enhancements and features can build off the pre-existing, rudimentary integration. In addition, the key architectural design patterns structuring the relationship between our software modules, such as the publisher-subscriber relationship between state-producing modules and the manager-worker relationship between the Raspberry Pi and the xcore, will also be implemented early into Senior Design 2. This phased approach to integration will assist in rapid iterative development, minimizing the risk of compatibility issues during later stages of the project.

Chapter 10

Administrative Content

The following chapter covers the administrative content related to the mixer project. Topics include budget estimates, timelines, and division of labor. Numbers and figures are still tentative at this stage of the project.

10.1 Budget Estimates

The budget estimate for our project is shown in Table 10.1 below. The parts are nearly finalized, however, some slight modifications may still be made based on further research.

Item	Quantity	Price Per Unit	Adjusted Price
Raspberry Pi 5	1	\$80	\$80
Raspberry Pi Touchscreen	1	\$60	\$60
Raspberry Pi Display Cable	1	\$1	\$1
Mean Well RT-65B	1	\$25	\$25
PCM1865 ADC	2	\$3	\$6
TLV73333PDBVR LDO	1	\$1	\$1
TPS73618DBVTG4 LDO	1	\$2	\$2
TLV71209DBVT LDO	1	\$1	\$1
NE5532AD Op Amp	2	\$1	\$2
LT8338 Boost Converter	1	\$6	\$6
Si5351A-B-GT Clock IC	1	\$3	\$3
AD SSM2019 Op Amp	4	\$5	\$20
TI OPA2134 Op Amp	2	\$5	\$10
PCM1780 DAC	1	\$2	\$2
TPA6110A2DGN Headphone Amp	1	\$2	\$2
PCB Estimated costs	1	\$50	\$50
Metal Casing	1	\$20	\$20
XU316-1024-TQ128 (DSP Chip)	1	\$15	\$15
XLR Combo Jack Connector	4	\$4	\$16
1/4" TRS Jack	4	\$1	\$4
Caps/Inductors/Resistors	100	\$0.1	\$10
Estimated Total:			\$336

Table 10.1: Estimated Cost of Production

The estimated cost of the digital mixer is roughly competitive with the sale price of several competitor products such as the Behringer Flow 8 (\$299) and the Mackie MobileMix 8 (\$229). Of course, the price of the components is not an entirely fair comparison to make with the off-the-shelf price of a developed product. However, it is encouraging to the team that the cost of the project itself does not significantly exceed the price of a typical mixer in this category.

10.2 SD1 Documentation timeline

Below is the timeline for tasks within Senior Design 1. Senior Design 1 is strictly related to research and theoretical design, and as such, the tasks are related to the SD1 document. Timelines are relatively loose and are more strictly dictated by the actual submission timelines posted online in Webcourses.

Task	Start Week	Est. Duration (Weeks)
Brainstorming/Preliminary Research	Start of Summer	12
Individual Research	8/19/2024	1
Finalize project scope	8/23/2024	1
Divide and Conquer	8/26/2024	2
30 page milestone	8/19/2024	3
60 page milestone	9/9/2024	3
90 page milestone	9/30/2024	3
120 page milestone	10/21/2024	3
Final draft milestone	11/11/2024	3
Review	11/21/2024	0.5

Table 10.2: Project Documentation Milestone

10.3 SD2 Design Timeline

Below is the timeline for tasks within Senior Design 2. Senior Design 2 contains the development and prototyping phase of the project, and the tasks reflect that. The tasks and timeline may be updated with further research. It may be wise to move the timelines up considerably to account for unforeseen issues in hardware design and development.

Task	Description	Start Week	Est. Duration (Weeks)
Interfaces between work units	Establish contracts between individual modules	1/6/2025	1
Power supply	Test proper power delivery to hardware components	1/13/2025	2
Mixer state protocol	A standardized protocol for inter-module state communication is drafted	1/20/2025	1
Audio work loop	Audio successfully passes through xcore unit	1/27/2025	3
Basic onboard GUI	Minimal GUI implemented for testing purposes	2/3/2025	2
Communications co-processing	Developed successful communication schematic for WiFi and Bluetooth capabilities	2/24/2025	3
Machine learning model	Successfully predicts optimal mixing parameters for audio inputs to enable real-time audio mixing	2/10/2025	4
Control & DSP integrated	Raspberry Pi can send control signals to xcore	3/3/2025	3
Compute Integration	xcore and Raspberry Pi properly communicate with each other through appropriate protocols.	3/31/2025	4

Table 10.3: Project Design Milestone

10.4 Work Distributions

Work distributions among the five team members are delineated in this section. Firstly, Michael Meyers is the hardware lead in the digital mixer project. He is responsible for schematic design in areas such as voltage regulation, AD/DA conversion, and power amplification. He is also responsible for the design of the PCB.

Joshua Yu is the system architect and integration lead as well as the DSP effects lead. He is responsible for the overall vision of the project as well as the integration of hardware and software systems. He acts as a liaison between the hardware and software members of the team. In terms of his specific technical responsibilities, he is responsible for schematic design for both the MCU itself (GPIO, clock source, flash, etc.) as well as the preamplification circuitry. He is also assisting Michael Meyers in PCB level design. Furthermore, he is writing the DSP effects library.

Juni Yeom is the lead software developer, responsible for the design and implementation of the system architecture. He is overseeing the integration of the disparate software work performed by team members into a cohesive system. He is tasked with implementing the central data management process to handle the

mixer's control and audio data, in collaboration with Sri Kankipati. Additionally, he leads the development of the user interface, ensuring intuitive control over mixer functions and real-time audio visualization.

Srinilai Kankipati is the device communications lead. He is responsible for the standardization of data transfer between both inter-device modules and external sources. He is also responsible for developing and configuring the interconnects between said devices. In terms of specific technical responsibilities, he is responsible for inter-device serial communication, the data transfer formats, and the development of the USB audio interface capabilities. Additionally, he is support for Juni Yeom, particularly in the realm of central state management and real-time performance. Finally, he is support for Joshua Yu in writing the DSP library.

Kevin Kurian is the auto-mix algorithm lead. He is responsible for researching state-of-the-art auto-mix machine learning models, creating models, and integrating them with the system. In terms of specific responsibilities, he is tasked with reading literature on the topic, researching and curating datasets, recreating and pruning model architectures for integration on embedded platforms, and evaluating model performance in the context of the mixer project.

A short summary table of responsibilities is listed below.

<u>Task</u>	Primary Person	Secondary Person
Hardware Design	Michael Meyers	Joshua Yu
DSP Design	Joshua Yu	Michael Meyers
Auto-mix Algorithm	Kevin Kurian	Srinilai Kankipati
Software Design	Juni Yeom	Srinilai Kankipati
Inter-device Communications	Srinilai Kankipati	Juni Yeom

Table 10.4: Duty Distribution

Chapter 11

Conclusion

Thus far, the project has been a rigorous exercise in hardware design, software design, and creating ideas which provide utility and creative benefit to musicians. The work has been challenging but rewarding, and much knowledge has been ac-

quired in the realms of audio hardware, MCU integration, MCU firmware, software design, and wireless networking.

11.1 Progress Report

The work done in Senior Design 1 sets the team up for success in Senior Design II. As it stands, substantial work has been done in the realms of both hardware and software. For hardware, the schematic has been finalized and the PCB is in the design process. The design requirements have been well established and the team is on track to have a manufactured PCB before the start of Senior Design 2.

Significant progress has been made in realizing our mixer's software as well. After a substantial amount of research and comparisons, we have fully realized our software design, ranging from our high-level architecture, down to specific implementation details (regarding choice of frameworks, protocols, design patterns, etc.). We have constructed working prototypes of most of our core software modules, including our user interface and state manager process, and have achieved limited integration between them. At this point, our remaining development time during Senior Design 2 can be fully devoted to the rapid implementation of our design.

11.2 Lessons Learned

11.2.1 Experience in Architecting Complex Systems

Working on a "greenfield project" such as our mixer provided invaluable experience in how to design complex systems from scratch. We learned how to decompose a complicated system with a diverse requirements into a clear architectural design, which we further broke down into actionable, individual tasks. We gained confidence in our ability to solve complex challenges utilizing our individual subject expertise and engineering mindset. Throughout our research process, we each learned how to immerse ourselves in the landscape of pre-existing work relevant to our sections of the project, informing our design decisions.

11.2.2 Importance of Modularization

Modularization has proven to be an essential principle, not only for software design but for managing any complex engineering project. We developed a mindset of treating each member's contributions as discrete modules: clearly defining what each module should accomplish and how it should interface with other modules. Modularization assisted greatly in managing our roles, responsibilities, and deliverables throughout our hardware and software design process. In addition, by limiting the exposure of the internal details of our modules, we reduced the amount of information that needed to be communicated to only the details of their interfaces, making our overall development process much more efficient.

11.2.3 Effective Communication

The value of the skill of effective communication was made very clear during our development process. Throughout the semester, we became increasingly competent at sharing exactly the information needed to ensure that interconnected components, developed by different members, functioned cohesively. Again, modularization was key to achieving this clarity, as it defined the boundaries and responsibilities of each team member's work. This ensured that we neither overlapped in our efforts nor left critical tasks unaddressed. Furthermore, it minimized the need for extensive back-and-forth during development, reducing deadlock in our development process.

Chapter A

Citations

Bibliography

- [1] E. T. Chourdakis and J. D. Reiss, “A machine-learning approach to application of intelligent artificial reverberation,” *Journal of the Audio Engineering Society*, 2017.
- [2] Behringer, “Behringer flow 8 8-channel digital mixer,” <https://www.gear4music.com/us/en/PA-DJ-and-Lighting/Behringer-FLOW-8-8-Channel-Digital-Mixer/3N67>, 2024.
- [3] Mackie, “Mackie mobile mixer,” <https://mackie.com/en/products/battery-powered/live-sound/Mobile-mix.html>, 2024.
- [4] Tascam, “Tascam model 16,” https://tascam.com/us/product/model_16, 2019.
- [5] Soundcraft, “Signature 22 mtk,” <https://www.soundcraft.com/en-US/products/signature-22-mtk>, 2015.
- [6] Audiostance, “The ultimate guide to audio connectors and cables,” <https://www.audiostance.com/audio-connectors-and-cables/>, 2020.
- [7] Schematron, “Xlr balanced female to male diagram,” <https://schematron.org/xlr-balanced-female-to-13-stereo-male-wiring-diagram.html>, 2018.
- [8] L. Fuston, “Understanding signal levels in audio gear,” <https://www.sweetwater.com/insync/understanding-signal-levels-audio-gear/>, 2022.
- [9] J. Vautin, “dbv, dbu, and dbm,” <https://jeffvautin.com/2009/06/dbv-dbua-and-dbm/>, 2009.
- [10] C. J. Heiduska, “What does line level mean?,” <http://www.ovnilab.com/articles/linelevel.shtml>, 2006.
- [11] R. Elliott, “Phantom power - what is it and how it works,” <https://sound-au.com/articles/p-48.htm>, 2021.

- [12] DJJules, “Build the four-channel ssm2019 phantom powered mic preamp,” <https://www.instructables.com/Build-the-Four-Channel-SSM2019-Phantom-Powered-Mic/>, 2018.
- [13] R. B. et al., *Signals and Systems*. Rice University, 2021.
- [14] Arrow, “Analog to digital converter types and basic functions,” <https://www.arrow.com/en/research-and-events/articles/analog-to-digital-adc-converter-types-and-basic-functions>, 2024.
- [15] A. Devices, “Understanding sar adcs: Their architecture and comparison with other adcs,” <https://www.analog.com/en/resources/technical-articles/successive-approximation-registers-sar-and-flash-adcs.html>, 2001.
- [16] A. Devices, “Sigma-delta adc tutorial,” <https://www.analog.com/en/resources/interactive-design-tools/sigma-delta-adc-tutorial.html>, 2024.
- [17] A. Devices, “Basic linear design,” <https://www.analog.com/media/en/training-seminars/design-handbooks/Basic-Linear-Design/Chapter6.pdf>, 2005.
- [18] J. Siau, “Interpreting thd measurements - think db not percent!,” https://benchmarkmedia.com/blogs/application_notes/interpreting_thd_measurements_think_db_not_percent, 2017.
- [19] M. C. Limited, “Fpga vs. microcontroller - what’s the difference?,” <https://www.mclpcb.com/blog/fpga-vs-microcontroller/>, 2023.
- [20] E. M. B. Consortium, “Eembc benchmark score viewer,” *Embedded Microprocessor Benchmark Consortium*, 2024.
- [21] beagleboard.org, “Design and specifications,” <https://docs.beagleboard.org/boards/beaglebone/ai-64/03-design-and-specifications.html#bbai64-design>, 2024.
- [22] Digikey, “Beagl-bone-ai-64,” <https://www.digikey.com/en/products/detail/texas-instruments/BEAGL-BONE-AI-64/21283880>, 2024.
- [23] Espressif, “Esp32-s3 series datasheet,” <https://www.espressif.com/sites/default/files/documents/Esp32-S3-Datasheet.pdf>
- [24] STMicroelectronics, “Stm32h747/757,” <https://www.st.com/en/microcontrollers-microprocessors/stm32h747-757.html>, 2024.
- [25] T. Instruments, “Tps7a53b 3a, low-input voltage, low-noise, high-accuracy, low-dropout (ldo) voltage regulator,” <https://www.ti.com/lit/ds/symlink/tps7a53b.pdf>, 2024.
- [26] T. Instruments, “Tps6282x 5.5-v, 1-a, 2-a, 3-a step-down converter family with 1

- [27] Renesas, “Linear vs. switching regulators,” <https://www.renesas.com/us/en/products/power-power-management/linear-vs-switching-regulators>, 2022.
- [28] S. Craze, “The pan law within your daw explained in detail,” <https://samplecraze.com/tutorials/the-pan-law/>, 2024.
- [29] GCore, “What is websocket? — how does it work?,” <https://gcore.com/learning/what-is-websocket/>, 2023.
- [30] M. contributors, “The websocket api (websockets),” https://developer.mozilla.org/en-US/docs/Web/API/WebSockets_API, 2024.
- [31] S. J. Bigelow, “What is rest api (restful api)?,” <https://www.techtarget.com/searchapparchitecture/definition/RESTful-API>, 2024.
- [32] O. Router, “What is mqtt? a practical introduction,” <https://www.opcrouter.com/what-is-mqtt/>, 2024.
- [33] Radware, “What is coap? understanding the constrained application protocol,” <https://www.radware.com/security/ddos-knowledge-center/ddospedia/coap/>, 2024.
- [34] M. Afaneh, “Bluetooth low energy (ble): A complete guide,” <https://novelbits.io/bluetooth-low-energy-ble-complete-guide/>, 2022.
- [35] C. Steinmetz, “Automaticmixingpapers: A collection of papers related to automatic mixing.” <https://github.com/csteinmetz1/AutomaticMixingPapers>, 2023.
- [36] B. De Man, J. Reiss, and R. Stables, “Ten years of automatic mixing,” *Workshop on Intelligent Music Production*, 2017.
- [37] F. Pachet and O. Delerue, “On-the-fly multi track mixing,” *Advances in Engineering Software*, 2000.
- [38] D. Moffat and M. B. Sandler, “Approaches in intelligent music production,” in *Arts*, vol. 8, p. 125, MDPI, 2019.
- [39] E. Perez-Gonzalez and J. Reiss, “Automatic gain and fader control for live mixing,” in *2009 IEEE Workshop on Applications of Signal Processing to Audio and Acoustics*, pp. 1–4, IEEE, 2009.
- [40] U. Rave, “Automixing and its applications in music production,” 2021.
- [41] E. Perez-Gonzalez and J. Reiss, “Automatic equalization of multichannel audio using cross-adaptive methods,” in *Audio Engineering Society Convention 127*, Audio Engineering Society, 2009.

- [42] B. De Man and J. D. Reiss, "A knowledge-engineered autonomous mixing system," in *Audio Engineering Society Convention 135*, Audio Engineering Society, 2013.
- [43] M. N. Lefford, G. Bromham, G. Fazekas, and D. Moffat, "Context aware intelligent mixing systems," in *Audio Engineering Society Volume 69*, Audio Engineering Society, 2021.
- [44] D. Koszewski, T. Görne, G. Korvel, and B. Kostek, "Automatic music signal mixing system based on one-dimensional wave-u-net autoencoders," *EURASIP Journal on Audio, Speech, and Music Processing*, vol. 2023, no. 1, p. 1, 2023.
- [45] B. Kolasinski, "A framework for automatic mixing using timbral similarity measures and genetic optimization," in *Audio Engineering Society Convention 124*, Audio Engineering Society, 2008.
- [46] E. M. Schmidt and Y. E. Kim, "Automatic multi-track mixing using linear dynamical systems," *Drexel University*, 2011.
- [47] A. L. Benito and J. D. Reiss, "Intelligent multitrack reverberation based on hinge-loss markov random fields," in *Audio Engineering Society Conference: 2017 AES International Conference on Semantic Audio*, Audio Engineering Society, 2017.
- [48] D. Giannoulis, M. Massberg, and J. D. Reiss, "Parameter automation in a dynamic range compressor," *Journal of the Audio Engineering Society*, vol. 61, no. 10, pp. 716–726, 2013.
- [49] A. Mason, N. Jillings, Z. Ma, J. D. Reiss, and F. Melchior, "Adaptive audio reproduction using personalized compression," in *Audio Engineering Society Conference: 57th International Conference: The Future of Audio Entertainment Technology—Cinema, Television and the Internet*, Audio Engineering Society, 2015.
- [50] I. Goodfellow, "Deep learning," 2016.
- [51] C. J. Steinmetz, S. S. Vanka, M. A. Martínez Ramírez, and G. Bromham, *Deep Learning for Automatic Mixing*. ISMIR, Dec. 2022.
- [52] C. J. Steinmetz and J. D. Reiss, "Deep learning for automatic mixing: Challenges and next steps." https://static1.squarespace.com/static/5554d97de4b0ee3b50a3ad52/t/618e84cf83e6fc75816e12bc/1636730122670/MDX_Workshop_at_ISMIR_Deep_learning_for_automatic_mixing.pdf, 2021. Workshop presented at ISMIR 2021, MDX Workshop.
- [53] M. Martinez Ramirez, D. Stoller, and D. Moffat, "A deep learning approach to intelligent drum mixing with the wave-u-net," in *Audio Engineering Society Volume 45*, Audio Engineering Society, 2021.

- [54] S. I. Mimirakis, K. Drossos, T. Virtanen, and G. Schuller, “Deep neural networks for dynamic range compression in mastering applications,” in *Audio Engineering Society Convention 140*, Audio Engineering Society, 2016.
- [55] S. H. Hawley, B. Colburn, and S. I. Mimirakis, “Signaltrain: Profiling audio compressors with deep neural networks,” *arXiv preprint arXiv:1905.11928*, 2019.
- [56] M. A. M. Ramirez and J. D. Reiss, “Modeling nonlinear audio effects with end-to-end deep neural networks,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 171–175, IEEE, 2019.
- [57] M. Martinez Ramirez, J. Reiss, *et al.*, “End-to-end equalization with convolutional neural networks,” *DAFx-2018*, 2018.
- [58] J. T. Colonel and J. Reiss, “Reverse engineering of a recording mix with differentiable digital signal processing,” *The Journal of the Acoustical Society of America*, vol. 150, no. 1, pp. 608–619, 2021.
- [59] D. Stoller, S. Ewert, and S. Dixon, “Wave-u-net: A multi-scale neural network for end-to-end audio source separation,” *arXiv preprint arXiv:1806.03185*, 2018.
- [60] C. J. Steinmetz, J. Pons, S. Pascual, and J. Serra, “Automatic multitrack mixing with a differentiable mixing console of neural audio effects,” in *ICASSP 2021-2021 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 71–75, IEEE, 2021.
- [61] M. A. Martínez-Ramírez, W.-H. Liao, G. Fabbro, S. Uhlich, C. Nagashima, and Y. Mitsufuji, “Automatic music mixing with deep learning and out-of-domain data,” *arXiv preprint arXiv:2208.11428*, 2022.
- [62] S. S. Vanka, C. Steinmetz, J.-B. Rolland, J. Reiss, and G. Fazekas, “Diff-mst: Differentiable mixing style transfer,” *arXiv preprint arXiv:2407.08889*, 2024.
- [63] L. Gabrielli, S. Tomassetti, S. Squartini, C. Zinato, *et al.*, “Introducing deep machine learning for parameter estimation in physical modelling,” in *Proceedings of the 20th international conference on digital audio effects*, 2017.
- [64] G. S. D. K. C. D. Jacob Larson, James S. Denford, “The evolution of artificial intelligence (ai) spending by the u.s. government,” <https://www.brookings.edu/articles/the-evolution-of-artificial-intelligence-ai-spending-by-the-u-s-government/>, 2024.
- [65] E. Griffith, “Investors pour \$27.1 billion into a.i. start-ups, defying a downturn,” *The New York Times*, 2024.
- [66] K. Houser, “Llms are a dead end to agi, says françois chollet,” *Freethink*, 2024.

- [67] M. McArdle, “Ai is coming for the professional class. expect outrage — and fear.,” <https://www.washingtonpost.com/opinions/2024/04/29/ai-professional-class-low-skill-jobs/>, 2024.
- [68] T. Corporation, “Input and output circuits for that preamplifier ics,” <https://thatcorp.com/datasheets/dn140.pdf>, 2024.
- [69] A. Devices, “Self-contained audio preamplifier ssm2019,” <https://www.analog.com/media/en/technical-documentation/data-sheets/SSM2019.pdf>, 2011.
- [70] JLCPCB, “Understanding analog and digital ground in pcb design,” <https://jlcppcb.com/blog/understanding-analog-and-digital-ground-in-pcb-design>, 2024.

Chapter B

Copyright

When using copyrighted content, we sought direct approval from the companies to follow fair use laws. As such, proof for used figures is shown below.

Approval for Figure 3.4:

Aug 22, 2015 at 5:49 history asked

Fraisse

CC BY-SA 3.0

You are free to:

Share — copy and redistribute the material in any medium or format for any purpose, even commercially.

Adapt — remix, transform, and build upon the material for any purpose, even commercially.

The licensor cannot revoke these freedoms as long as you follow the license terms.

Under the following terms:



Attribution — You must give [appropriate credit](#), provide a link to the license, and [indicate if changes were made](#). You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.



ShareAlike — If you remix, transform, or build upon the material, you must distribute your contributions under the [same license](#) as the original.

No additional restrictions — You may not apply legal terms or [technological measures](#) that legally restrict others from doing anything the license permits.

Approval for Figure 3.10:

⌚ Your attachment has been added successfully

Hi Michael,

Good Day.

Yes, you may cite the requested document for your project we only ask that you include "Courtesy of Texas Instruments Inc." under the image.

I hope this solution I have provided helps you. If any clarification particular to this request is needed, please reply to this email within the next 5 days and let me know before the case is closed by the system.

To close this case, just click the accept link together with this update or inform us that this case can be closed.

Best Regards,

Ray Vincent Aranzaso

Texas Instruments | Customer Support Center

Approval for Figures 3.5 and 3.8:

Hi Michael,

My name is Alyssa Scarpulla, and I am an attorney from ADI's Intellectual Property department.

Thank you for reaching out and asking for permission to use the two figures in your paper. This email serves as permission to use the figures in your paper on ADC technology, please include the following disclaimer under each image "Analog Devices, Inc. ("ADI"), © 2024. All Rights Reserved. These images are reproduced with permission by ADI."

Please let me know if you have any additional questions.

Best,

Alyssa Scarpulla

Associate Counsel, Intellectual Property

Email alyssa.scarpulla@analog.com



analog.com |

Approval for Figures 3.9 and 3.11:

Hello Michael,

Here is the response I got from a member of our legal team.

Michael can use those two graphics in his research paper as long as he includes the following attribution on each page that contains the material: "Reproduced with Permission. © Renesas Electronics Corporation or one of its subsidiaries. All Rights Reserved."

Best regards,
Cindy



10/11/2024 09:15 PM

Approval for Figure 3.3:



Phil Nass <phil_nass@sweetwater.com>

To: Michael Meyers



...

Tue 10/8/2024 5:56 PM

Retention: UCF Delete after 10 Years (10 years) Expires: Fri 10/6/2034 5:56 PM

Oh I gotcha. Yeah that's not a problem. I would just make sure you cite where it came from, but that's no worries. Its just a standard power diagram, so no worries.

...

Attempt at approval for Figure 3.7:

Approval for Figure 3.16:

Request for Image Usage in School Paper



Michael Meyers

To: microcontrollerslabhub@gmail.com



Thu 10/3/2024 1:44 PM



Retention: UCF Delete after 10 Years (10 years) Expires: Sun 10/1/2034 1:44 PM

To whom it may concern,

My name is Michael Meyers and I am involved in writing a Senior Design paper as my capstone project at the University of Central Florida (UCF).

I am reaching out to see if I am able to use the figure "Binary weighted resistor DAC circuit" (Found on [Binary Weighted Resistor DAC](#) ([microcontrollerslab.com](#))) in this paper as a demonstration of how the technology works.

Thank you in advance,

Michael Meyers

Re: Request for Image Usage in School Paper

Inbox x



miniDSP TechSupport <info@minidsp.com>

to me ▾

Hi Joshua

Thanks for your interest in our products. You are welcome to use the diagram of a biquad filter from our page

Have a nice day !

Hoping this info helps and feel free to contact us if you have further questions.

Best Regards
miniDSP DevTeam

Chapter C

LLM Citations

Prompt and Output for figure 5.1. Output from ChatGPT model 4o.

Can you generate an image of a schematic of a preamp circuit with the SSM2019 IC?

Here is the schematic of the preamp circuit using the SSM2019 IC as described. Let me know if you need any further details or adjustments!

Diagram:

Schematic Labels:

- Pin V
- SSM2019 SSM2019
- SSM2019 SSM2019
- Pin 2 IN-2 IN-2
- IN-2 IN-2
- IN-3 IN-3
- pin 3 Pin 3
- Pin 5 Pin 5
- IN + IN +
- 15V-copret -15V-copret
- Output coppét Output coppét
- 1n-15v -1n-15v
- PN- PN-
- Output -IN-0V Output -IN-0V
- COPAL COPAL
- IN-2 IN-2
- IN-3 IN-3
- 1N-5 -1N-5
- 1IN-1V -1IN-1V
- OUT OUT